



Relational Representation Learning with Graph Neural Networks

Matthias Fey

matthias.fey@tu-dortmund.de

   /rusty1s

Roadmap

1. Motivation & Introduction
2. Common Operators
3. Efficient Computation of Graph Neural Networks
4. Relation to the Weisfeiler-Lehman Algorithm
5. Recent Extensions
 - Hierarchical Models
 - Deep(er) Models
 - Unsupervised Models
 - Generative Models
6. Conclusion

Motivation & Introduction

Motivation

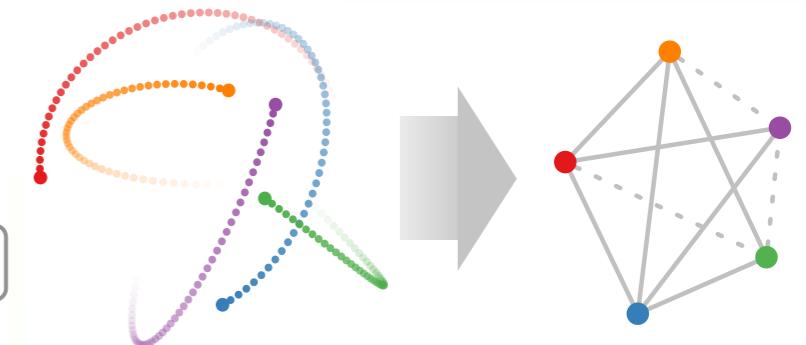
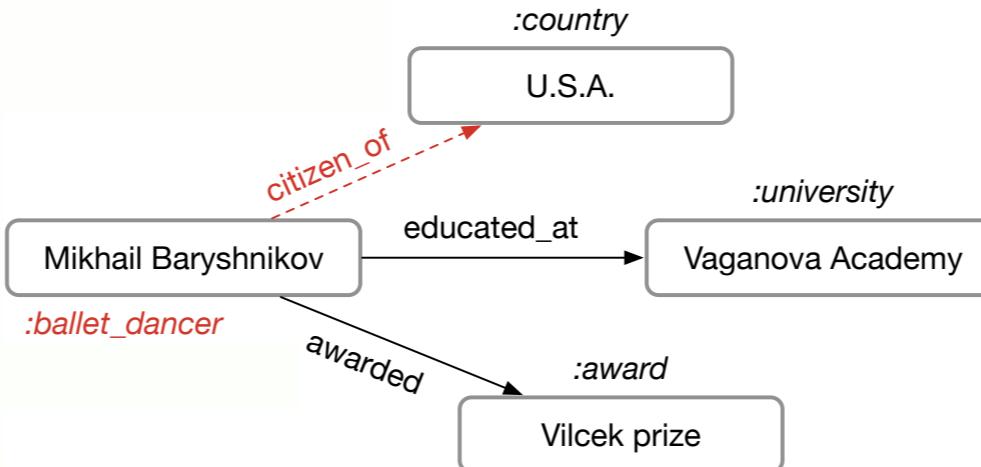
Enable learning on irregularly structured data, e.g., graphs, point clouds, and manifolds:

borrow from/generalize successful concepts on Euclidean domains, e.g., convolution

Enforce relational inductive biases into the model: enable learning about entities, relations, and rules for composing them

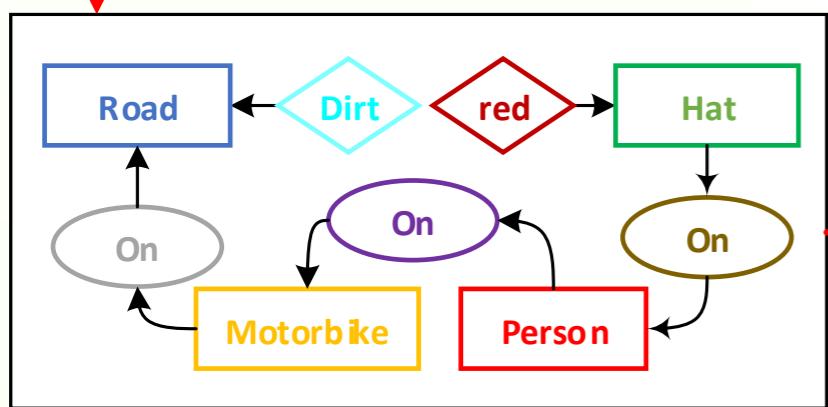
known as either **Geometric Deep Learning** or **Relational Representation Learning**

A Non-Exhaustive List of Applications



Dynamics in
Physical Systems

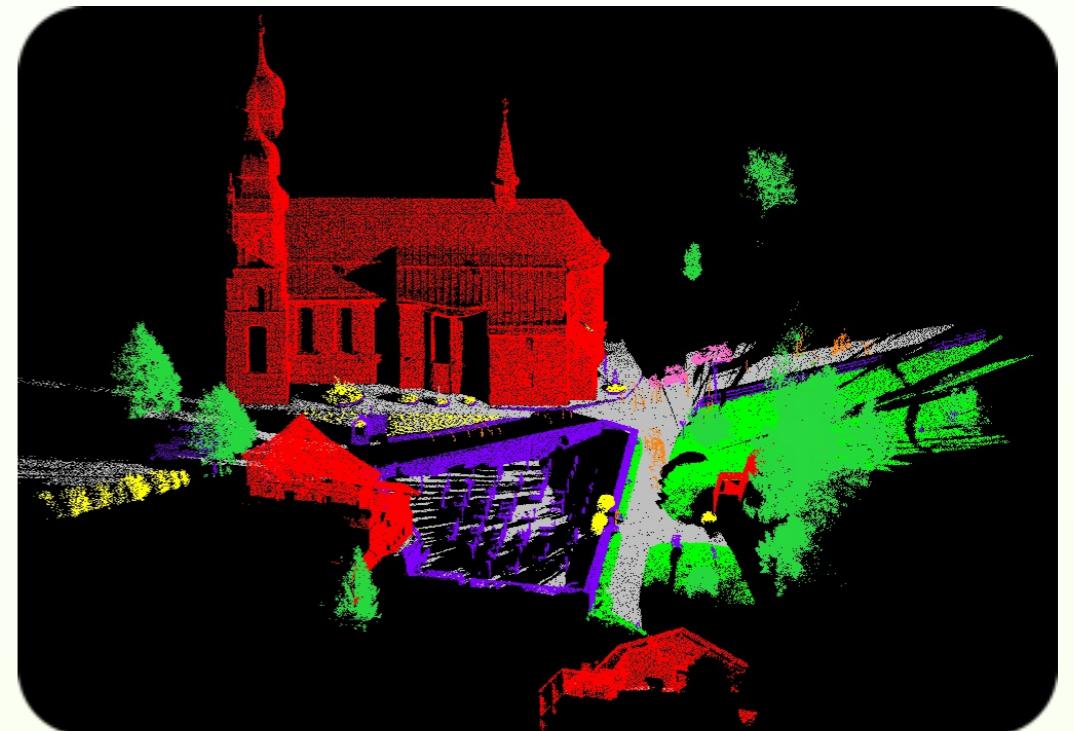
Reasoning in
Knowledge Graphs



Visual Scene Understanding



Autonomous Driving

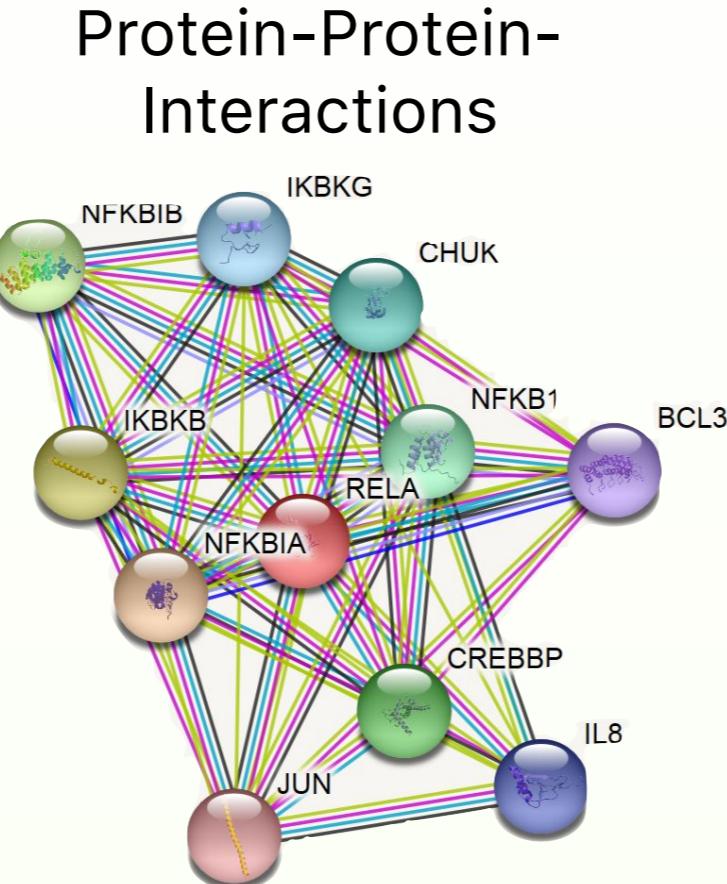


Point Cloud Analysis

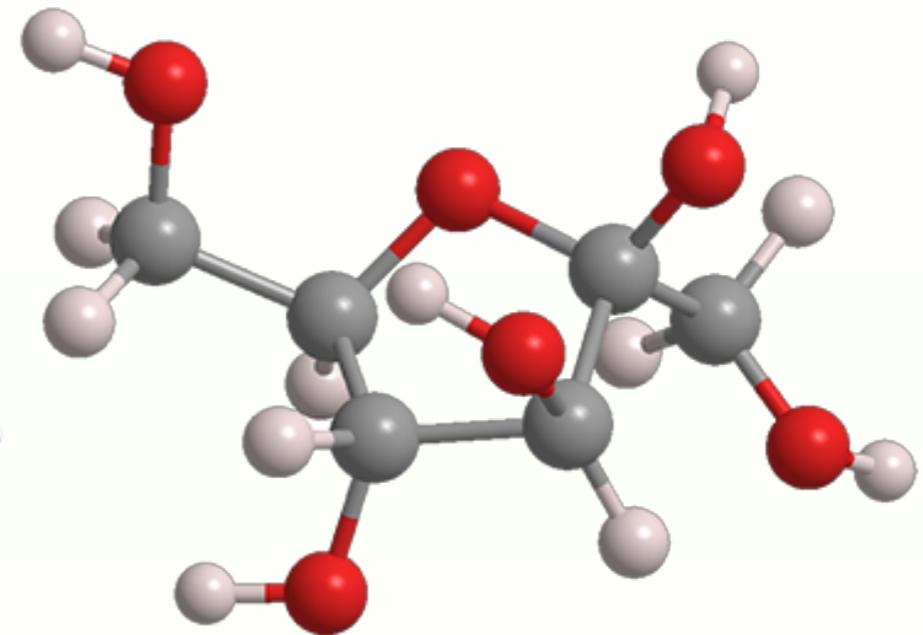
A Non-Exhaustive List of Applications



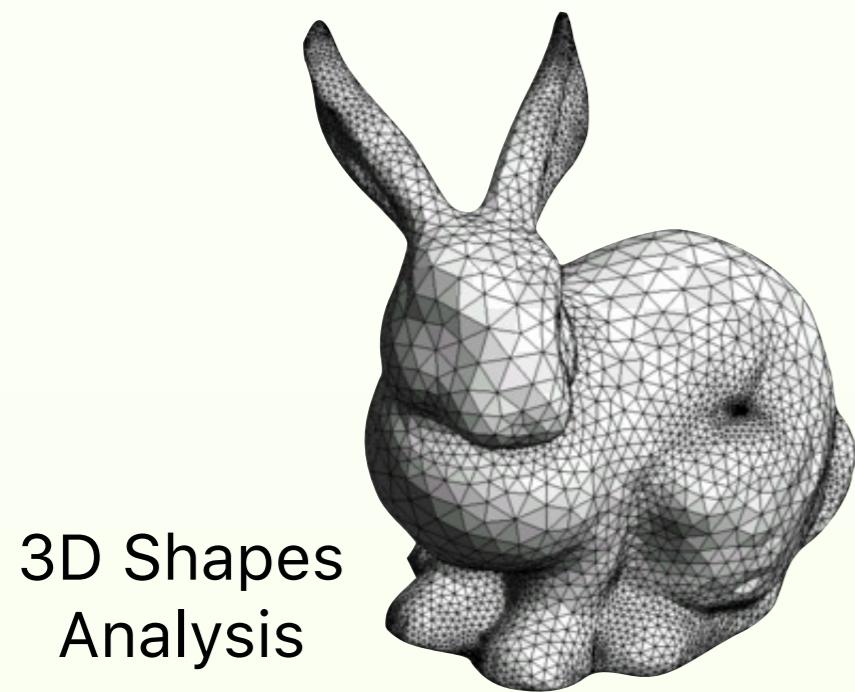
Social Networks



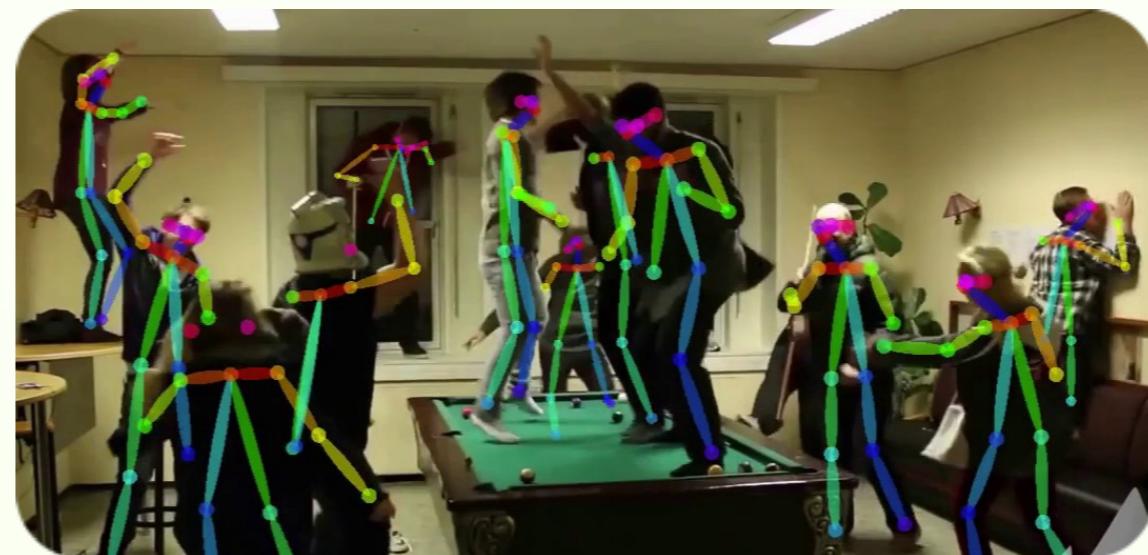
Protein-Protein-Interactions



Prediction of molecular properties
and drug discovery



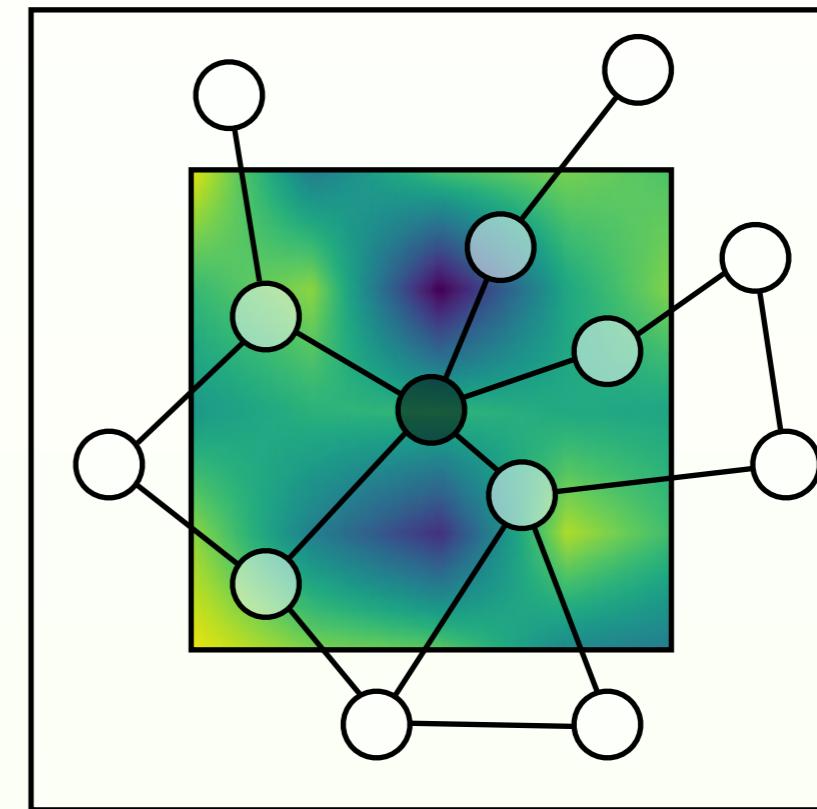
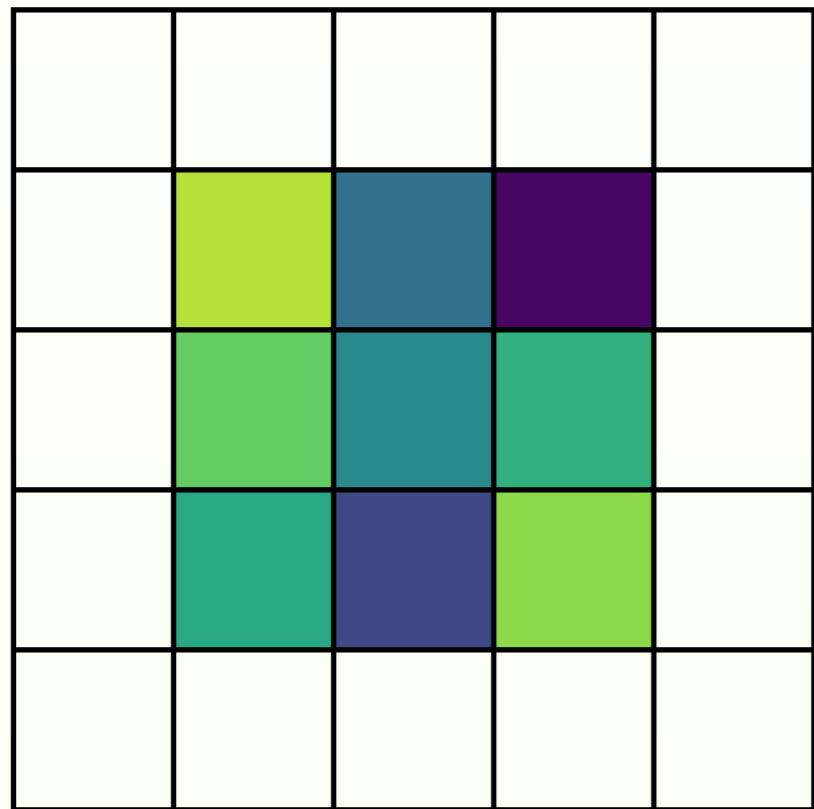
3D Shapes
Analysis



Human Pose Estimation

Graph Neural Networks (GNN)

Generalization of the convolution operator to graphs



dynamic positions

dynamic number of neighbors

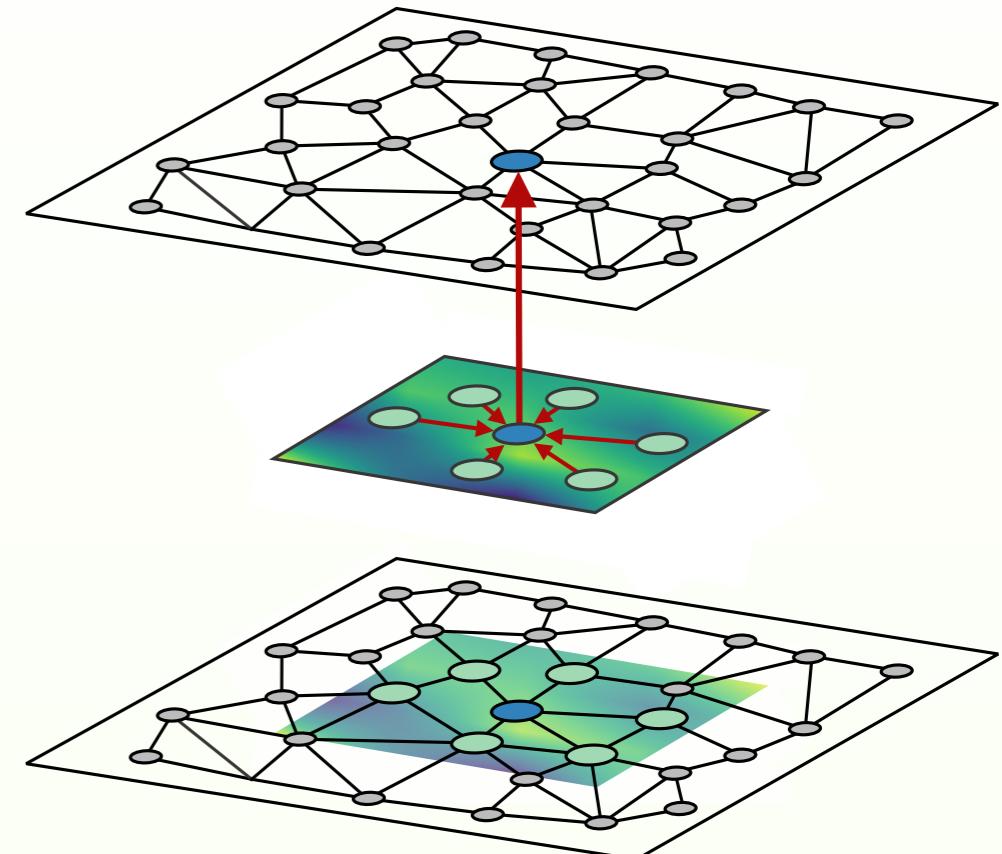
→ From discrete to continuous filters

Graph Neural Networks (GNNs)

Graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

Node features $\vec{h}_v^{(t)} \in \mathbb{R}^F$ in layer t

GNNs can be expressed as a neighborhood aggregation or message passing scheme:



$$\vec{m}_v^{(t)} = \text{AGGREGATE}^{(t)} \left(\left\{ \vec{h}_w^{(t-1)} : w \in \mathcal{N}(v) \right\} \right)$$

permutation-
invariant

$$\vec{h}_v^{(t)} = \text{COMBINE}^{(t)} \left(\vec{h}_v^{(t-1)}, \vec{m}_v^{(t)} \right)$$

Common Operators

Operators for Learning on Graphs

$$\vec{h}_v^{(t)} = \sigma \left(\Theta^{(t)} \sum_{w \in \mathcal{N}(v) \cup \{v\}} C_{v,w}^{(t-1)} \vec{h}_w^{(t-1)} \right)$$

Non-
Linearity Trainable
parameters Added
self-loops Normalization
coefficient

Normalization can be either ...

... static: $C_{v,w}^{(t)} := 1$

... structure-dependent: $C_{v,w}^{(t)} := |\mathcal{N}(v)|^{-1}$

... data-dependent, aka Attention

Xu et al.: How Powerful are Graph Neural Networks? (ICLR 2019)

Hamilton et al.: Inductive Representation Learning on Large Graphs (NIPS 2017)

Kipf and Welling: Semi-Supervised Classification with Graph Convolutional Networks (ICLR 2017)

Velickovic et al.: Graph Attention Networks (ICLR 2018)

Operators for Learning on Point Clouds

PointNet(++):

$$\vec{h}_v^{(t)} = \max_{w \in \mathcal{N}(v) \cup \{v\}} \phi_{\Theta}^{(t)} \left(\vec{h}_w^{(t-1)}, \vec{p}_w - \vec{p}_v \right)$$

↑
trainable
function, e.g., MLP

↑
relative Cartesian
coordinates

Neighborhood is typically given by ball query or
 K nearest neighbor search.

Formulation can be easily extended to handle
arbitrary edge features!

Qi et al.: Deep Learning on Point Sets for 3D Classification and Segmentation (CVPR 2017)

Qi et al.: PointNet++: Deep Hierarchical Feature Learning on Point Sets in a Metric Space (NIPS 2017)

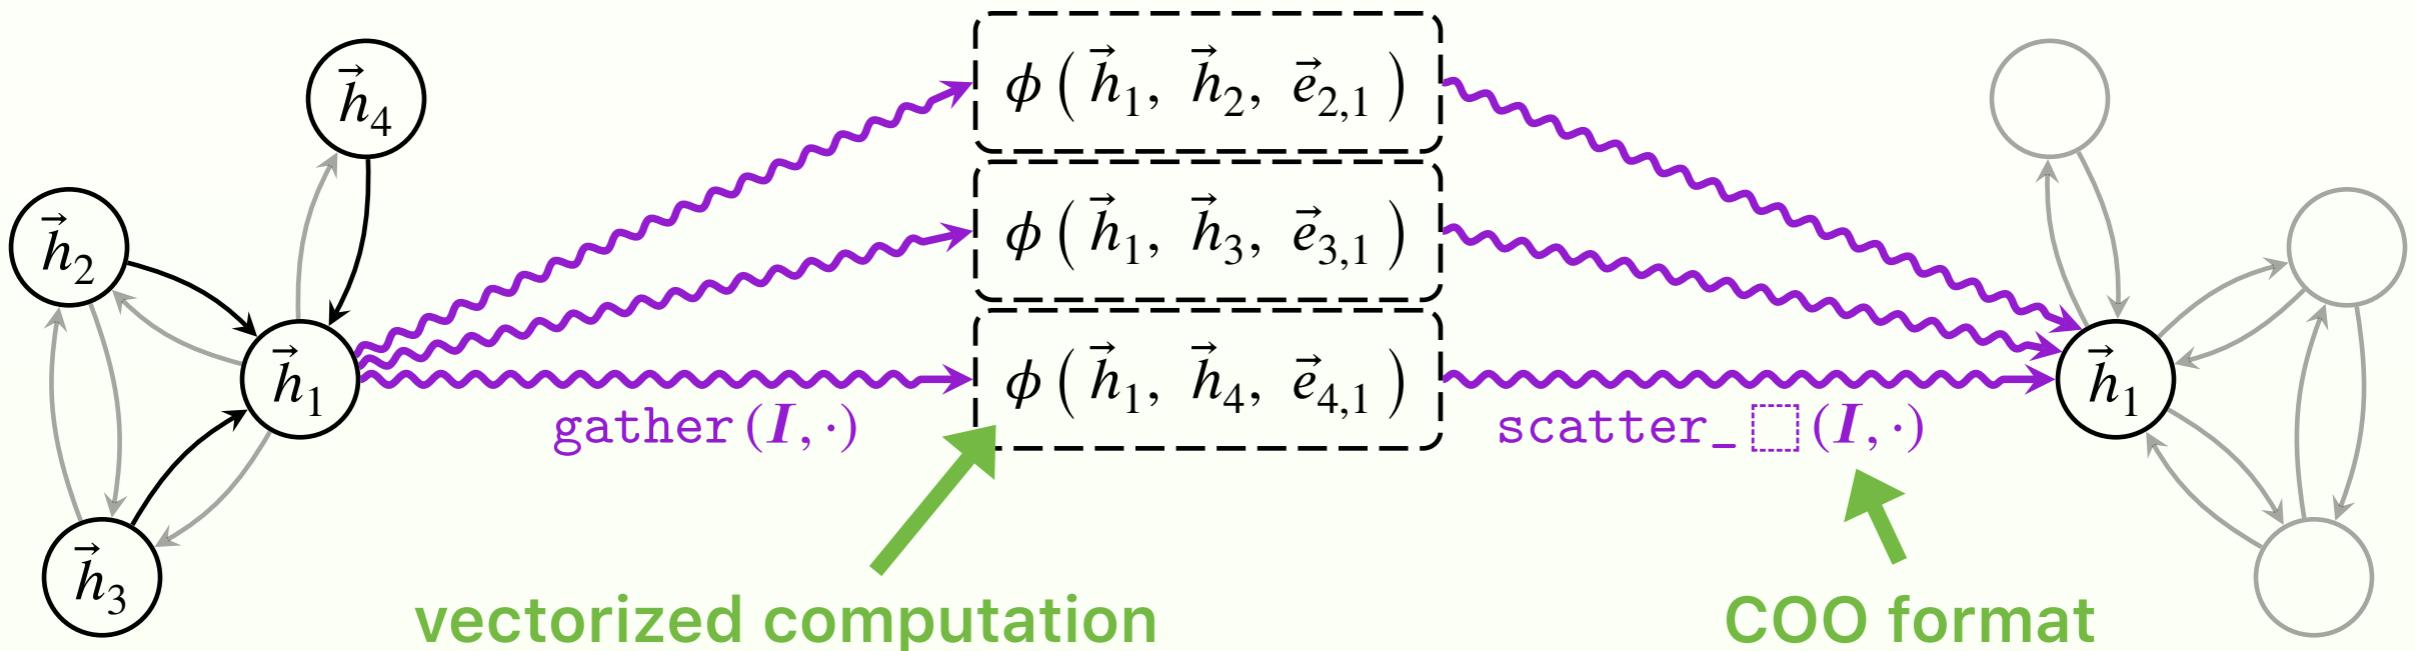
Fey et al.: SplineCNN: Fast Geometric Deep Learning with Continuous B-Spline Kernels (CVPR 2018)

Efficient Computation of Graph Neural Networks

Efficient Parallelization of GNNs

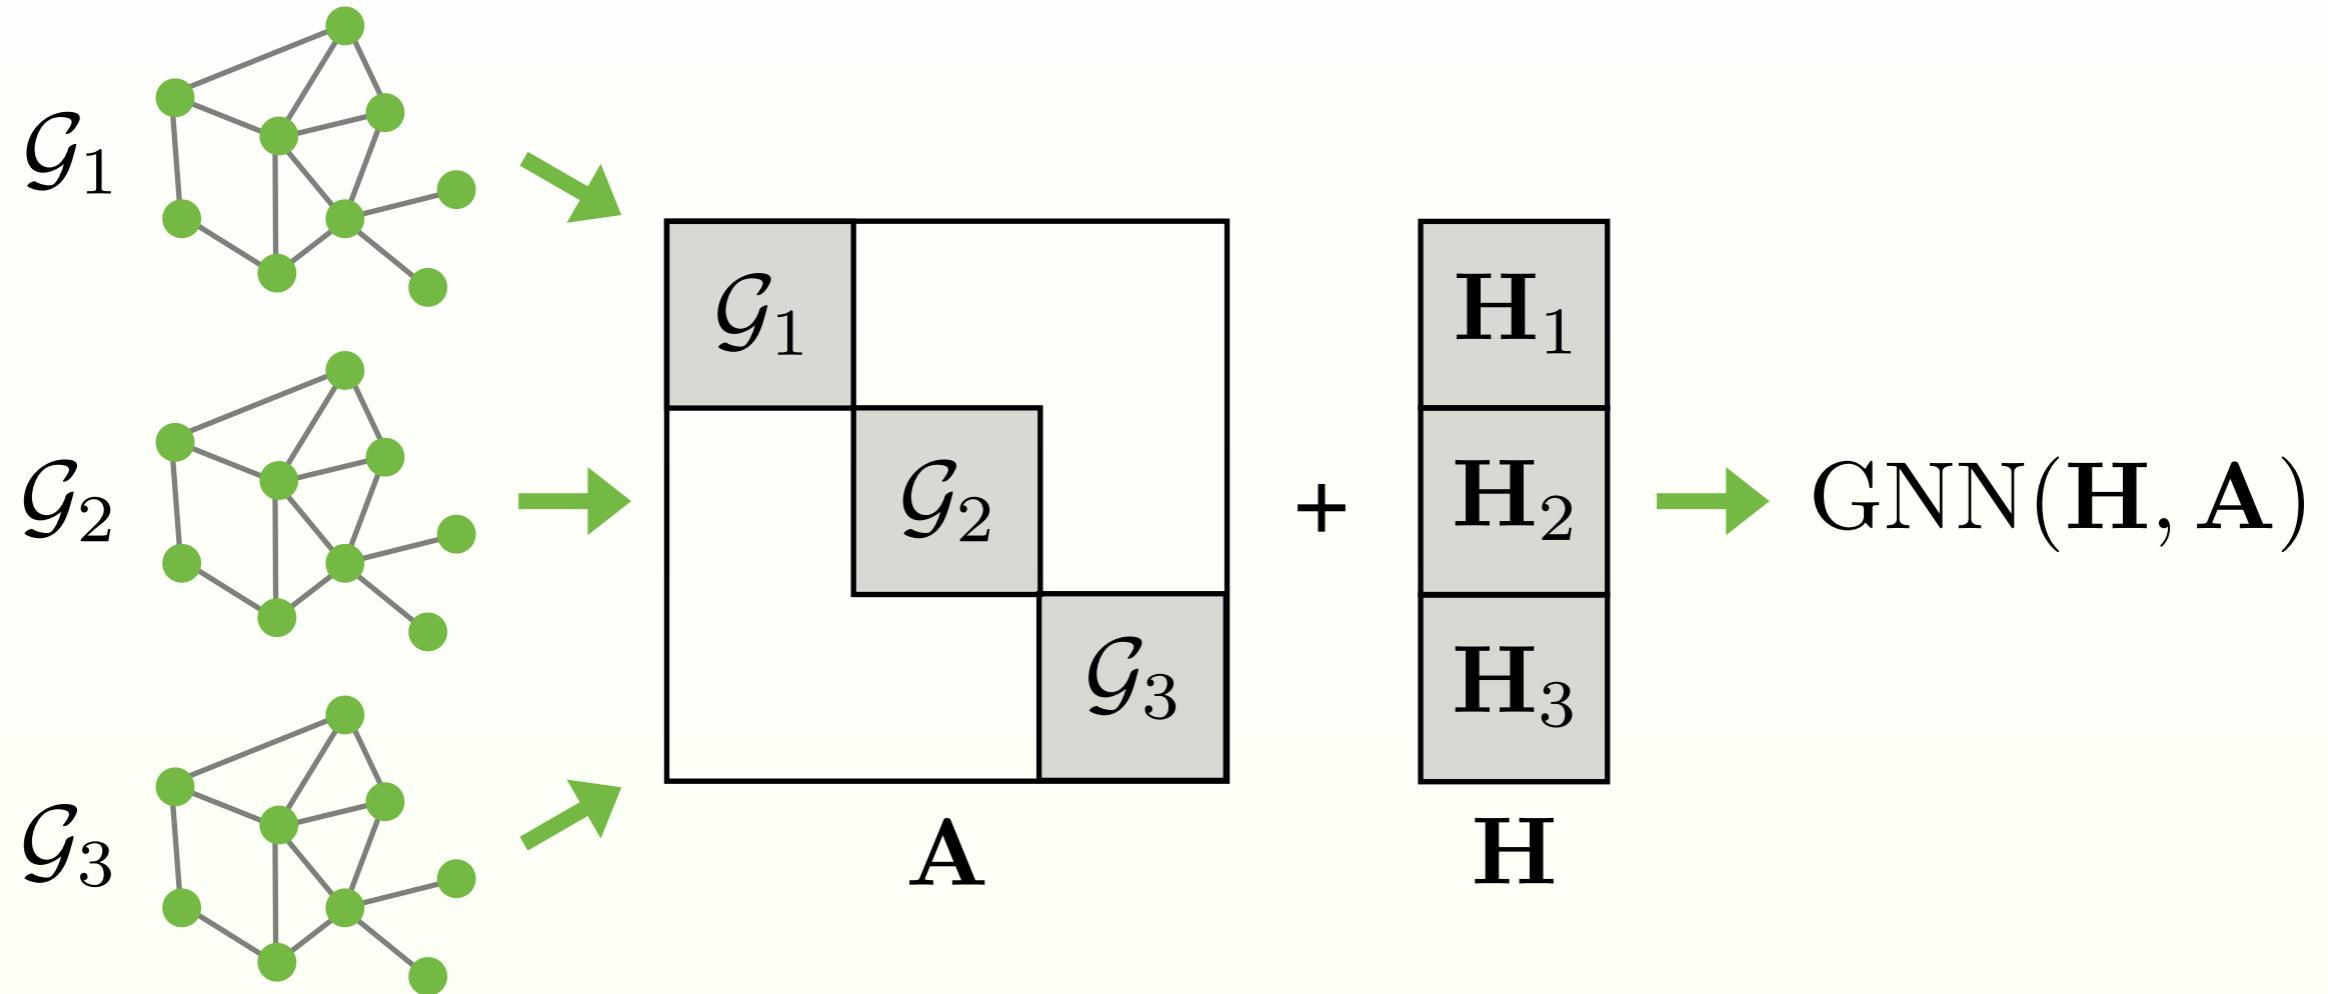
$$\text{AGGREGATE}^{(t)}(v) = \square_{w \in \mathcal{N}(v)} \phi_{\Theta}^{(t)} \left(\vec{h}_v^{(t-1)}, \vec{h}_w^{(t-1)}, \vec{e}_{w,v}^{(t-1)} \right)$$

\square is placeholder for either sum, mean, or max



1. Gather node information into edge parallel space
2. Scatter edge information into node parallel space

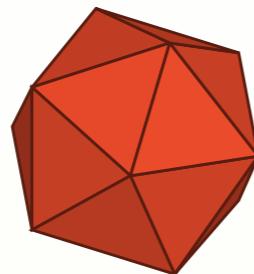
Mini-Batching in GNNs



No additional overhead due to sparse layout!

Input examples can have different size!

PyTorch Geometric: A Fast GNN Library

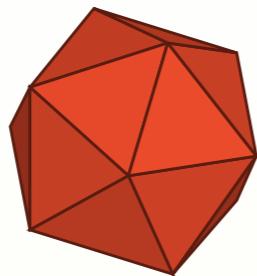


PyTorch
geometric

- ✓ provides over 25 GNN implementations
- ✓ access to over 100 benchmark datasets
- ✓ extendable via a simple Message Passing API
- ✓ leverages dedicated CUDA kernels
- ✓ supports multi-GPUs
- ✓ thoroughly documented

[https://github.com/rusty1s/
pytorch_geometric](https://github.com/rusty1s/pytorch_geometric)

PyTorch Geometric: A Fast GNN Library



PyTorch
geometric

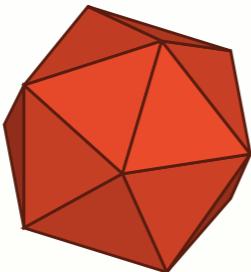
```
class MyOwnNet(Module):
    def __init__(self, in_channels, out_channels):
        self.conv1 = GCNConv(in_channels, 16)
        self.conv2 = GCNConv(16, out_channels)

    def forward(self, x, edge_index):
        x = self.conv1(x, edge_index)
        x = relu(x)
        x = self.conv2(x, edge_index)
        x = softmax(x, dim=1)
        return x
```

layer
initialization

network
execution flow

PyTorch Geometric: A Fast GNN Library



PyTorch geometric

```
class MyOwnConv(MessagePassing):
    def __init__(self, ...):
        super(MyOwnConv, self).__init__('add')add, mean or max
                                            ↗ aggregation

    def forward(self, x, edge_index):
        return self.propagate(edge_index, x=x)pass every-
thing needed to
construct messages

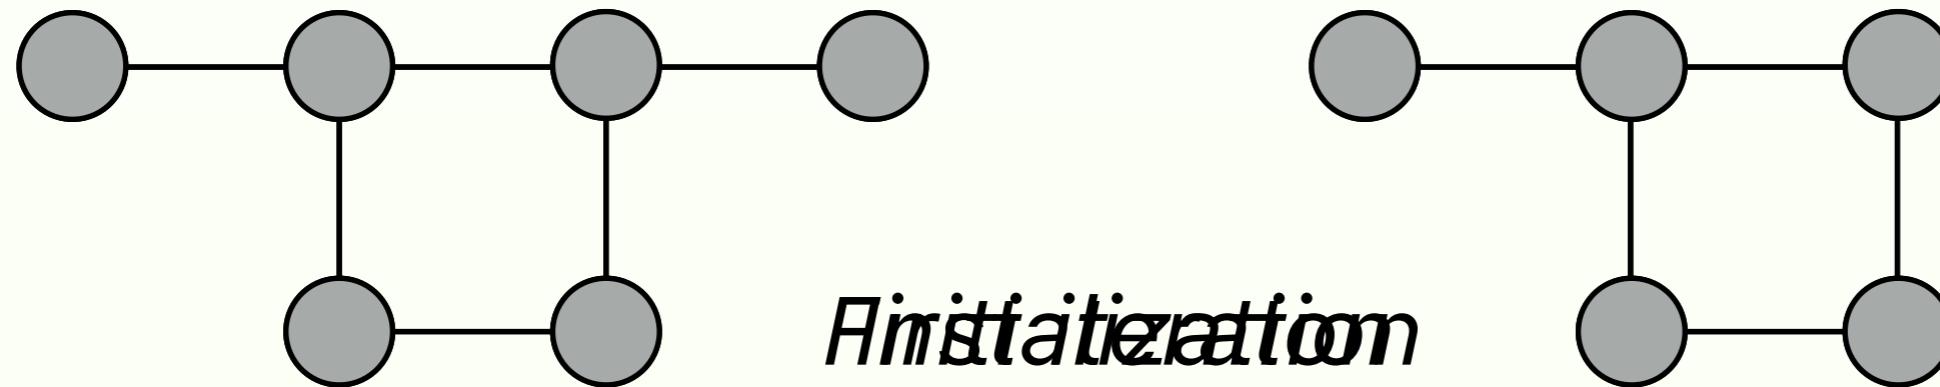
    def message(self, x_i, x_j):
        # Construct messages to aggregate.↑
Features get automatically mapped to
respective source and target nodes
```

Relation to the Weisfeiler-Lehman Algorithm

How Powerful are Graph Neural Networks?

WL-Test - A Graph Isomorphism Heuristic via iterative color refinement:

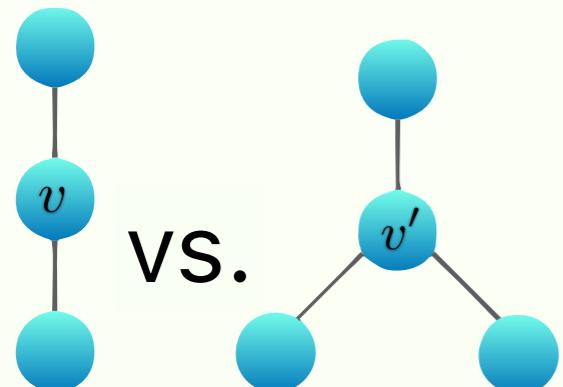
1. Aggregate colors from neighboring nodes
2. Hash aggregated colors into unique new colors



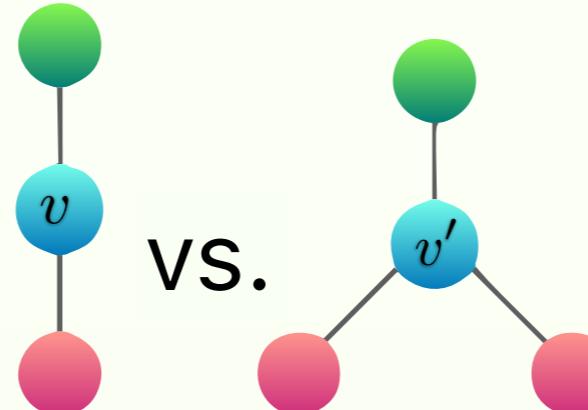
3. Graphs are non-isomorphic if color histograms differ after T rounds

How Powerful are Graph Neural Networks?

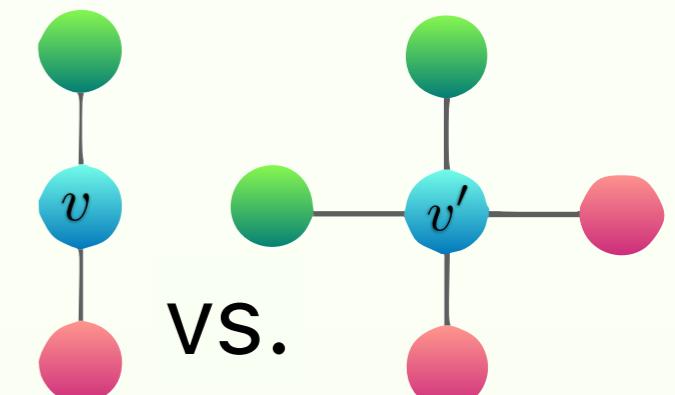
- GNNs are at most as powerful as the WL-Test in distinguishing graph structures
- GNNs are as powerful as the WL-Test *iff* AGGREGATE and COMBINE are injective
- Use sum instead of mean or max aggregation



mean/max fail



max fails



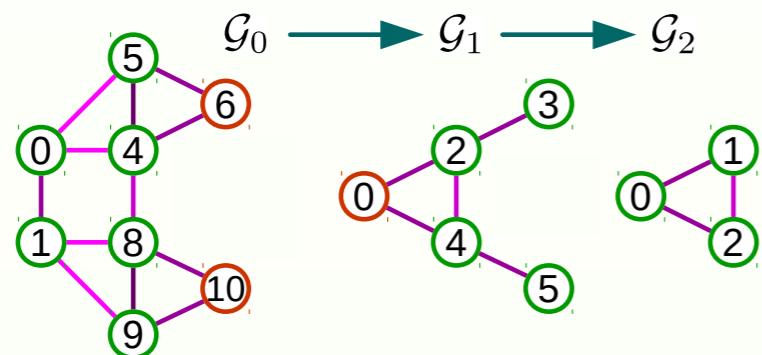
mean/max fail

Recent GNN Extensions

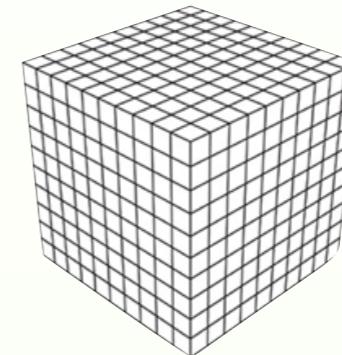
Hierarchical Models

Either via deterministic coarsening methods...

Graclus
(greedy pair matching)



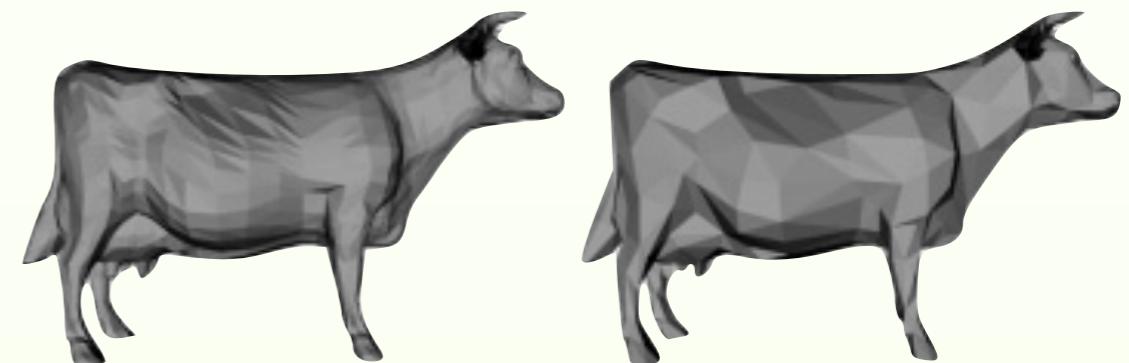
Voxel Grids
(pool_2d generalisation)



Iterative Sampling
of the most distant points



Mesh Downsampling
via quadric matrices



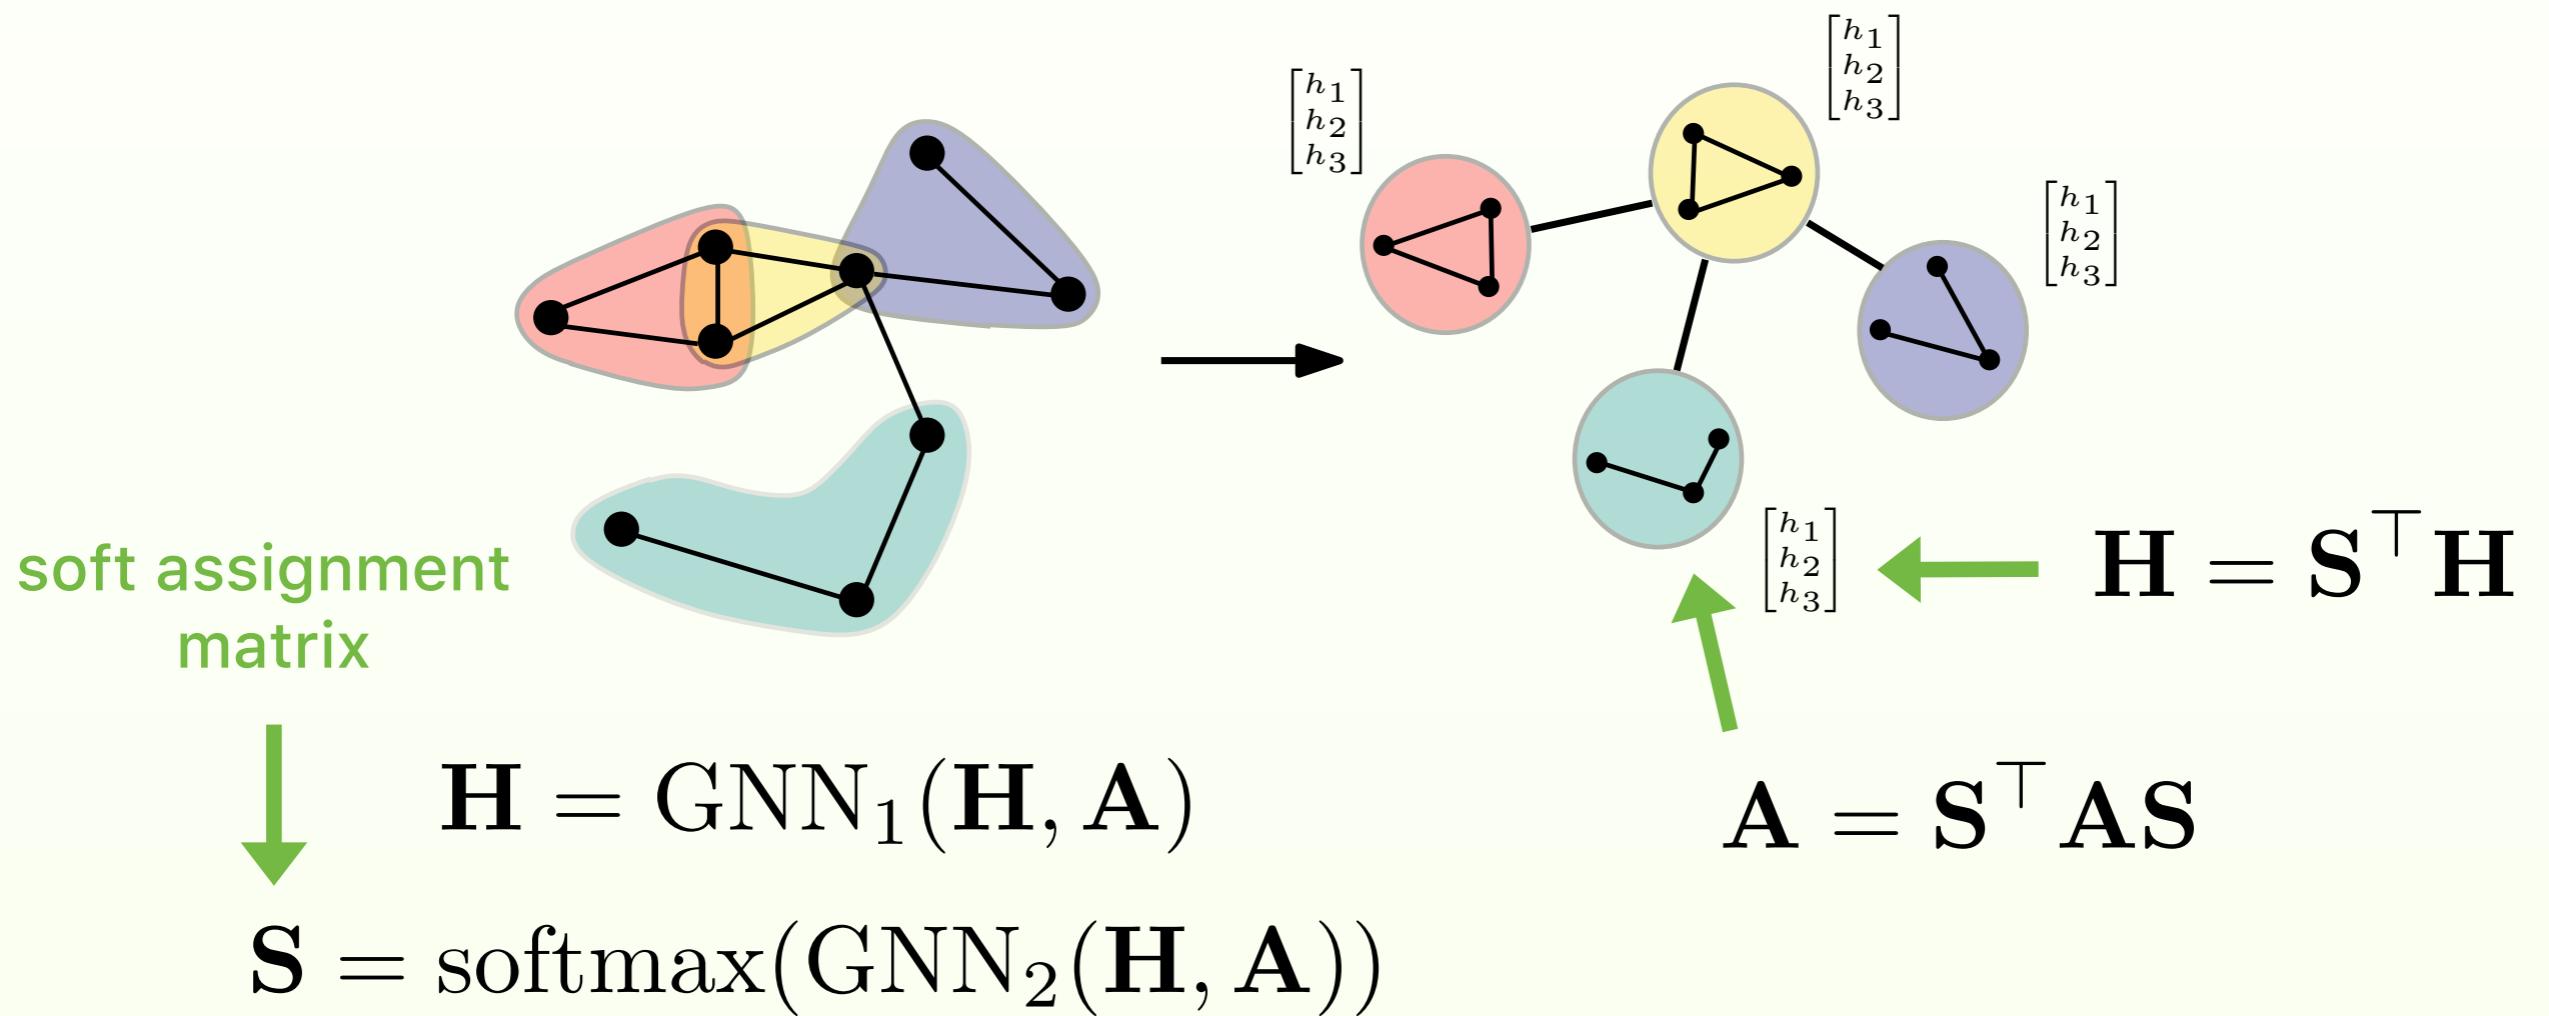
- Defferrard et al.: Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering (NIPS 2016)
Simonovsky and Komodakis: Dynamic Edge-Conditioned Filters in CNNs on Graphs (CVPR 2017)
Qi et al.: PointNet++ Deep Hierarchical Feature Learning on Point Sets in a Metric Space (NIPS 2017)
Ranjan et al.: Generating 3D Faces using Convolutional Mesh Autoencoders (ECCV 2018)

Hierarchical Models

... or differentiable pooling modules

DiffPool:

One GNN learns node embeddings
Another GNN learns cluster assignments

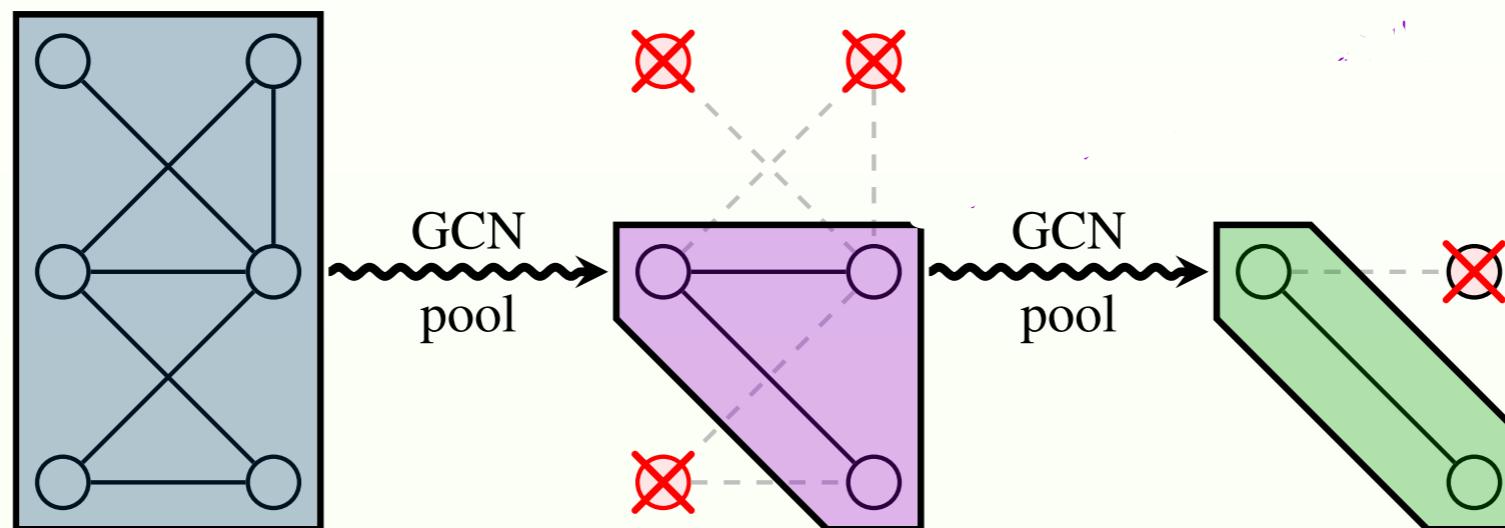


Hierarchical Models

... or differentiable pooling modules

Top-k Pooling:

Filter nodes based on scoring function



$$\vec{y} = \tanh(\mathbf{H} \cdot \vec{\theta})$$

$$\vec{i} = \text{top}_k(\vec{y})$$

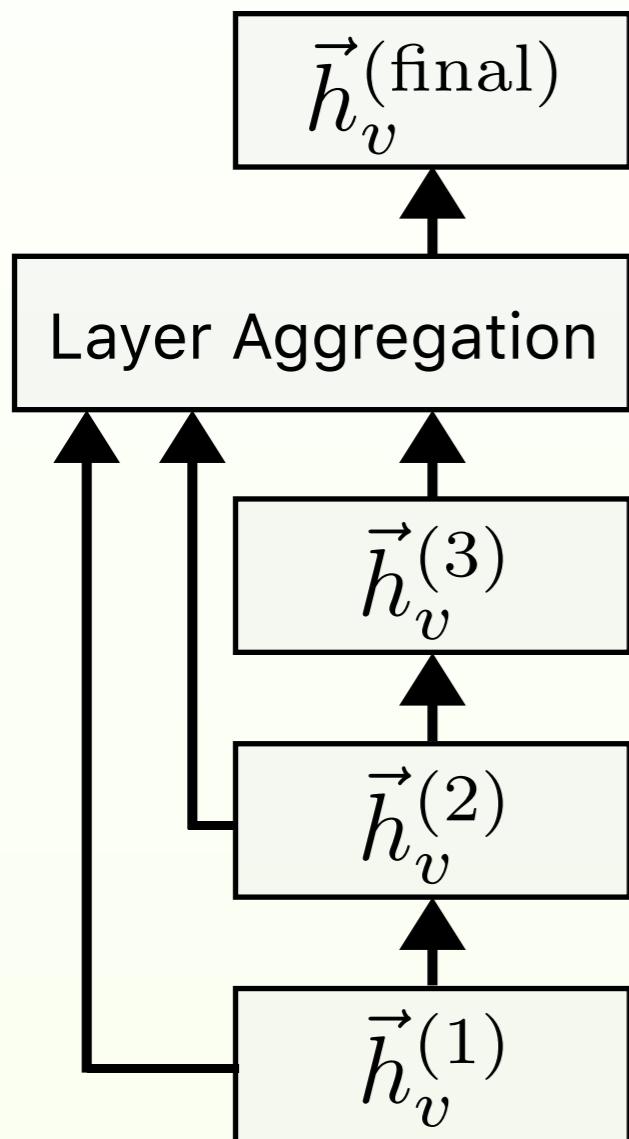
$$\mathbf{H} = (\mathbf{H} \odot \vec{y})_{\vec{i}}$$

ensure gradient computation

Deep(er) Models

"Washed out" representations after just a few layers

Idea: "Jump back" to earlier representations



Concatenation: $\vec{h}_v^{(1)} \parallel \dots \parallel \vec{h}_v^{(T)}$

Pooling:

$$\max\left(\vec{h}_v^{(1)}, \dots, \vec{h}_v^{(T)}\right)$$

Weighted Sum:

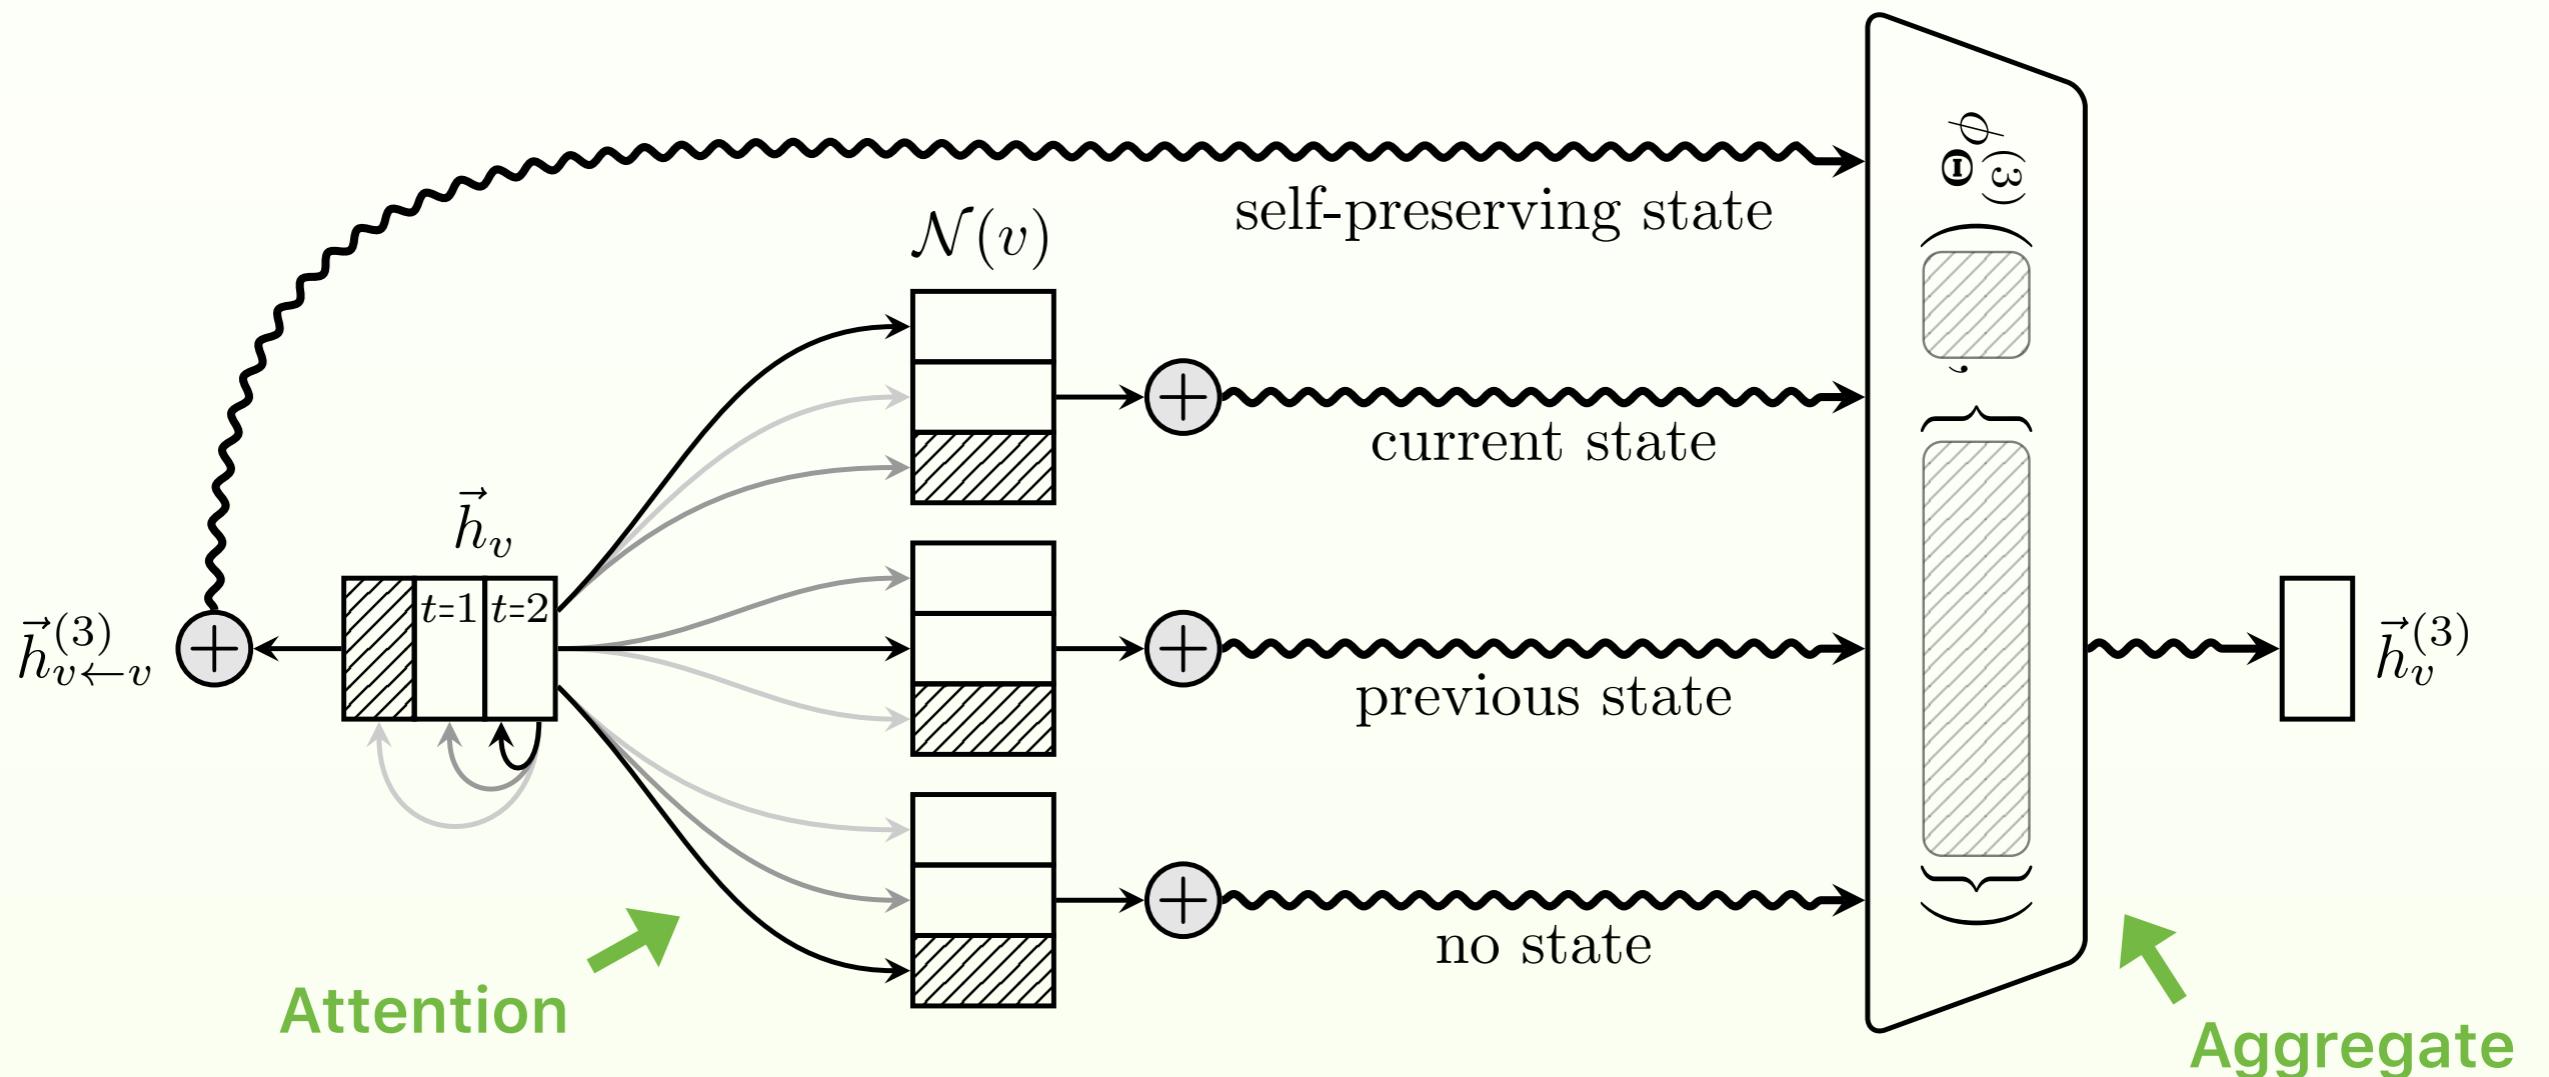
$$\sum_{t=1}^T \alpha_v^{(t)} \vec{h}_v^{(t)}$$

Attention ↙

Deep(er) Models

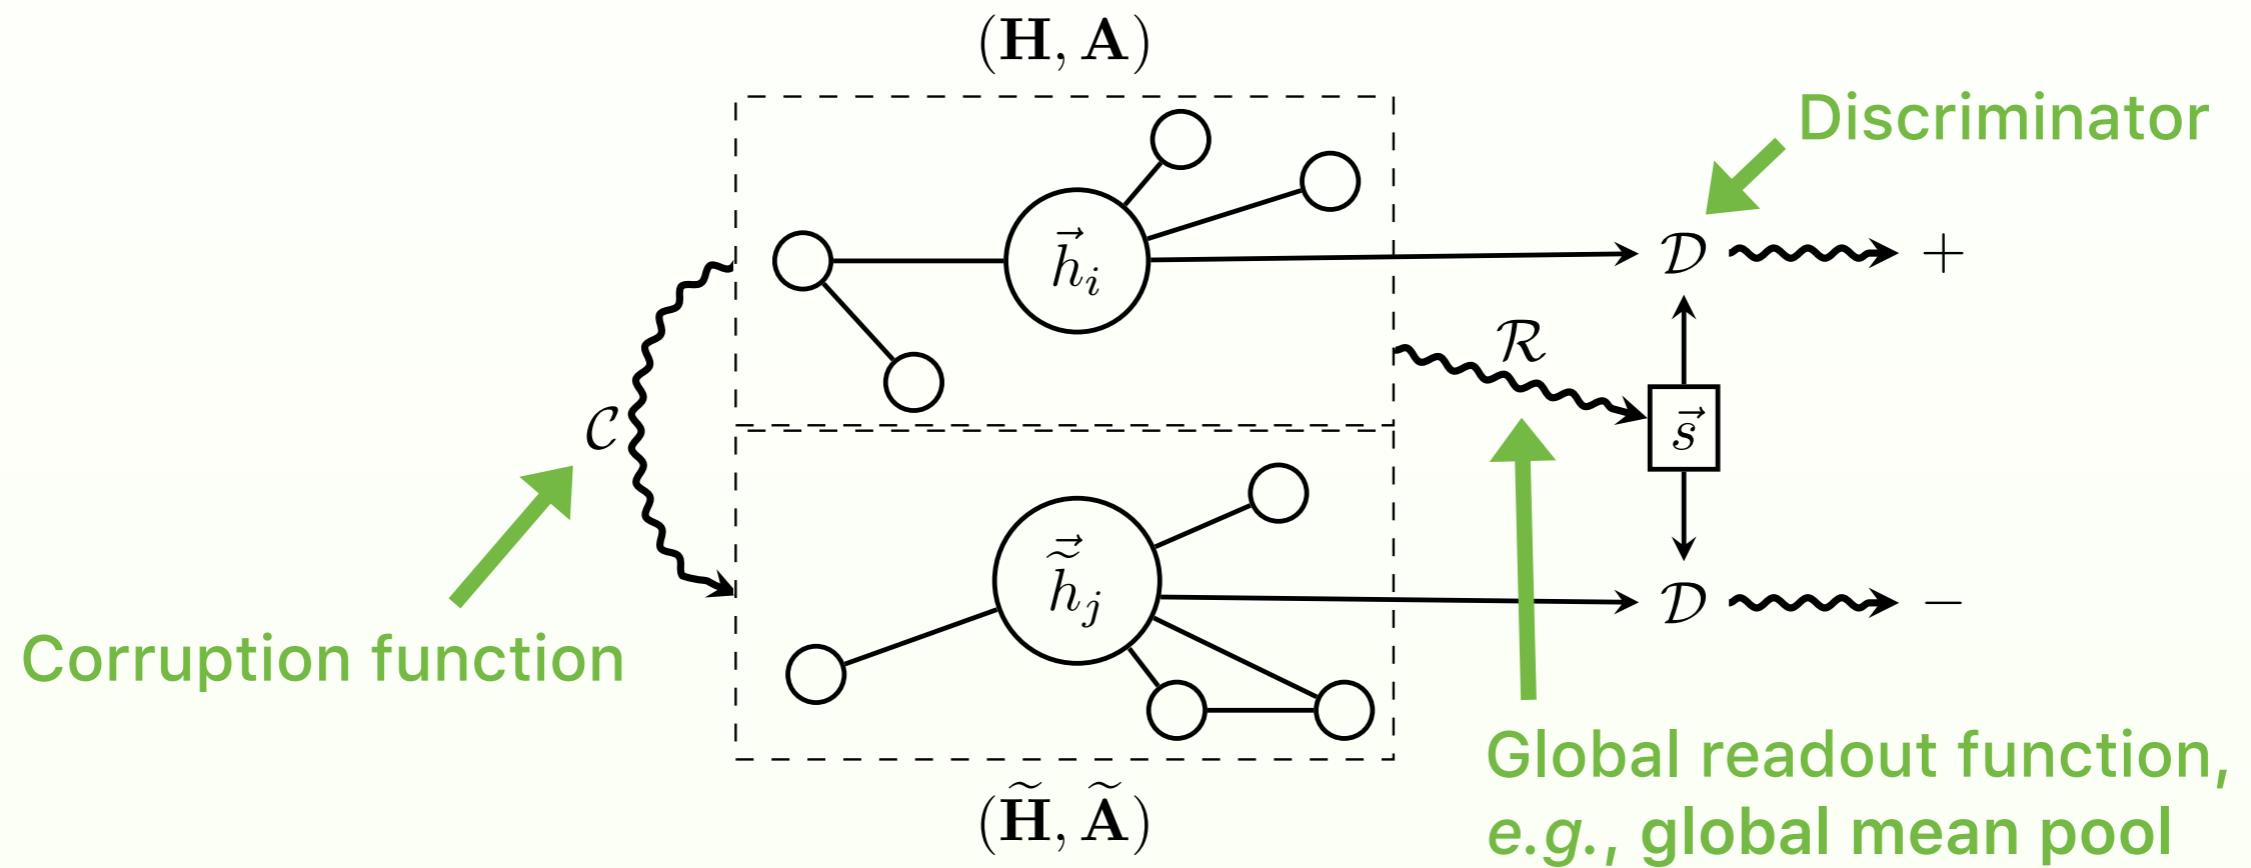
Extension: Allow jumps directly *while aggregating*

- Aggregate more global information in one branch while falling back to more local information in others



Unsupervised Models

based on Message Passing-based models



Maximize mutual information:

$$\sum_{v \in \mathcal{V}} \mathbb{E}_{(\mathbf{H}, \mathbf{A})} \left[\log \mathcal{D} \left(\vec{h}_v, \vec{s} \right) \right] + \mathbb{E}_{(\tilde{\mathbf{H}}, \tilde{\mathbf{A}})} \left[\log \left(1 - \mathcal{D} \left(\tilde{\vec{h}}_v, \vec{s} \right) \right) \right]$$

Generative Models

(Variational/Adversarially Regularized) Autoencoder:
GNN encoder and inner product/MLP decoder

→ expensive graph matching procedure, prefixed size

Generative Adversarial Nets for Graphs:
MLP generator and GNN discriminator

→ likelihood-free, but still prefixed in size

Auto-regressive Graph Generation:
Decomposing of graph generation into a sequence
of node and edge formations

→ not prefixed in size, but inherently sequential

Kipf and Welling: Variational Graph Auto-Encoders (NIPS-W 2016)

Pan et al.: Adversarially Regularized Graph Autoencoder for Graph Embedding (IJCAI 2018)

Simonovsky and Komodakis: GraphVAE: Towards Generation of Small Graphs Using VAEs (arXiv/1802.03480)

Cao and Kipf: MolGAN: An Implicit Generative Model for Small Molecular Graphs (ICML-W 2018)

You et al.: Generating Realistic Graphs with Deep Auto-regressive Models (ICML 2018)

Conclusion

GNNs have interesting applications in many different areas

GNNs follow a simple message passing scheme to learn *localized* embeddings

GNNs can be very efficiently implemented

GNNs can be enhanced by and used for various concepts

GNNs are really hot right now 🔥🔥🔥

