



# Tutorial on Graph Neural Networks

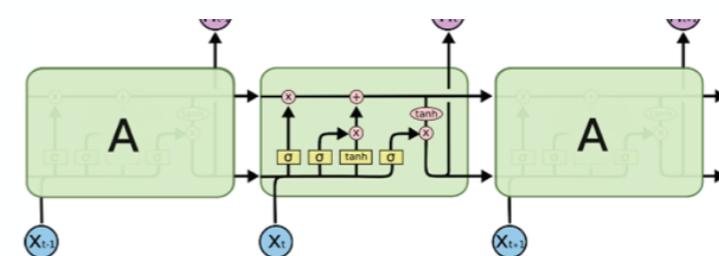
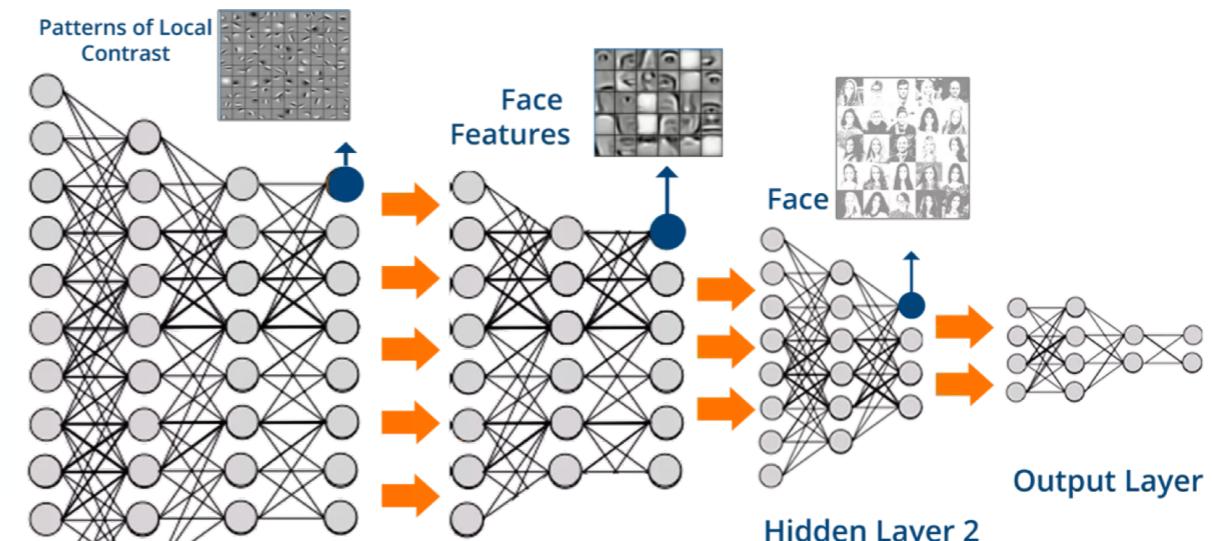
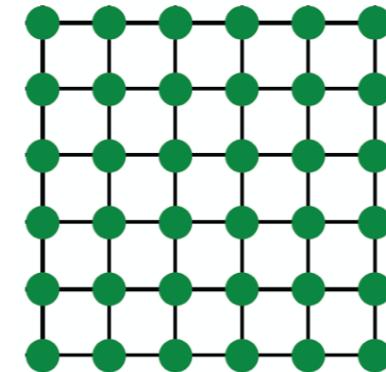
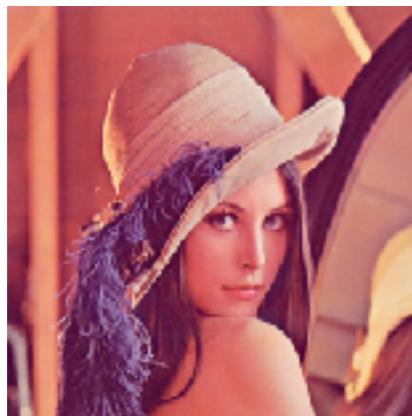
Workshop on Geometry and Machine Learning with Applications to  
Biomedical Engineering - University College London



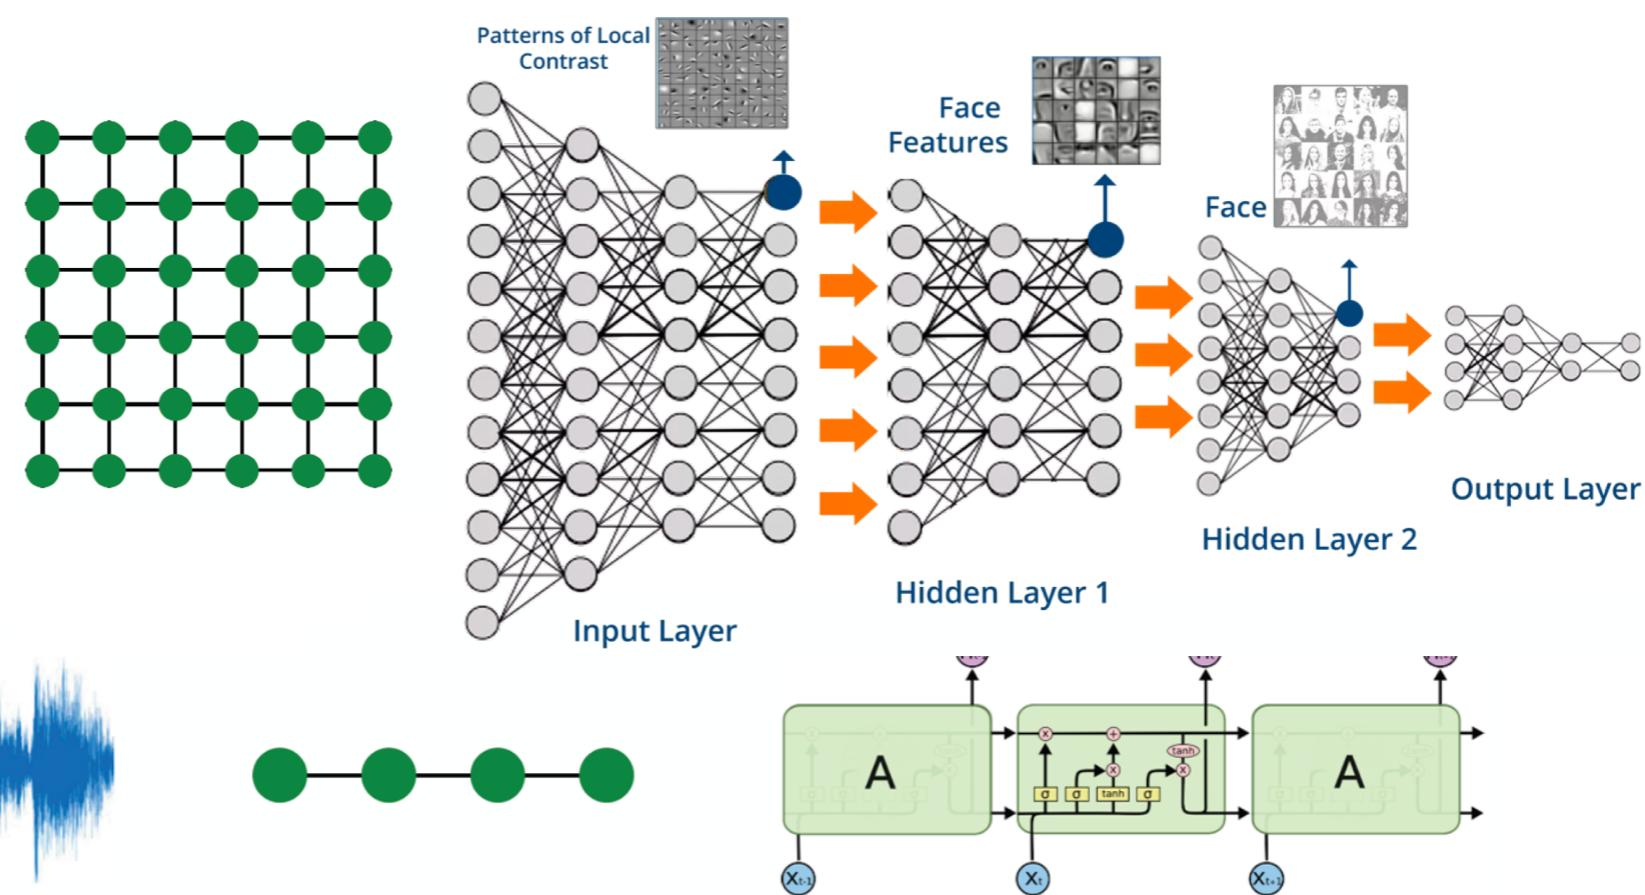
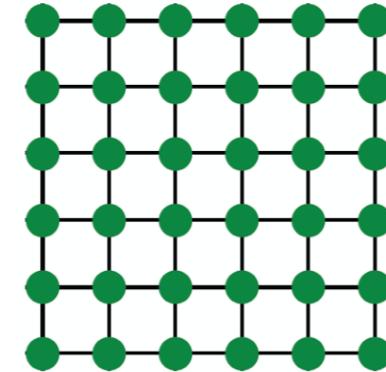
Matthias Fey

- ✉ matthias.fey@udo.edu
- ☁ rusty1s.github.io
- 🐱 🐦 /rusty1s

The modern deep learning toolbox is designed for simple sequences and grids



The modern deep learning toolbox is designed for simple sequences and grids

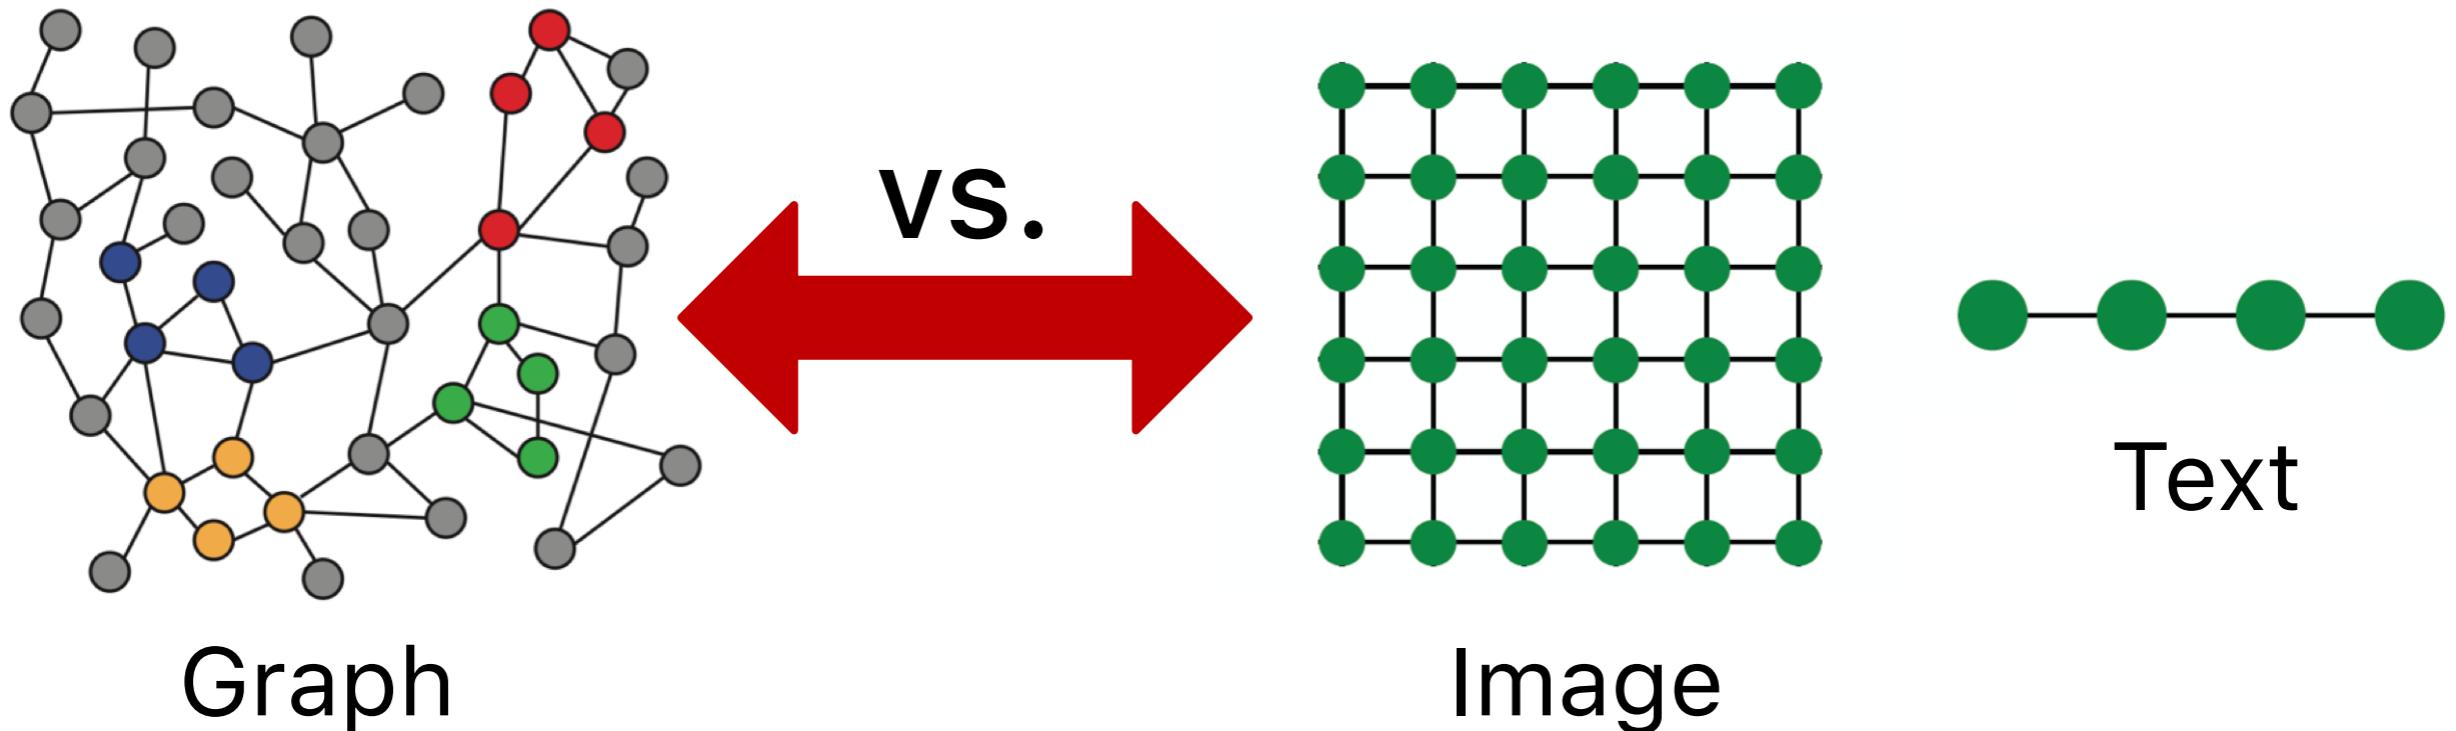


**Not everything can be represented as such!**

How can we develop neural networks that are much more broadly applicable?

(Mostly) everything can be represented as a graph

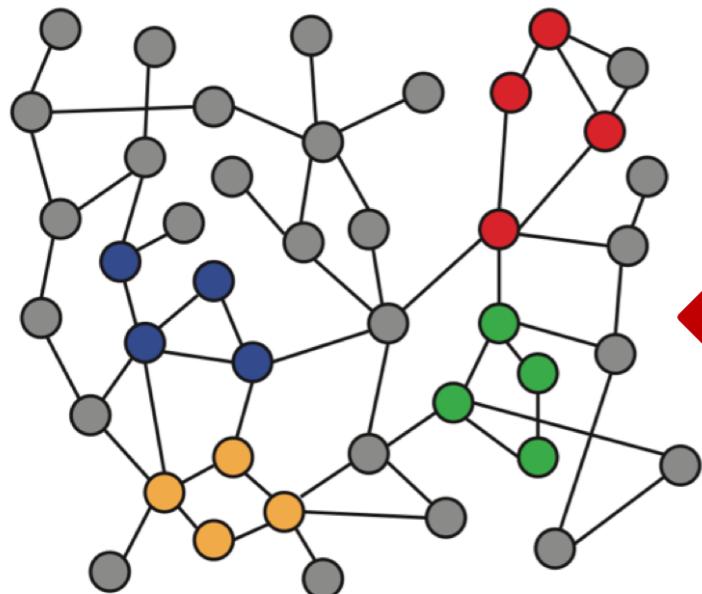
- ▶ Going beyond classic neural networks via Graph Neural Networks



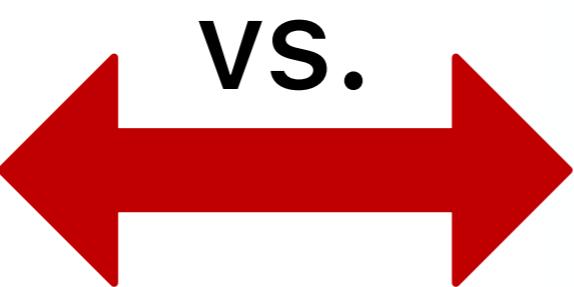
How can we enable learning on arbitrary graphs?

## Networks are complex!

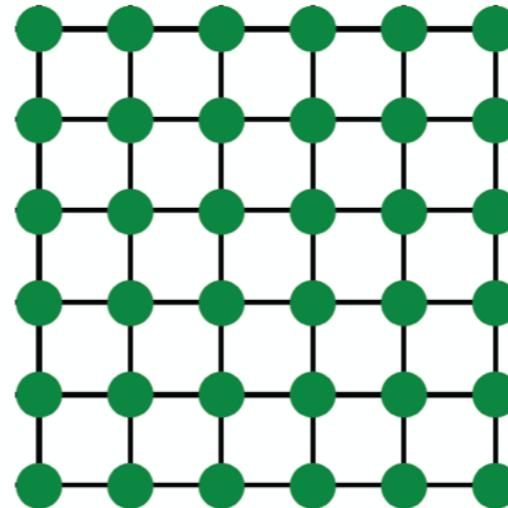
- ▶ Arbitrary size and complex topological structure
- ▶ No fixed node ordering or reference point
- ▶ Often dynamic
- ▶ Multimodal node and edge features



Graph



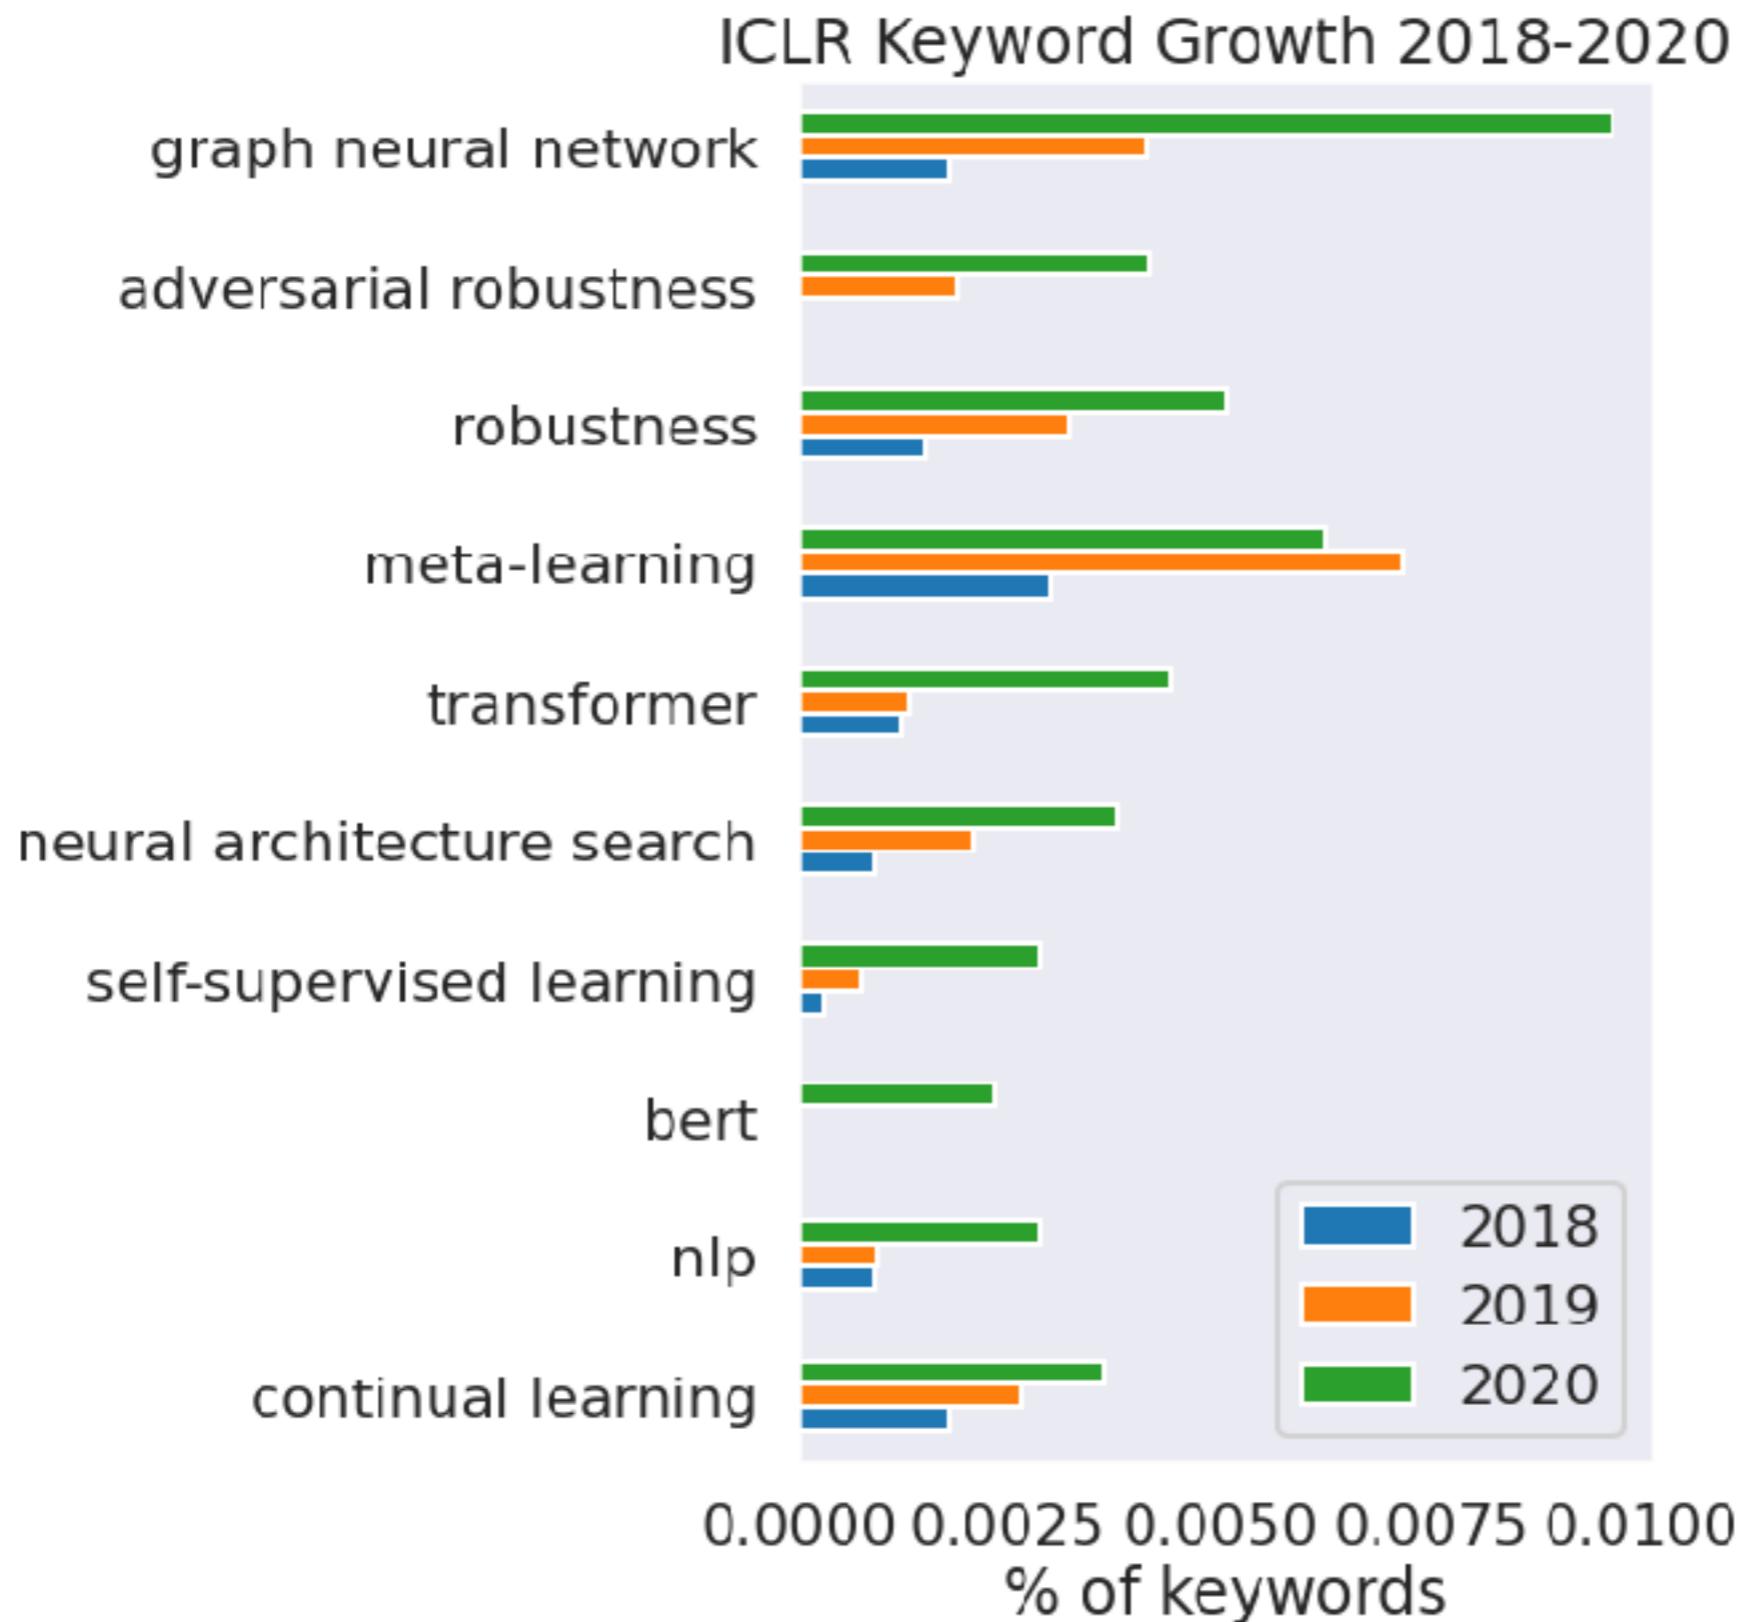
vs.



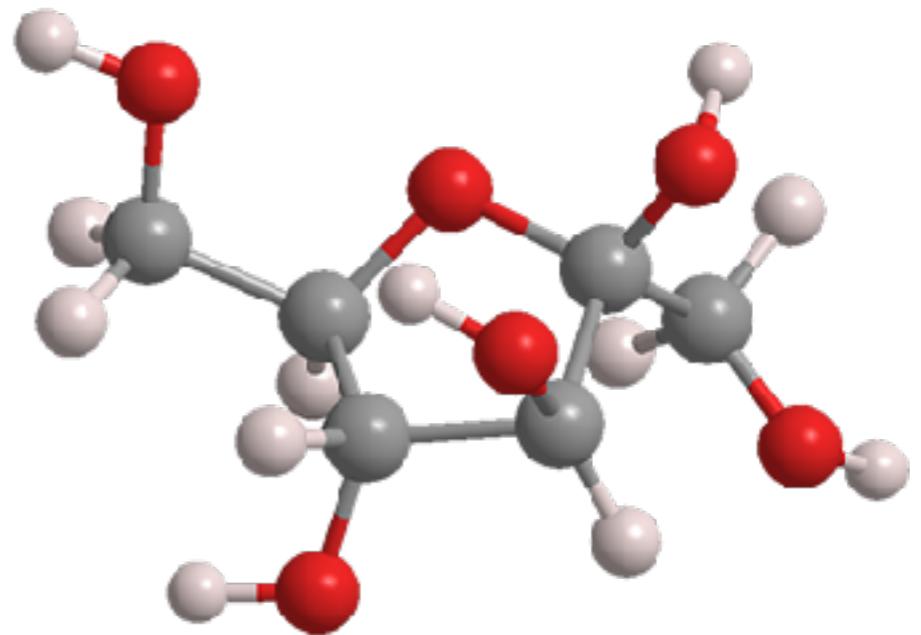
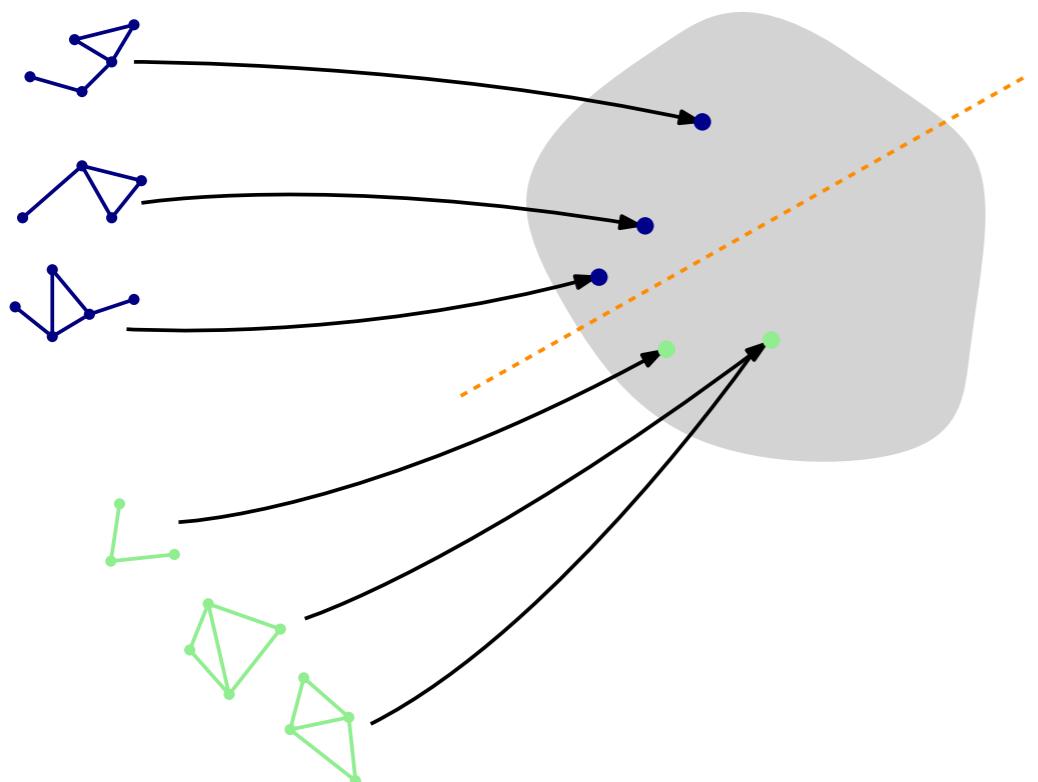
Image



Text

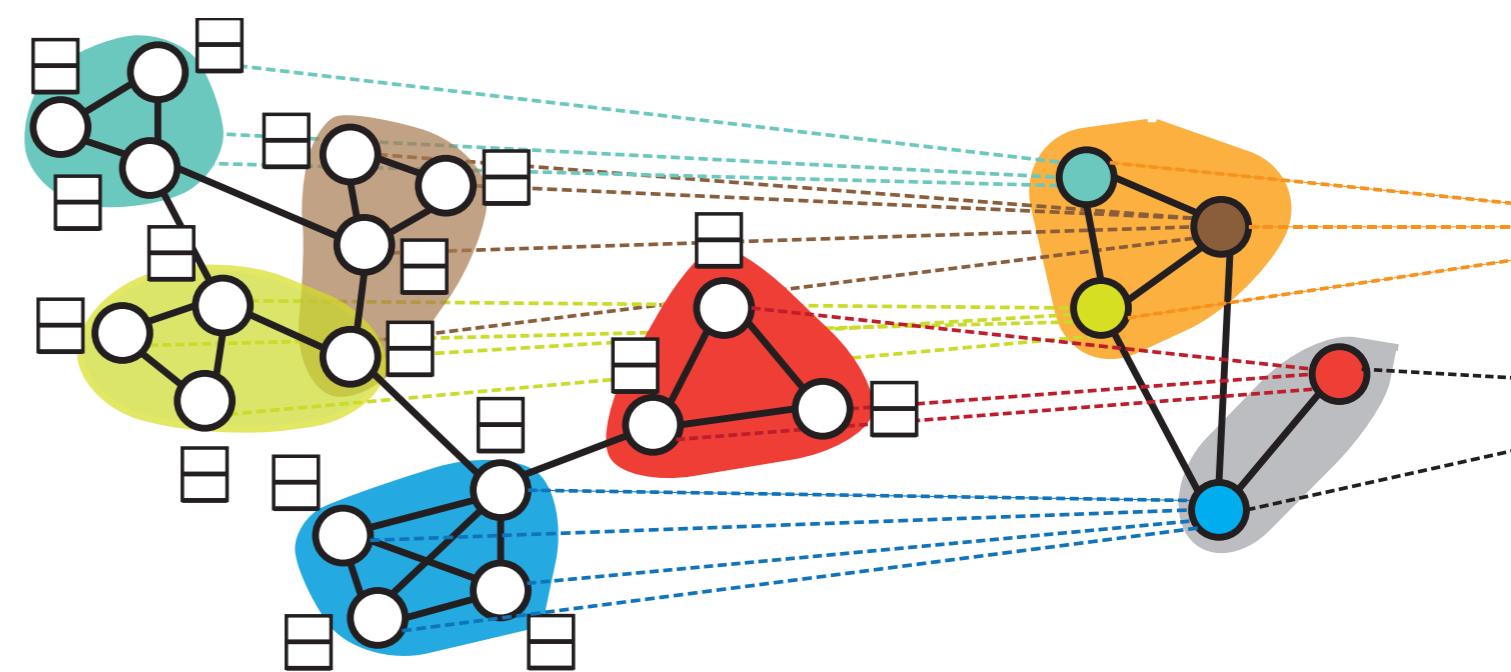
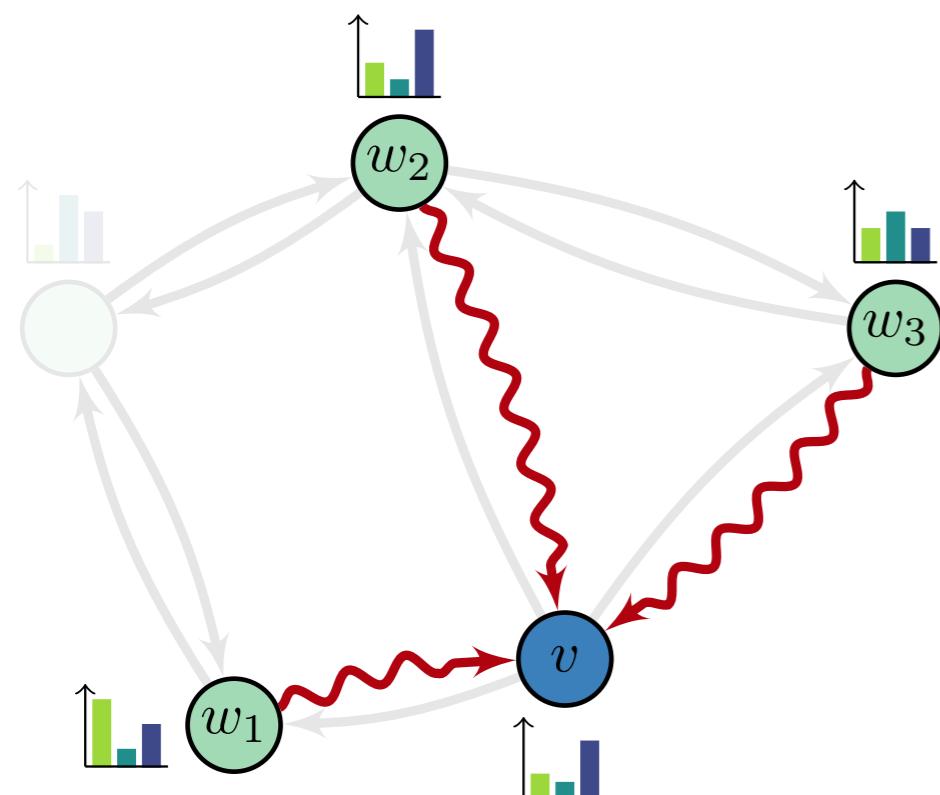


# 1. Learning on Graphs and its Applications



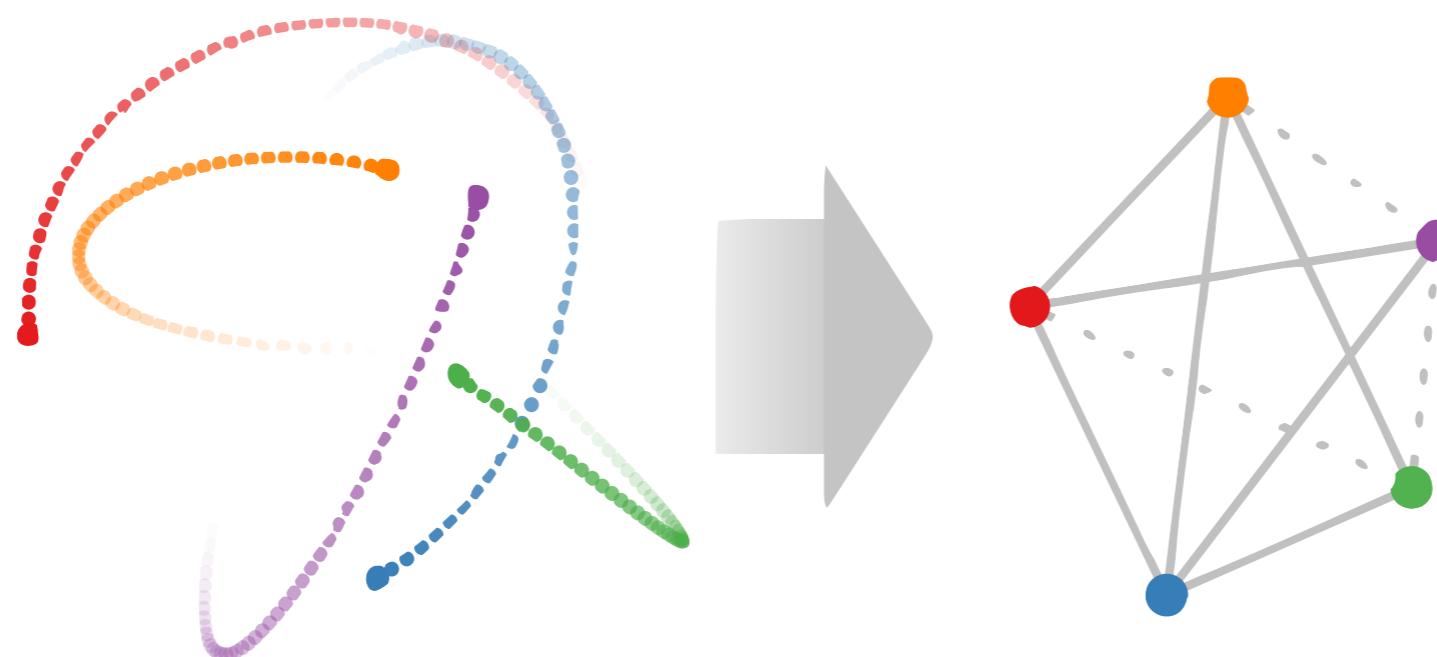


1. Learning on Graphs and its Applications
2. What are Graph Neural Networks (GNNs)?



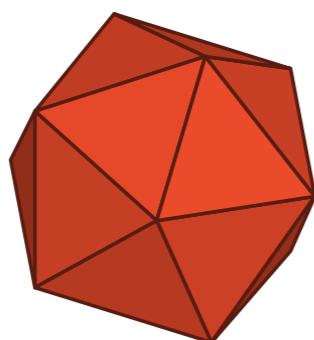


1. Learning on Graphs and its Applications
2. What are Graph Neural Networks (GNNs)?
3. Relational Structure Discovery
  - ▶ *How can we discover objects and their relations?*





1. Learning on Graphs and its Applications
2. What are Graph Neural Networks (GNNs)?
3. Relational Structure Discovery
  - ▶ *How can we discover objects and their relations?*
4. Implementing Graph Neural Networks



PyTorch  
geometric

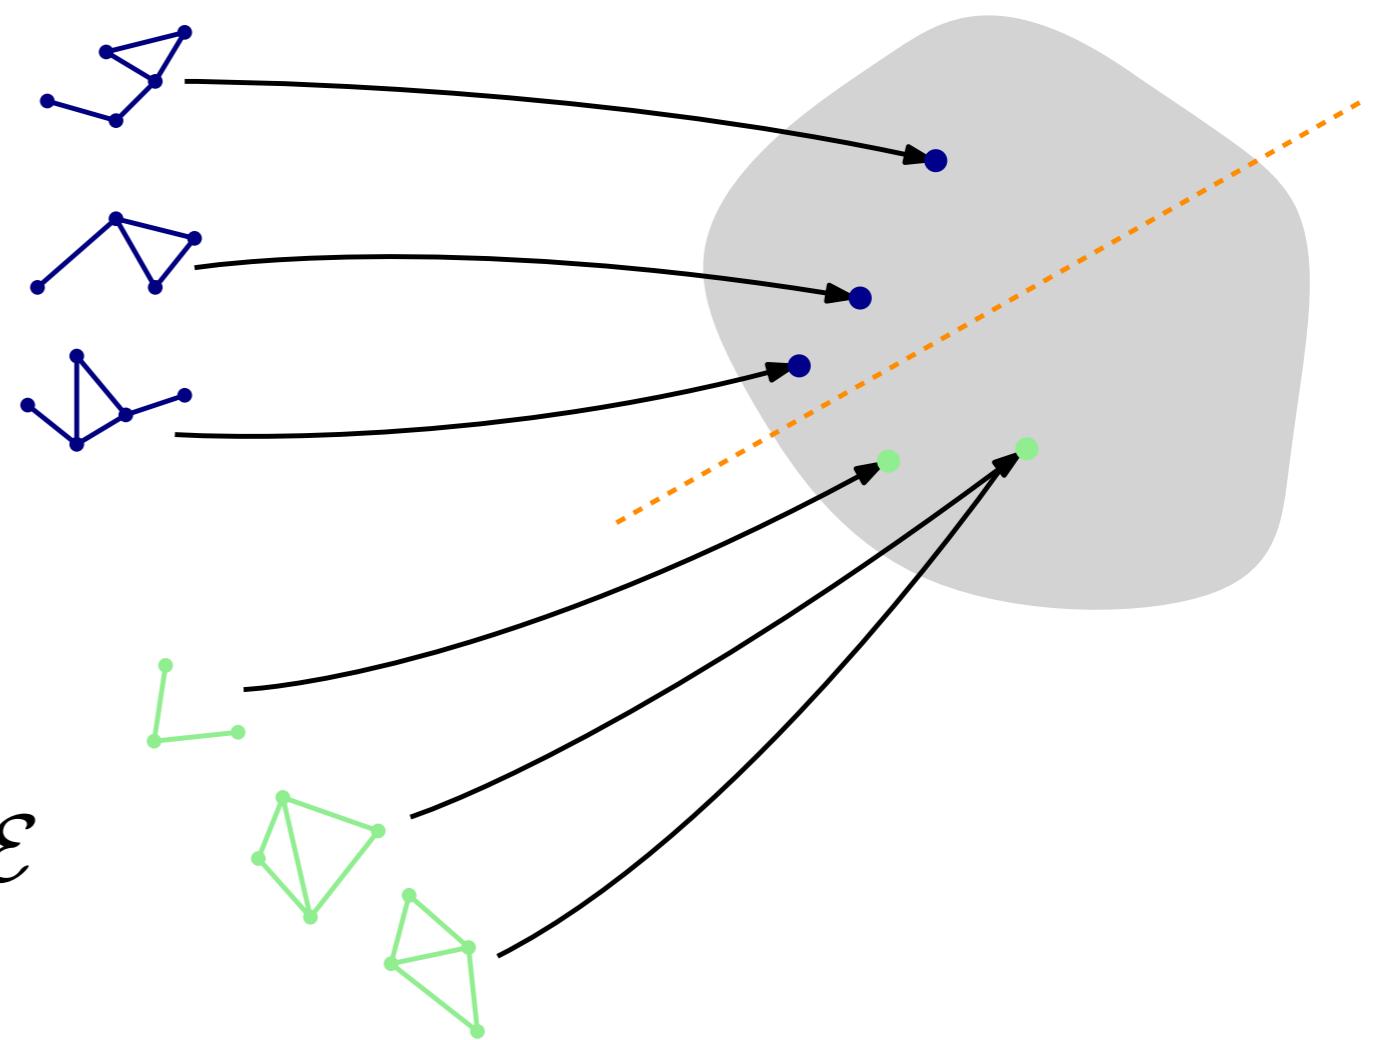


**Aim:** Map nodes/graphs into an embedding space, in which they can be distinguished for a given task

**Graph:**  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$   
or  $\mathbf{A} \in \{0, 1\}^{|\mathcal{V}| \times |\mathcal{V}|}$

**Node features:**  
 $\mathbf{x}_v \in \mathbb{R}^F$  for all  $v \in \mathcal{V}$

**Edge features:**  
 $\mathbf{e}_{v,w} \in \mathbb{R}^d$  for all  $(v, w) \in \mathcal{E}$



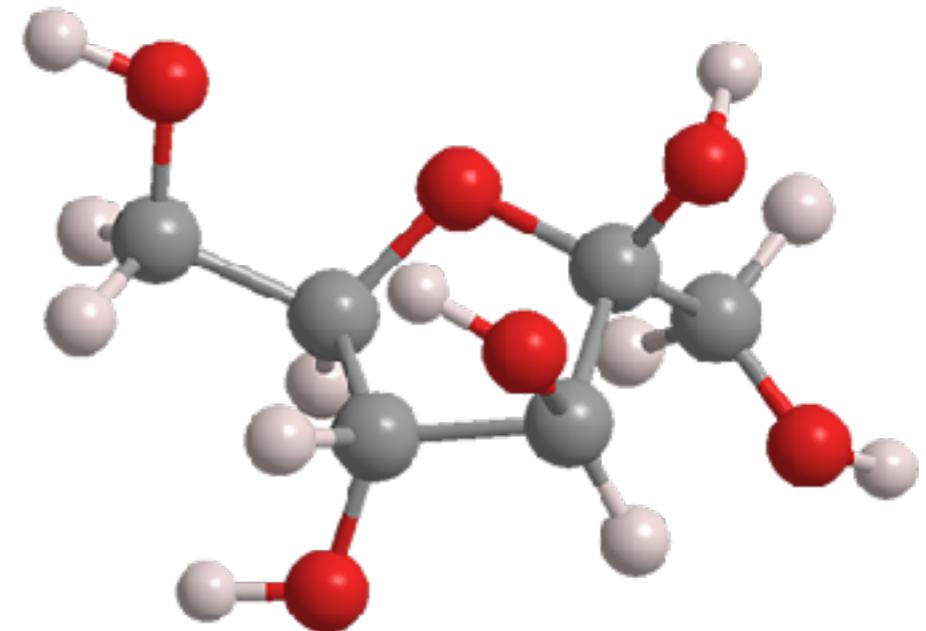
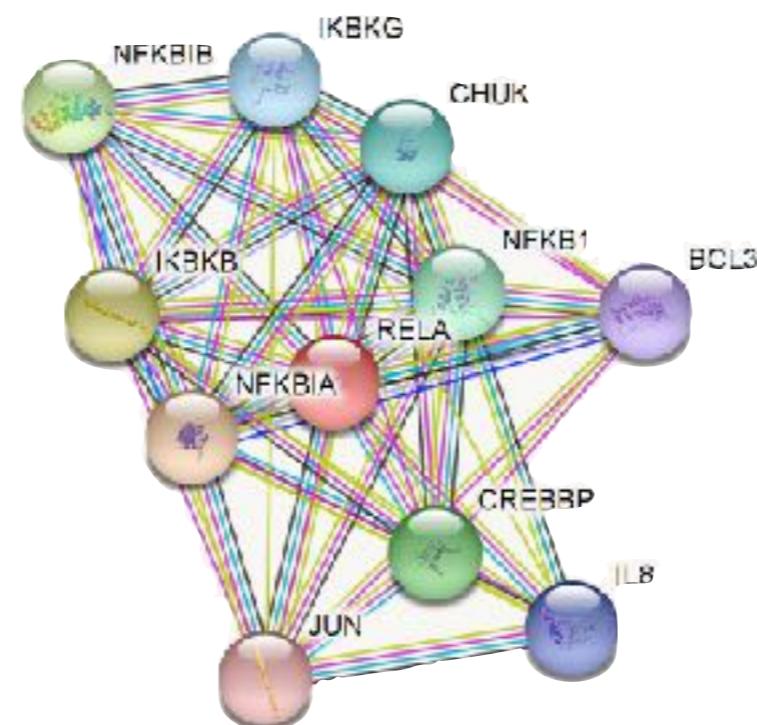


Social Networks



Social Networks

### Protein-Protein- Interactions

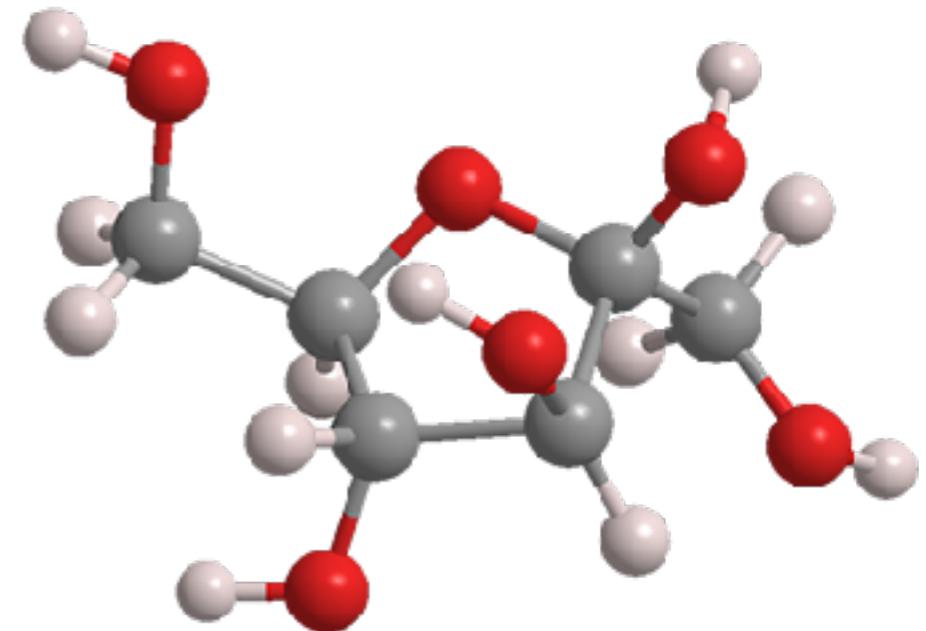
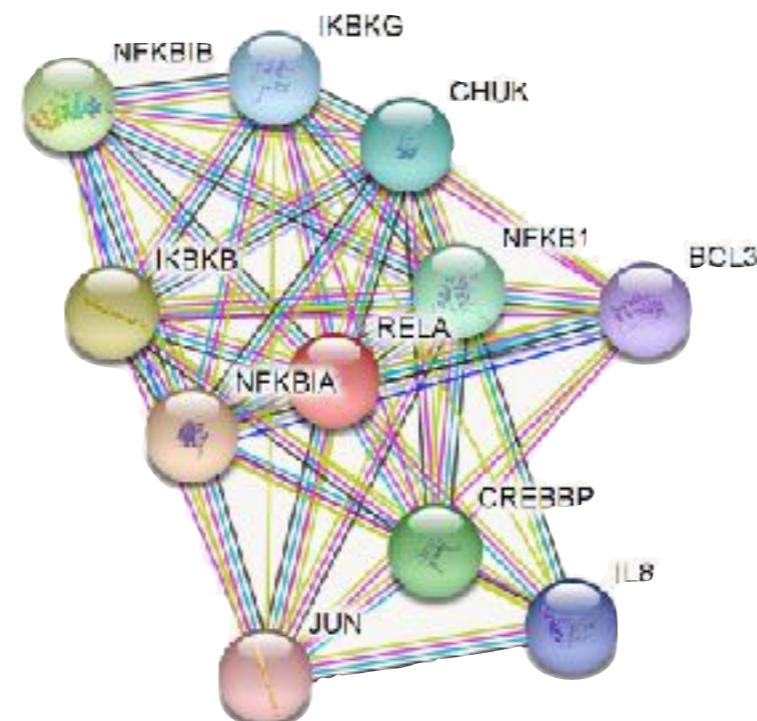


Predicting molecular properties

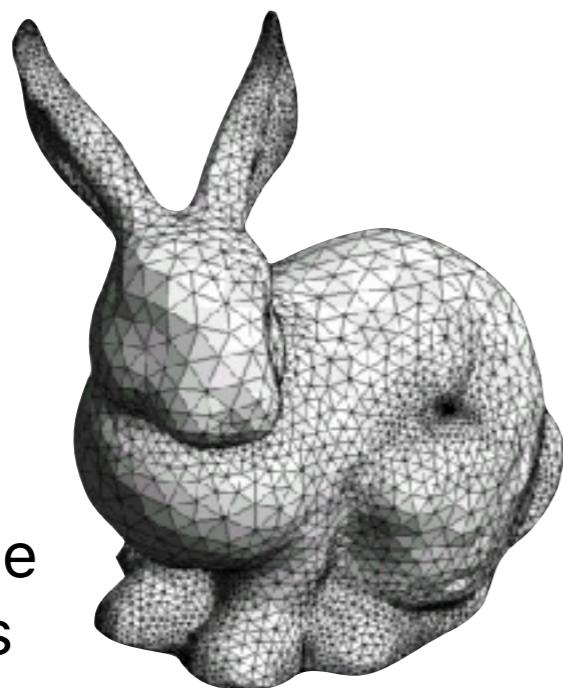


Social Networks

### Protein-Protein-Interactions



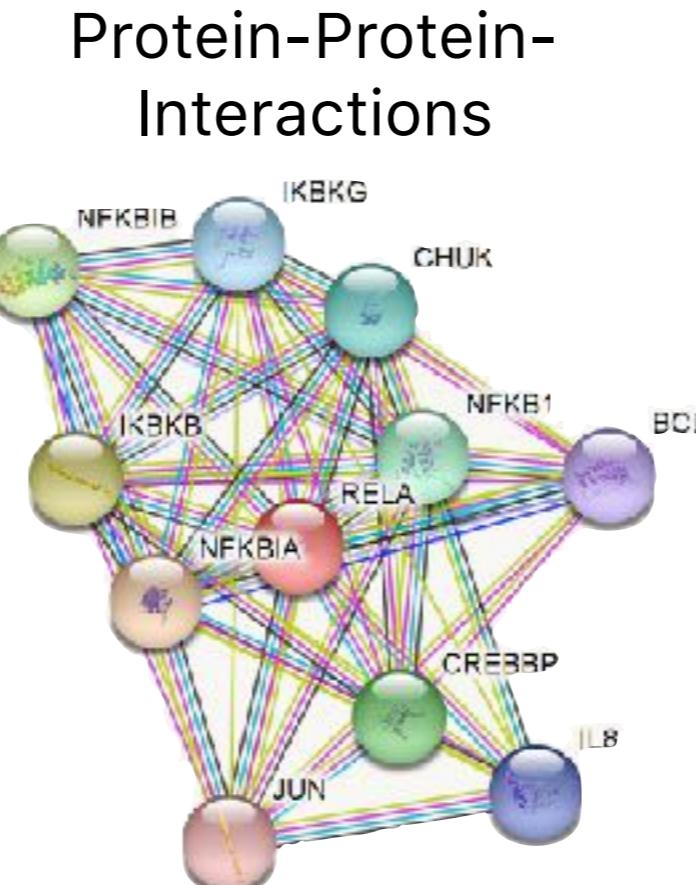
Predicting molecular properties



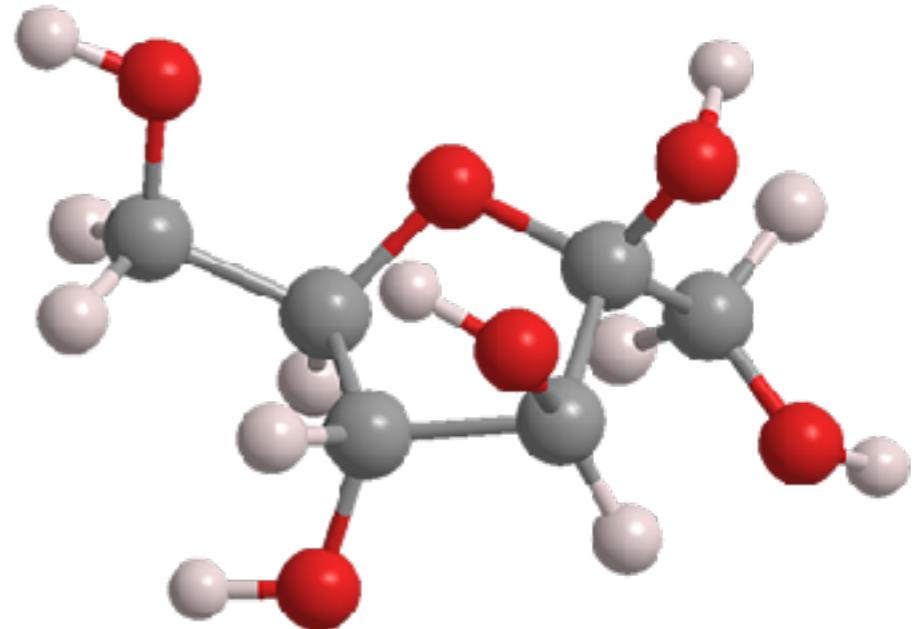
3D Shape  
Analysis



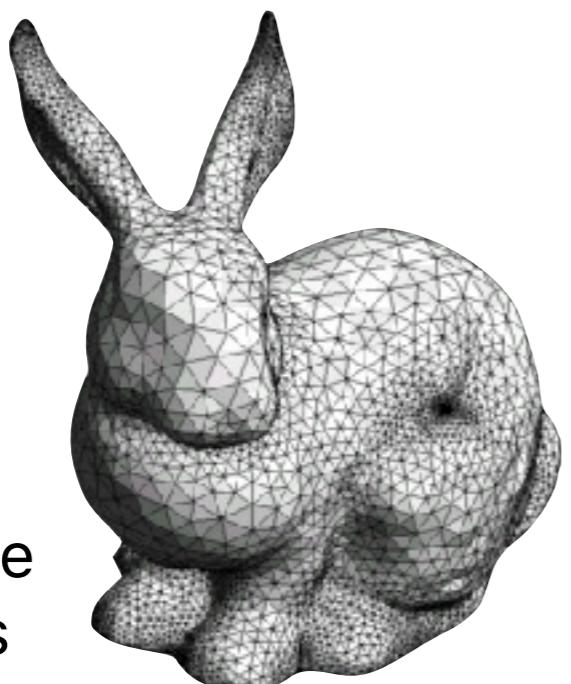
Social Networks



Protein-Protein-  
Interactions



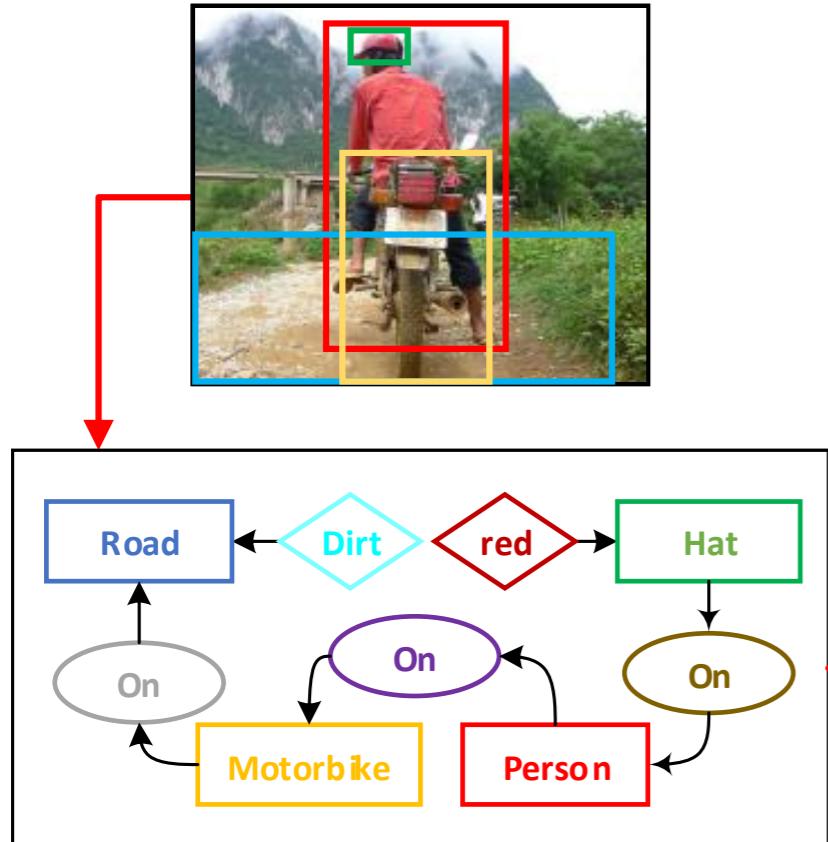
Predicting molecular properties



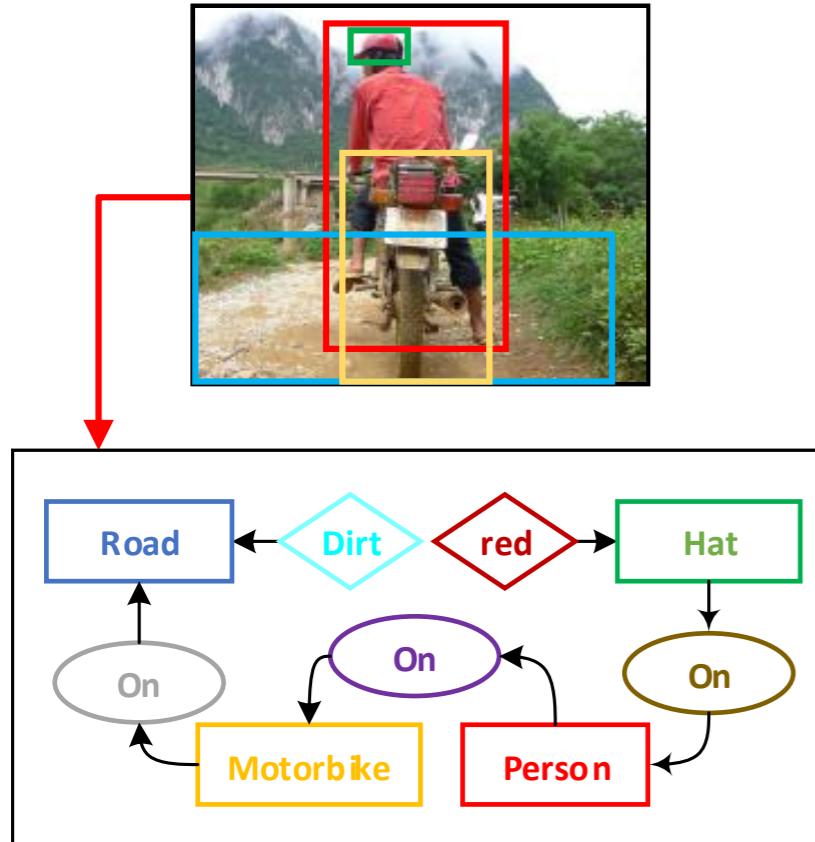
3D Shape  
Analysis



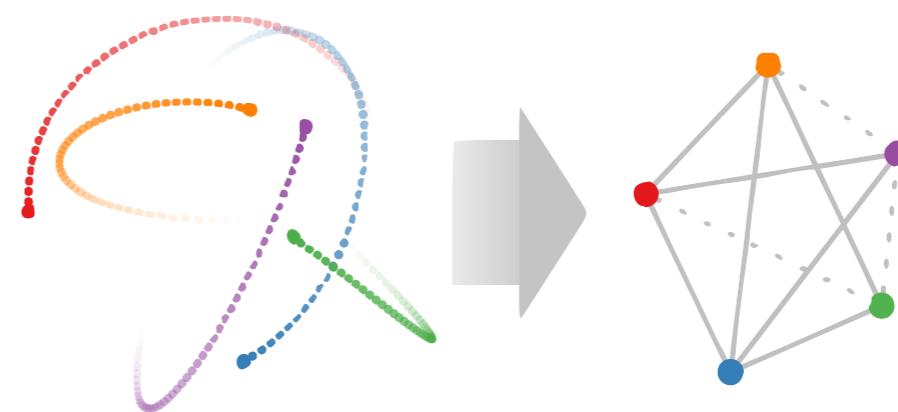
Human Pose Estimation



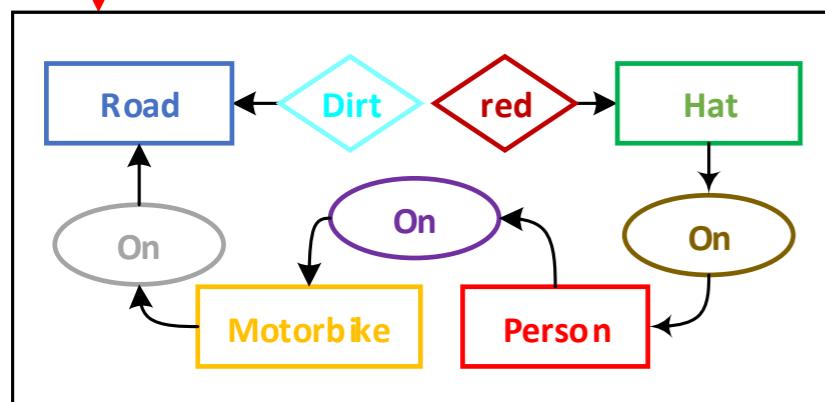
Understanding Scenes



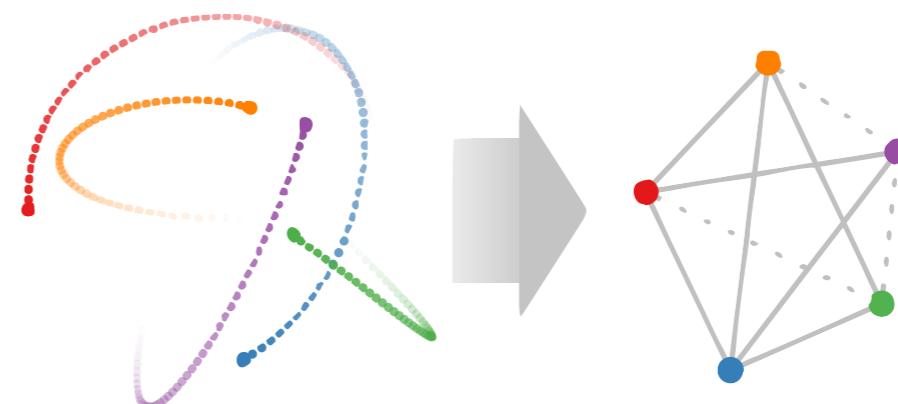
Understanding Scenes



Understanding the dynamics in  
physical systems

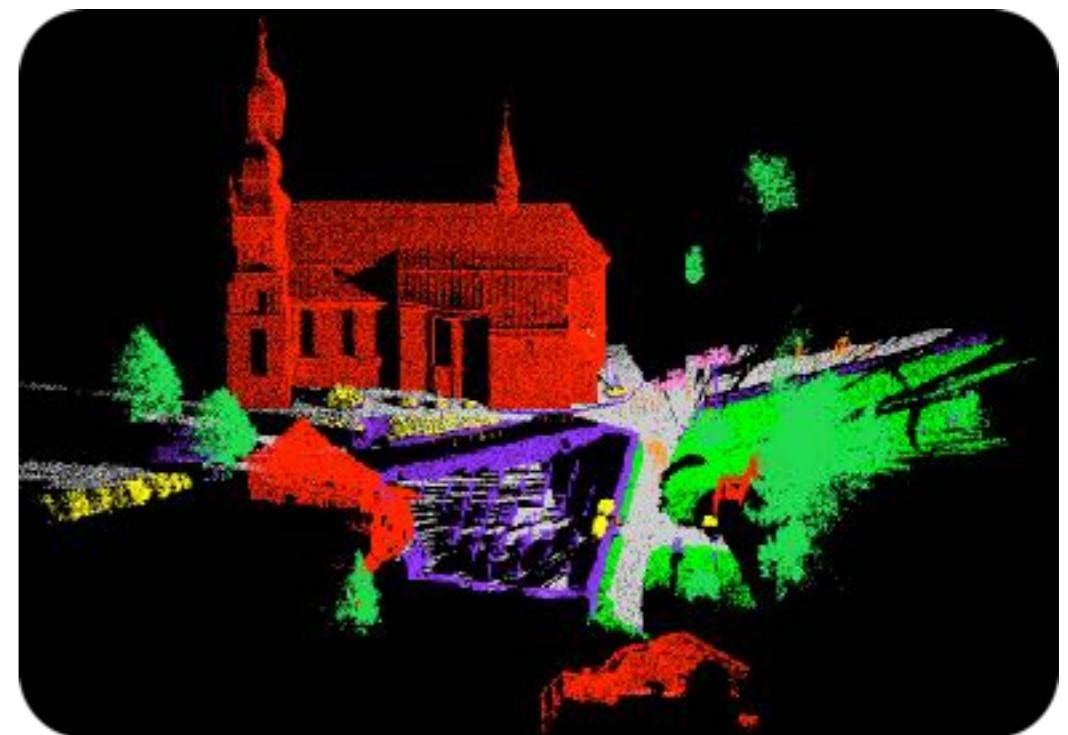


Understanding Scenes



Understanding the dynamics in physical systems

Autonomous driving



Segmenting Point Clouds



Graph Neural Networks borrow ideas from classical neural networks and generalize them to graphs:

### Locality

- ▶ Learn local patterns of the graph

### Weight sharing

- ▶ Learn universally applicable patterns of the graph



Graph Neural Networks borrow ideas from classical neural networks and generalize them to graphs:

### Locality

- ▶ Learn local patterns of the graph

### Weight sharing

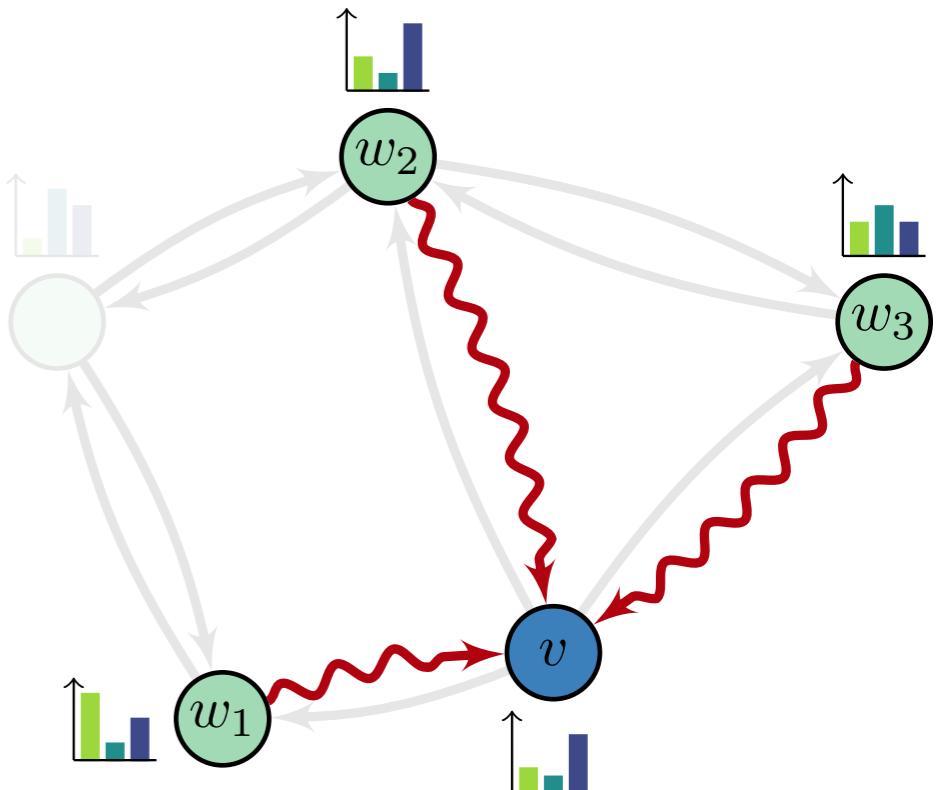
- ▶ Learn universally applicable patterns of the graph

A Graph Neural Network updates node representations by repeatedly transforming and aggregating neighboring node representations

# Message Passing Scheme

1. Each neighbor sends a message:

$$\mathbf{m}_{w,v}^{(\ell)} = \text{MESSAGE}(\mathbf{h}_v^{(\ell-1)}, \mathbf{h}_w^{(\ell-1)})$$



## Message Passing Scheme

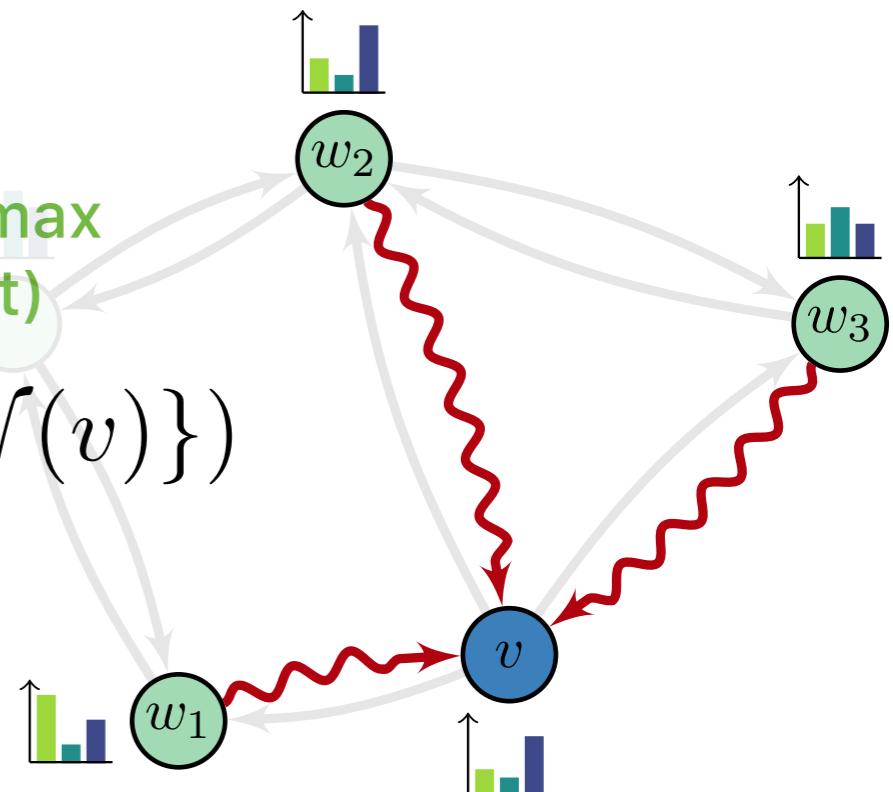
1. Each neighbor sends a message:

$$\mathbf{m}_{w,v}^{(\ell)} = \text{MESSAGE}(\mathbf{h}_v^{(\ell-1)}, \mathbf{h}_w^{(\ell-1)})$$

2. Messages are aggregated across all neighbors:

typically sum, mean or max  
(permutation-invariant)

$$\mathbf{a}_v^{(\ell)} = \text{AGGREGATE}(\{\mathbf{m}_{w,v}^{(\ell)} : w \in \mathcal{N}(v)\})$$



## Message Passing Scheme

1. Each neighbor sends a message:

$$\mathbf{m}_{w,v}^{(\ell)} = \text{MESSAGE}(\mathbf{h}_v^{(\ell-1)}, \mathbf{h}_w^{(\ell-1)})$$

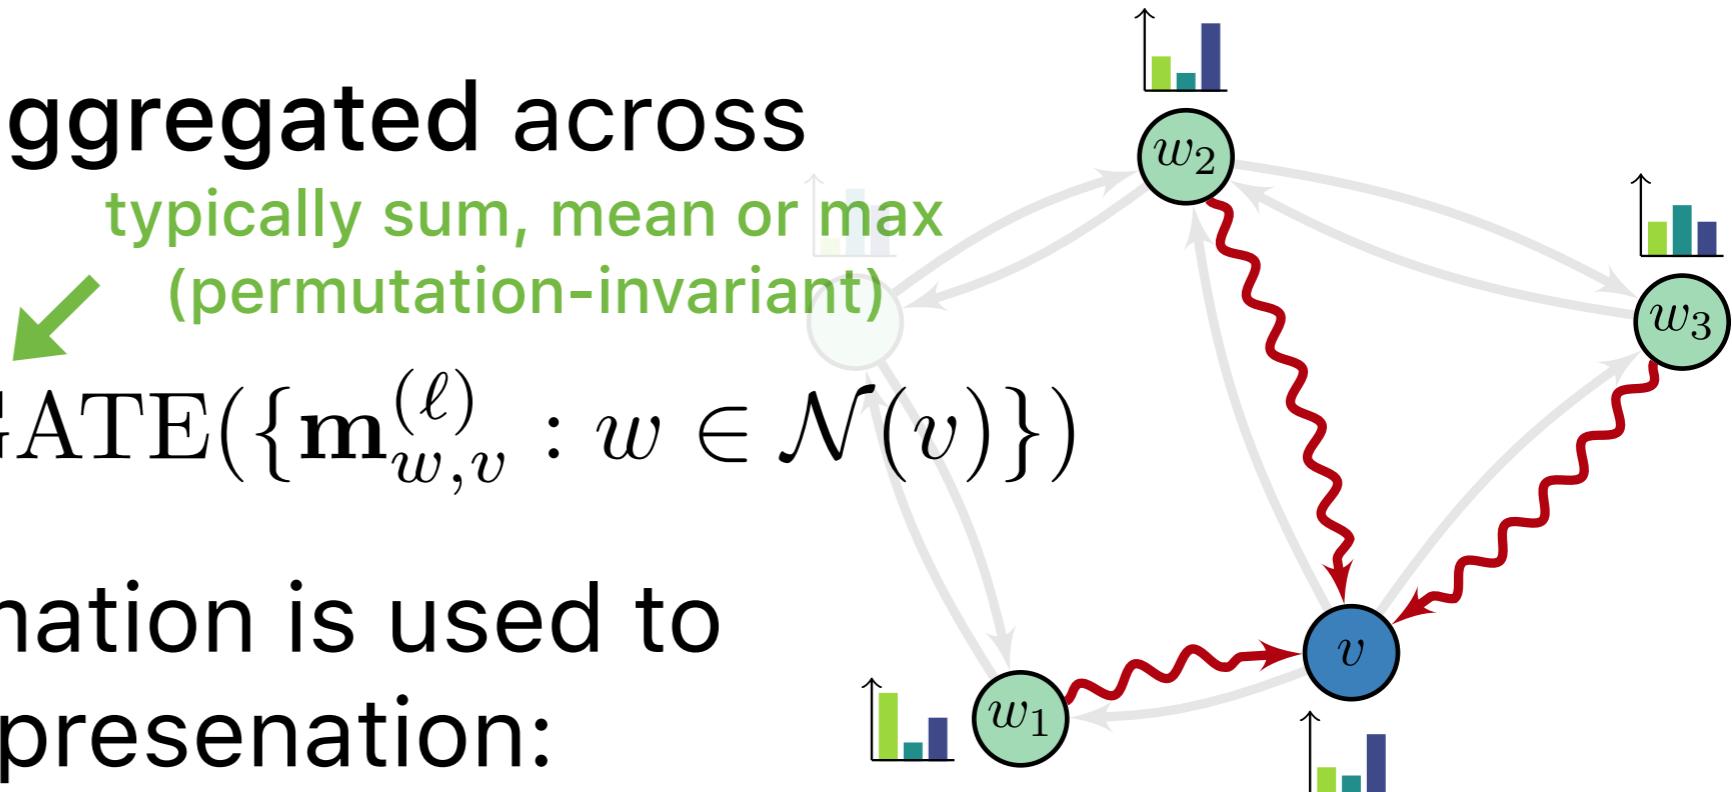
2. Messages are aggregated across all neighbors:

typically sum, mean or max  
(permutation-invariant)

$$\mathbf{a}_v^{(\ell)} = \text{AGGREGATE}(\{\mathbf{m}_{w,v}^{(\ell)} : w \in \mathcal{N}(v)\})$$

3. Neighbor information is used to update node representation:

$$\mathbf{h}_v^{(\ell)} = \text{UPDATE}(\mathbf{h}_v^{(\ell-1)}, \mathbf{a}_v^{(\ell)})$$

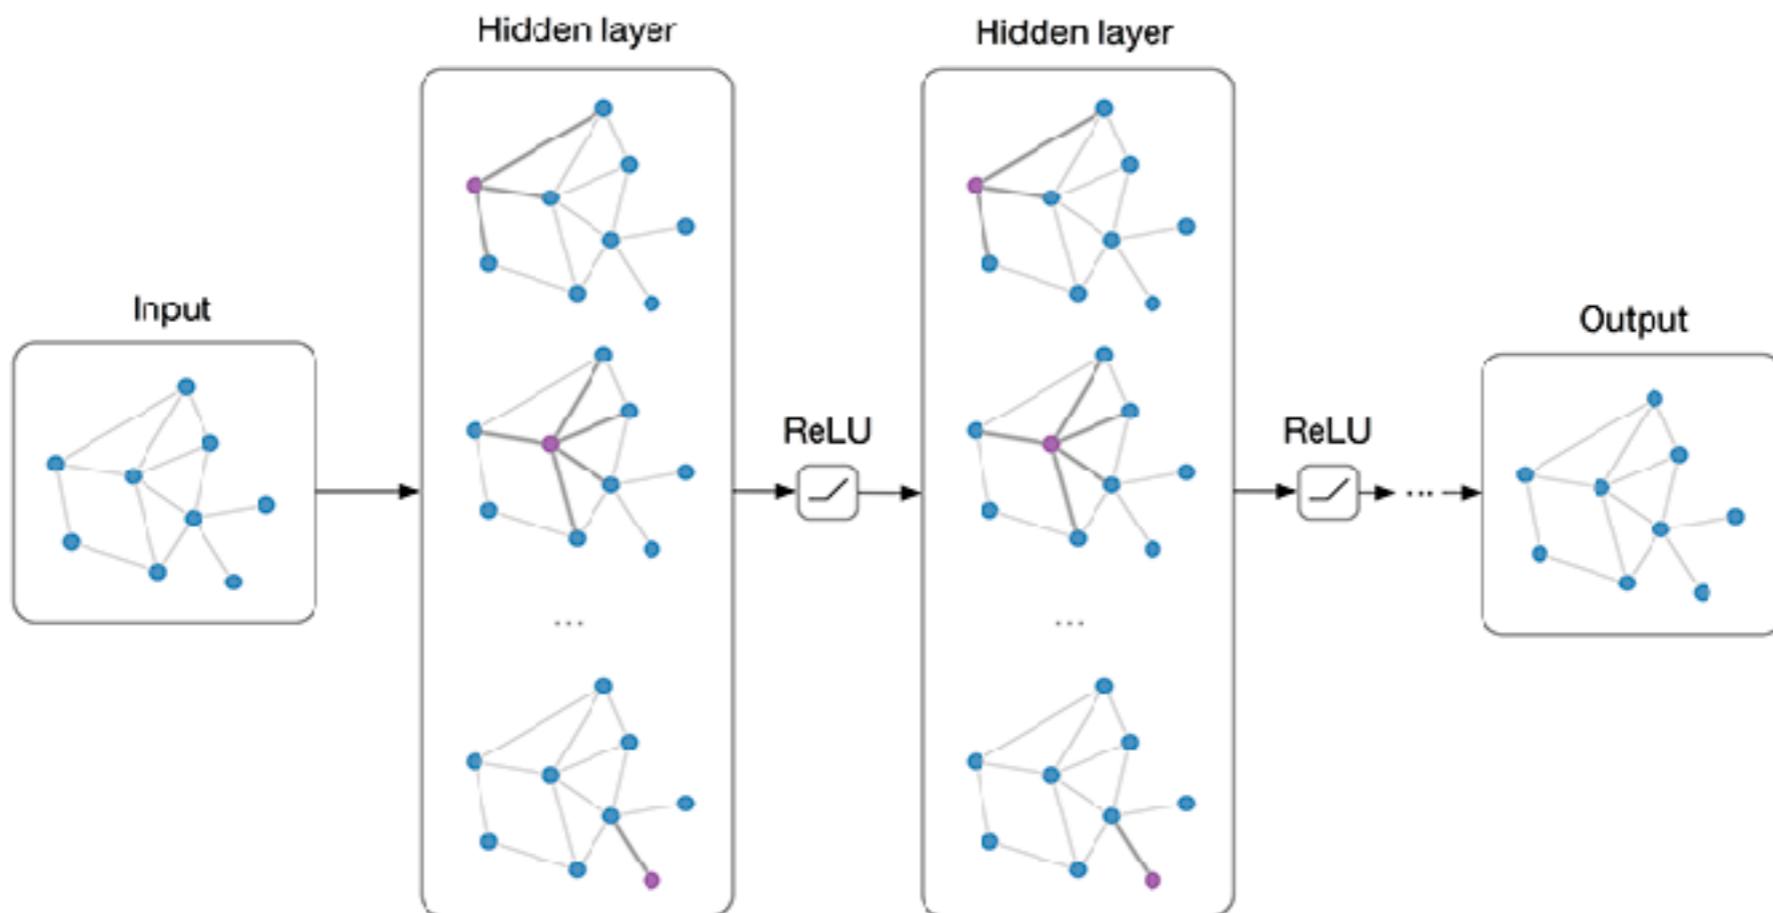


Message passing functions are trainable and differentiable

Each node uses the same set of shared parameters

Message Passing operators can be stacked:

After  $L$  layers, a node has aggregated  $L$ -hop neighborhood





Graph Neural Networks compute localized node embeddings  $\mathbf{h}_v^{(L)} \in \mathbb{R}^{|\mathcal{V}| \times h}$  for a given task:

### Node Classification:

Prediction based on node embedding

*End-to-end training!*

$$\phi(\mathcal{G}, v) = \text{MLP}(\mathbf{h}_v^{(L)})$$

### Graph Classification:

Prediction based on global readout of node embeddings

$$\phi(\mathcal{G}) = \text{MLP}\left(\sum_{v \in \mathcal{V}} \mathbf{h}_v^{(L)}\right)$$

### Link Prediction:

Prediction based on pairs of node embeddings

$$\phi(\mathcal{G}, v, w) = \text{MLP}(\mathbf{h}_v^{(L)}, \mathbf{h}_w^{(L)})$$



$$\mathbf{h}_v^{(\ell)} = \sigma \left( \mathbf{W}_1 \mathbf{h}_v^{(\ell-1)} + \mathbf{W}_2 \sum_{w \in \mathcal{N}(v)} C_{v,w} \mathbf{h}_w^{(\ell-1)} \right)$$

Non-  
Linearity

Skip-  
connection

Trainable  
parameters

Normalization  
coefficient

Xu et al.: How Powerful are Graph Neural Networks? (ICLR '19)

Hamilton et al.: Inductive Representation Learning on Large Graphs (NIPS '17)

Kipf and Welling: Semi-Supervised Classification with Graph Convolutional Networks (ICLR '17)

Velickovic et al.: Graph Attention Networks (ICLR '18)



$$\mathbf{h}_v^{(\ell)} = \sigma \left( \mathbf{W}_1 \mathbf{h}_v^{(\ell-1)} + \mathbf{W}_2 \sum_{w \in \mathcal{N}(v)} C_{v,w} \mathbf{h}_w^{(\ell-1)} \right)$$

Non-  
Linearity      Skip-  
connection      Trainable  
parameters      Normalization  
coefficient

Normalization can be either ...

... static:  $C_{v,w} = 1$

Xu et al.: How Powerful are Graph Neural Networks? (ICLR '19)

Hamilton et al.: Inductive Representation Learning on Large Graphs (NIPS '17)

Kipf and Welling: Semi-Supervised Classification with Graph Convolutional Networks (ICLR '17)

Velickovic et al.: Graph Attention Networks (ICLR '18)



$$\mathbf{h}_v^{(\ell)} = \sigma \left( \mathbf{W}_1 \mathbf{h}_v^{(\ell-1)} + \mathbf{W}_2 \sum_{w \in \mathcal{N}(v)} C_{v,w} \mathbf{h}_w^{(\ell-1)} \right)$$

Non-  
Linearity      Skip-  
connection      Trainable  
parameters      Normalization  
coefficient

Normalization can be either ...

... static:  $C_{v,w} = 1$

... structure-dependent:  $C_{v,w} = |\mathcal{N}(v)|^{-1}$

Xu et al.: How Powerful are Graph Neural Networks? (ICLR '19)

Hamilton et al.: Inductive Representation Learning on Large Graphs (NIPS '17)

Kipf and Welling: Semi-Supervised Classification with Graph Convolutional Networks (ICLR '17)

Velickovic et al.: Graph Attention Networks (ICLR '18)



$$\mathbf{h}_v^{(\ell)} = \sigma \left( \mathbf{W}_1 \mathbf{h}_v^{(\ell-1)} + \mathbf{W}_2 \sum_{w \in \mathcal{N}(v)} C_{v,w} \mathbf{h}_w^{(\ell-1)} \right)$$

Non-  
Linearity      Skip-  
connection      Trainable  
parameters      Normalization  
coefficient

Normalization can be either ...

... static:  $C_{v,w} = 1$

... structure-dependent:  $C_{v,w} = |\mathcal{N}(v)|^{-1}$

... data-dependent (learned), aka Attention

Xu et al.: How Powerful are Graph Neural Networks? (ICLR '19)

Hamilton et al.: Inductive Representation Learning on Large Graphs (NIPS '17)

Kipf and Welling: Semi-Supervised Classification with Graph Convolutional Networks (ICLR '17)

Velickovic et al.: Graph Attention Networks (ICLR '18)



The message passing scheme is really flexible:

It can model anisotropic transformations:

$$\text{MESSAGE}(\mathbf{h}_v^{(\ell)}, \mathbf{h}_w^{(\ell)}) = \text{MLP}(\mathbf{h}_v^{(\ell)}, \mathbf{h}_w^{(\ell)} - \mathbf{h}_v^{(\ell)})$$

Wang et al.: Dynamic Graph CNN for Learning on Point Clouds? ('19)

Schlichtkrull et al.: Modeling Relational Data with Graph Convolutional Networks ('17)

Fey et al.: SplineCNN: Fast Geometric Deep Learning with Continuous B-Spline Kernels (CVPR '18) 18



The message passing scheme is really flexible:

It can model anisotropic transformations:

$$\text{MESSAGE}(\mathbf{h}_v^{(\ell)}, \mathbf{h}_w^{(\ell)}) = \text{MLP}(\mathbf{h}_v^{(\ell)}, \mathbf{h}_w^{(\ell)} - \mathbf{h}_v^{(\ell)})$$

It can model different edge relation types:

$$\text{MESSAGE}(\mathbf{h}_v^{(\ell)}, \mathbf{h}_w^{(\ell)}) = \mathbf{W}_{R(v,w)} \mathbf{h}_w^{(\ell)}$$

Wang et al.: Dynamic Graph CNN for Learning on Point Clouds? ('19)

Schlichtkrull et al.: Modeling Relational Data with Graph Convolutional Networks ('17)

Fey et al.: SplineCNN: Fast Geometric Deep Learning with Continuous B-Spline Kernels (CVPR '18) 18



The message passing scheme is really flexible:

It can model anisotropic transformations:

$$\text{MESSAGE}(\mathbf{h}_v^{(\ell)}, \mathbf{h}_w^{(\ell)}) = \text{MLP}(\mathbf{h}_v^{(\ell)}, \mathbf{h}_w^{(\ell)} - \mathbf{h}_v^{(\ell)})$$

It can model different edge relation types:

$$\text{MESSAGE}(\mathbf{h}_v^{(\ell)}, \mathbf{h}_w^{(\ell)}) = \mathbf{W}_{R(v,w)} \mathbf{h}_w^{(\ell)}$$

It can incorporate edge features:

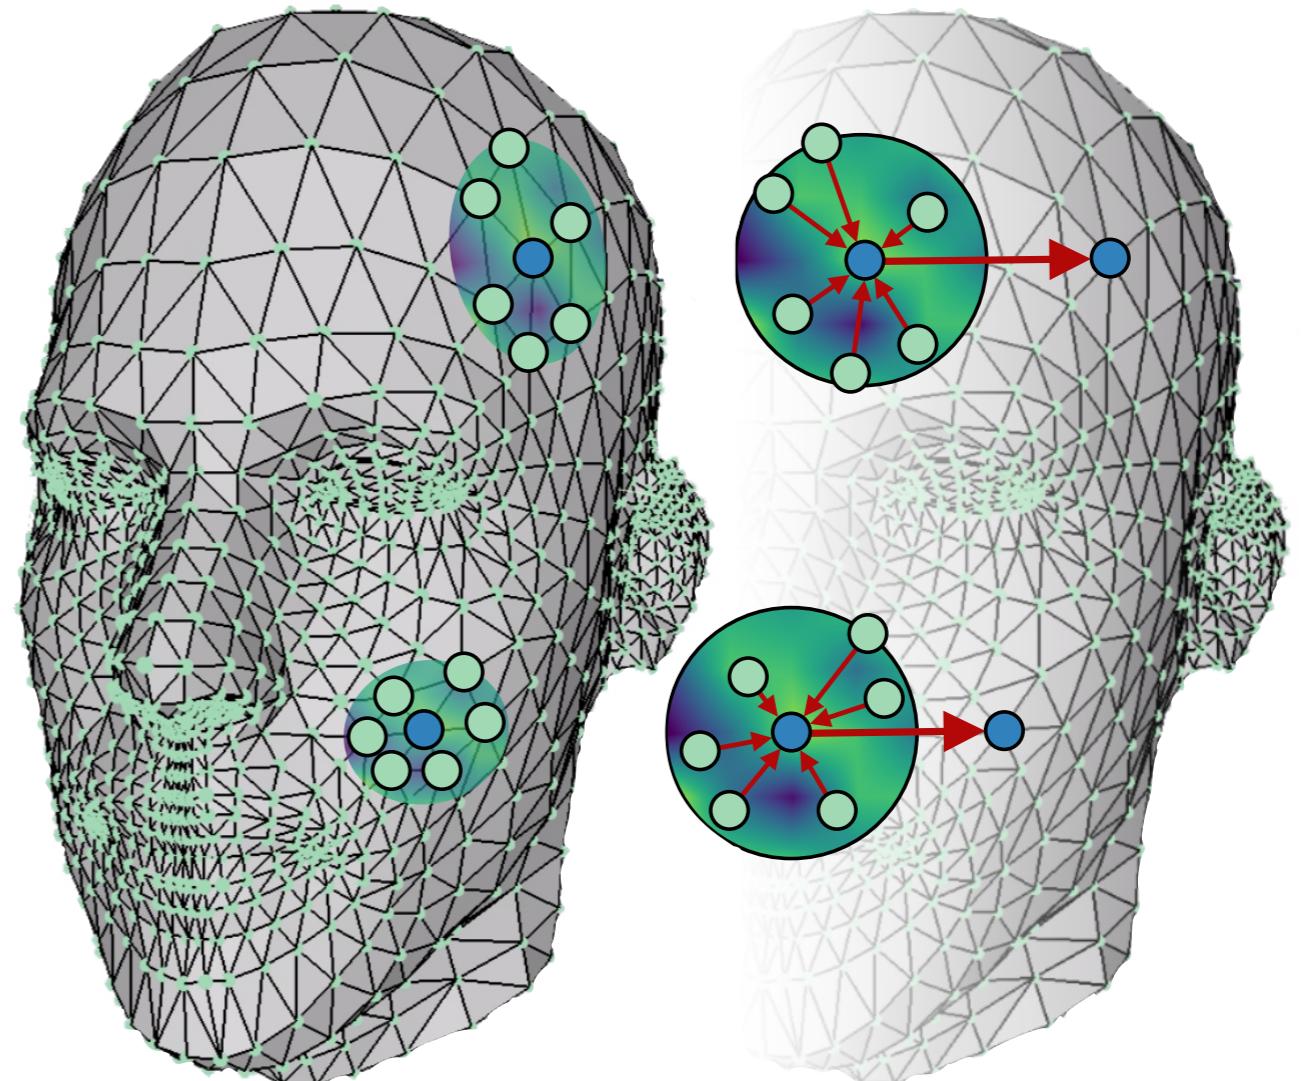
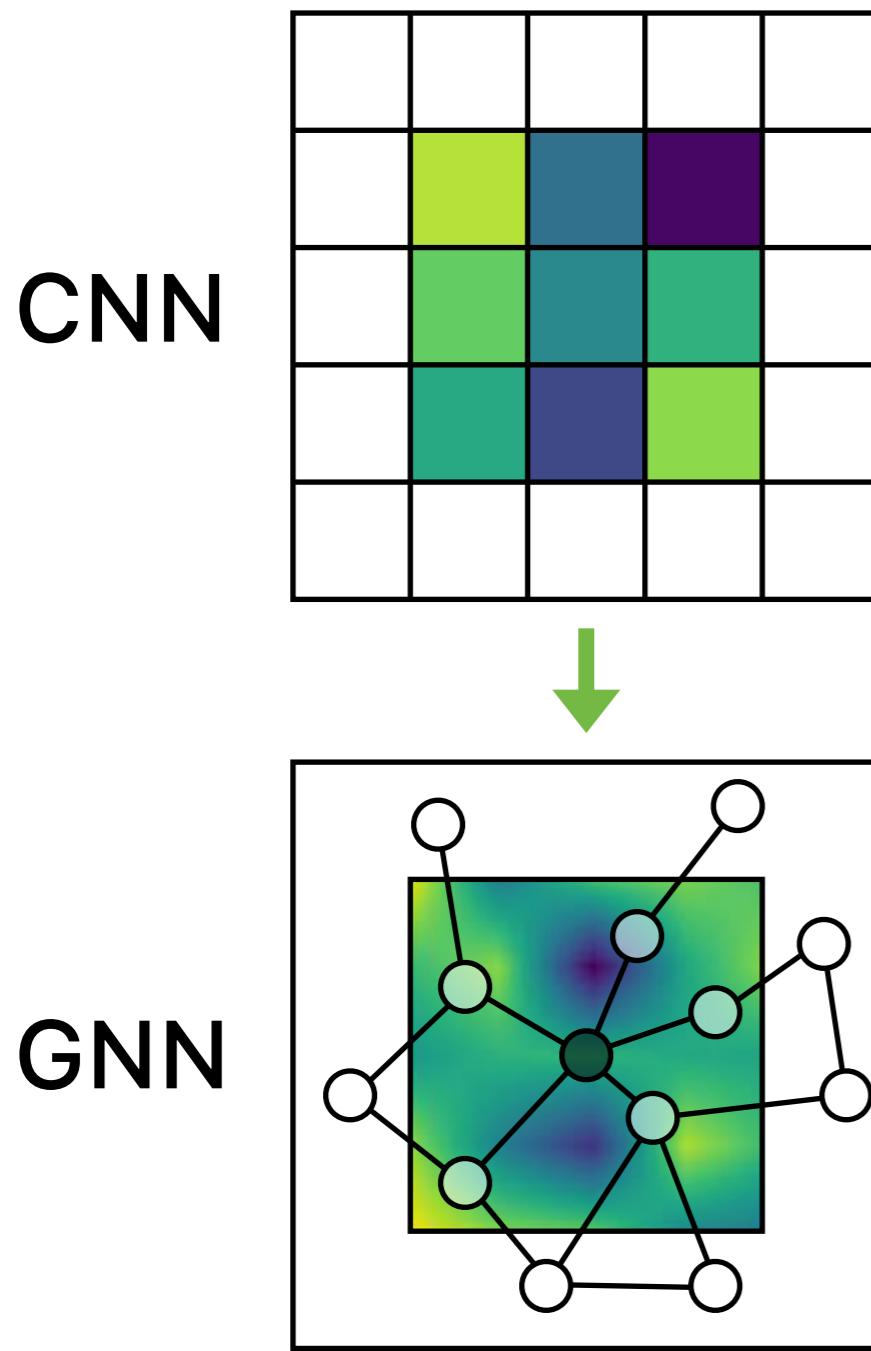
$$\text{MESSAGE}(\mathbf{h}_w^{(\ell)}, \mathbf{e}_{w,v}) = \text{MLP}(\mathbf{e}_{w,v}) \mathbf{h}_w^{(\ell)}$$

Wang et al.: Dynamic Graph CNN for Learning on Point Clouds? ('19)

Schlichtkrull et al.: Modeling Relational Data with Graph Convolutional Networks ('17)

Fey et al.: SplineCNN: Fast Geometric Deep Learning with Continuous B-Spline Kernels (CVPR '18) 18

$$\text{MESSAGE}(\mathbf{h}_w^{(\ell)}, \mathbf{e}_{w,v}) = \text{MLP}(\mathbf{e}_{w,v})\mathbf{h}_w^{(\ell)}$$





**Assumption:** We know what the nodes and edges are

Kipf and Welling: Graph Auto-Encoders (NIPS-W '16)

Schlichtkrull et al.: Knowledge Graph Completion ('17)

Kipf et al.: Neural Relational Inference for Interacting Systems (ICML '18)

Kipf et al.: Contrastive Learning of Structured World Models (ICLR '20)



**Assumption:** We know what the nodes and edges are

What if we have incomplete knowledge about relations?

- ▶ *Link prediction (graph completion)*
- ▶ *Infer links from scratch (relational inference)*

Kipf and Welling: Graph Auto-Encoders (NIPS-W '16)

Schlichtkrull et al.: Knowledge Graph Completion ('17)

Kipf et al.: Neural Relational Inference for Interacting Systems (ICML '18)

Kipf et al.: Contrastive Learning of Structured World Models (ICLR '20)



**Assumption:** We know what the nodes and edges are

What if we have incomplete knowledge about relations?

- ▶ *Link prediction (graph completion)*
- ▶ *Infer links from scratch (relational inference)*

What if we do not even know about nodes?

- ▶ *Object / entity / event discovery*
- ▶ *Grouping / clustering of low-level features*

Kipf and Welling: Graph Auto-Encoders (NIPS-W '16)

Schlichtkrull et al.: Knowledge Graph Completion ('17)

Kipf et al.: Neural Relational Inference for Interacting Systems (ICML '18)

Kipf et al.: Contrastive Learning of Structured World Models (ICLR '20)



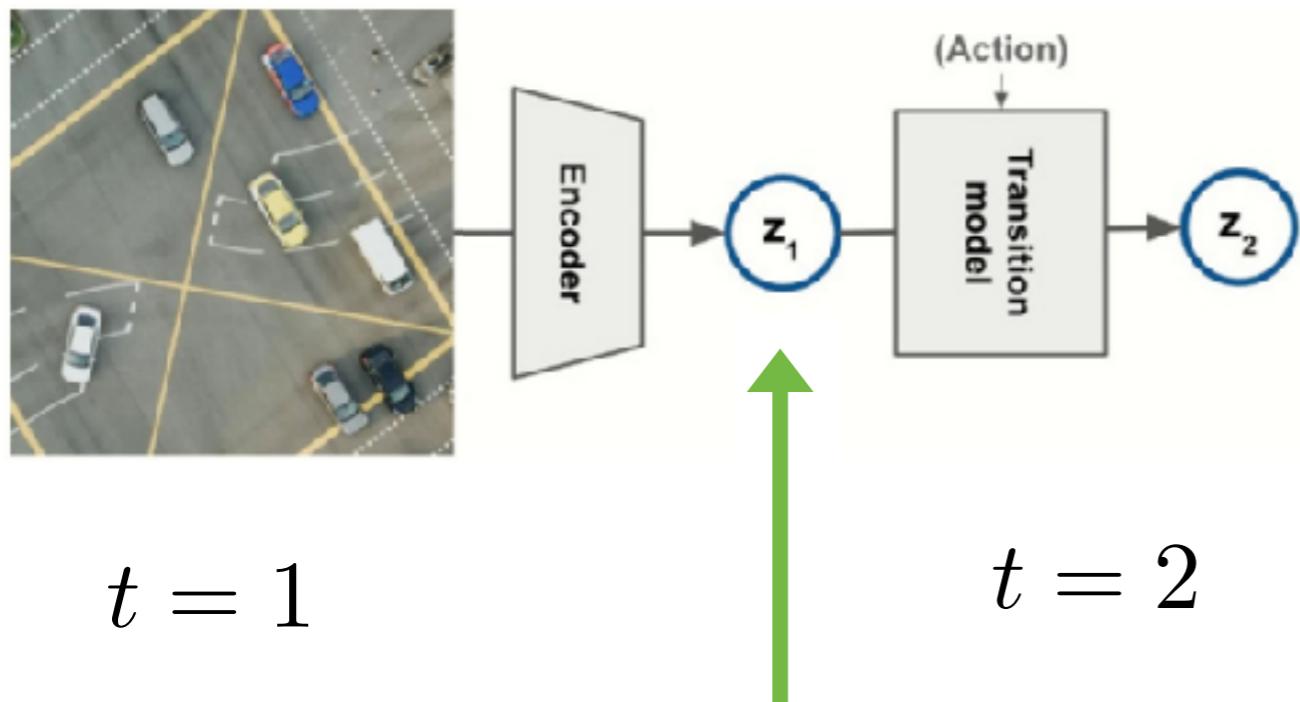
What if we do not even know about nodes?





What if we do not even know about nodes?

## "Unstructured" World Model



Representation is  
fully distributed  
and unstructured

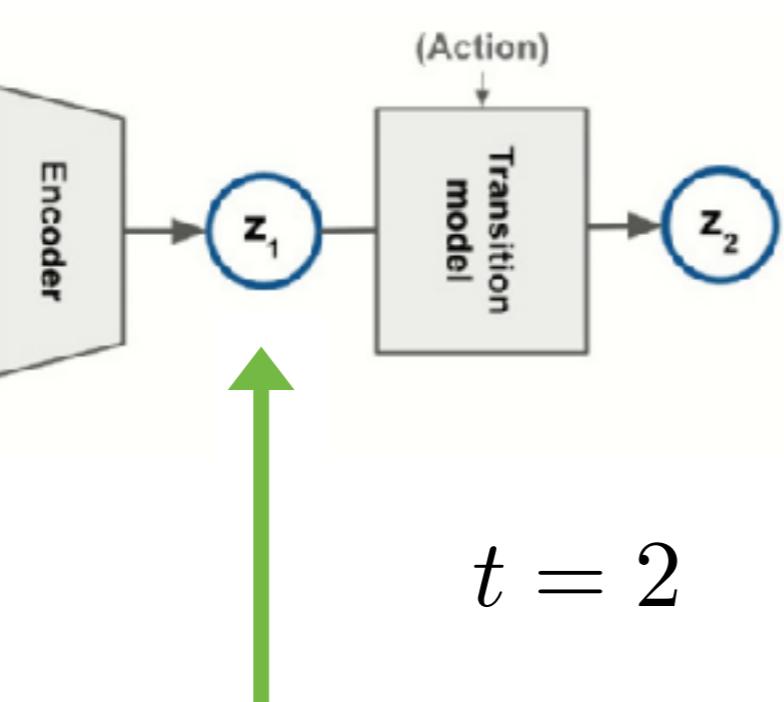


What if we do not even know about nodes?

### "Unstructured" World Model



$t = 1$



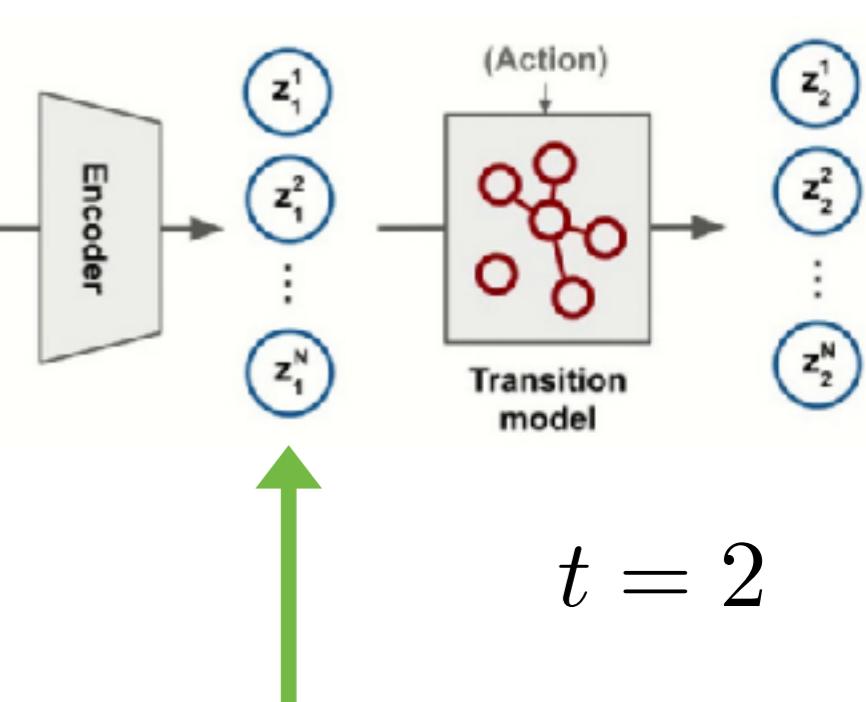
$t = 2$

Representation is  
fully distributed  
and unstructured

### "Structured" World Model



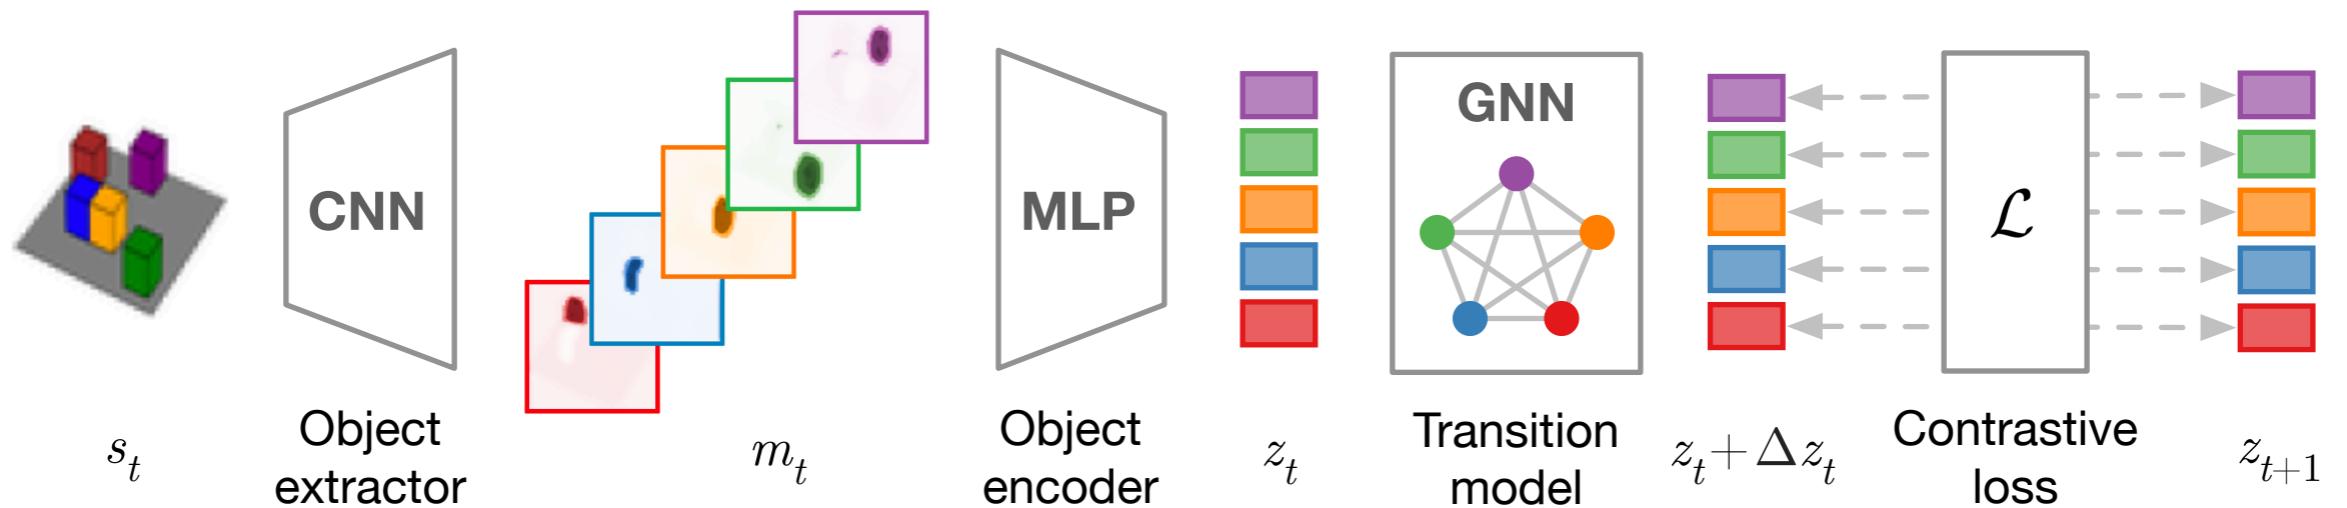
$t = 1$



$t = 2$

Representation is  
fully distributed, but  
structured as a set

## Fully self-supervised "structured" world models



### Encoder:

- ▶ "Object discovery" module that maps observations to a set of vectors

### Transition model:

- ▶ Models pairwise interactions between objects using a GNN

**Training:** Using a contrastive loss in latent space based on a random state in the dataset







Graph Neural Networks introduce their own unique challenges regarding implementation

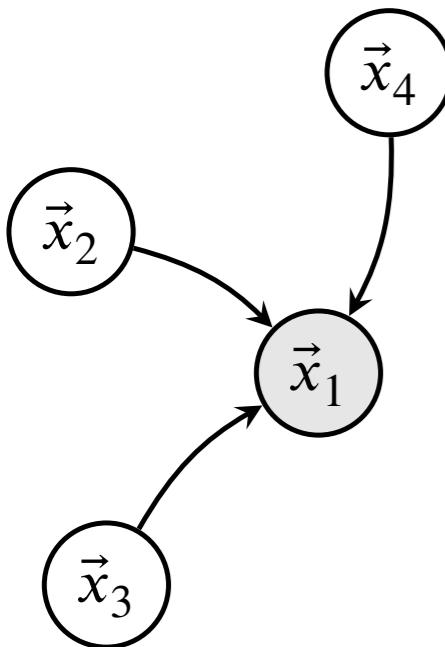
Most of the deep learning frameworks are designed for dense representations!

- ▶ *How can we achieve effective parallelization on irregular and sparse graphs?*
- ▶ *How can we provide mini-batching capabilities for graphs of potentially different sizes?*

Given a *sparse graph*  $\mathcal{G} = (\mathbf{X}, (\mathbf{I}, \mathbf{E}))$  with

$$\mathbf{X} = \begin{bmatrix} \vec{x}_1 \\ \vec{x}_2 \\ \vec{x}_3 \\ \vec{x}_4 \end{bmatrix}$$

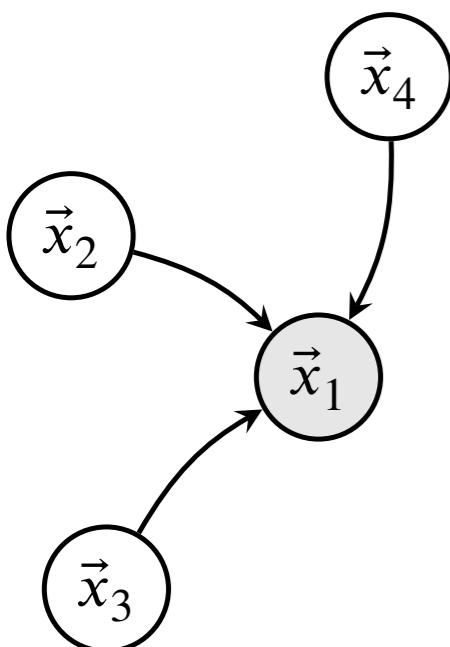
- ▶ **node features**  $\mathbf{X} \in \mathbb{R}^{|\mathcal{V}| \times F}$
- ▶ **edge indices**  $\mathbf{I} \in \{1, \dots, N\}^{2 \times |\mathcal{E}|}$
- ▶ *optional* **edge features**  $\mathbf{E} \in \mathbb{R}^{|\mathcal{E}| \times D}$



$$\mathbf{I} = \begin{bmatrix} 2 & 1 \\ 3 & 1 \\ 4 & 1 \end{bmatrix}^\top$$

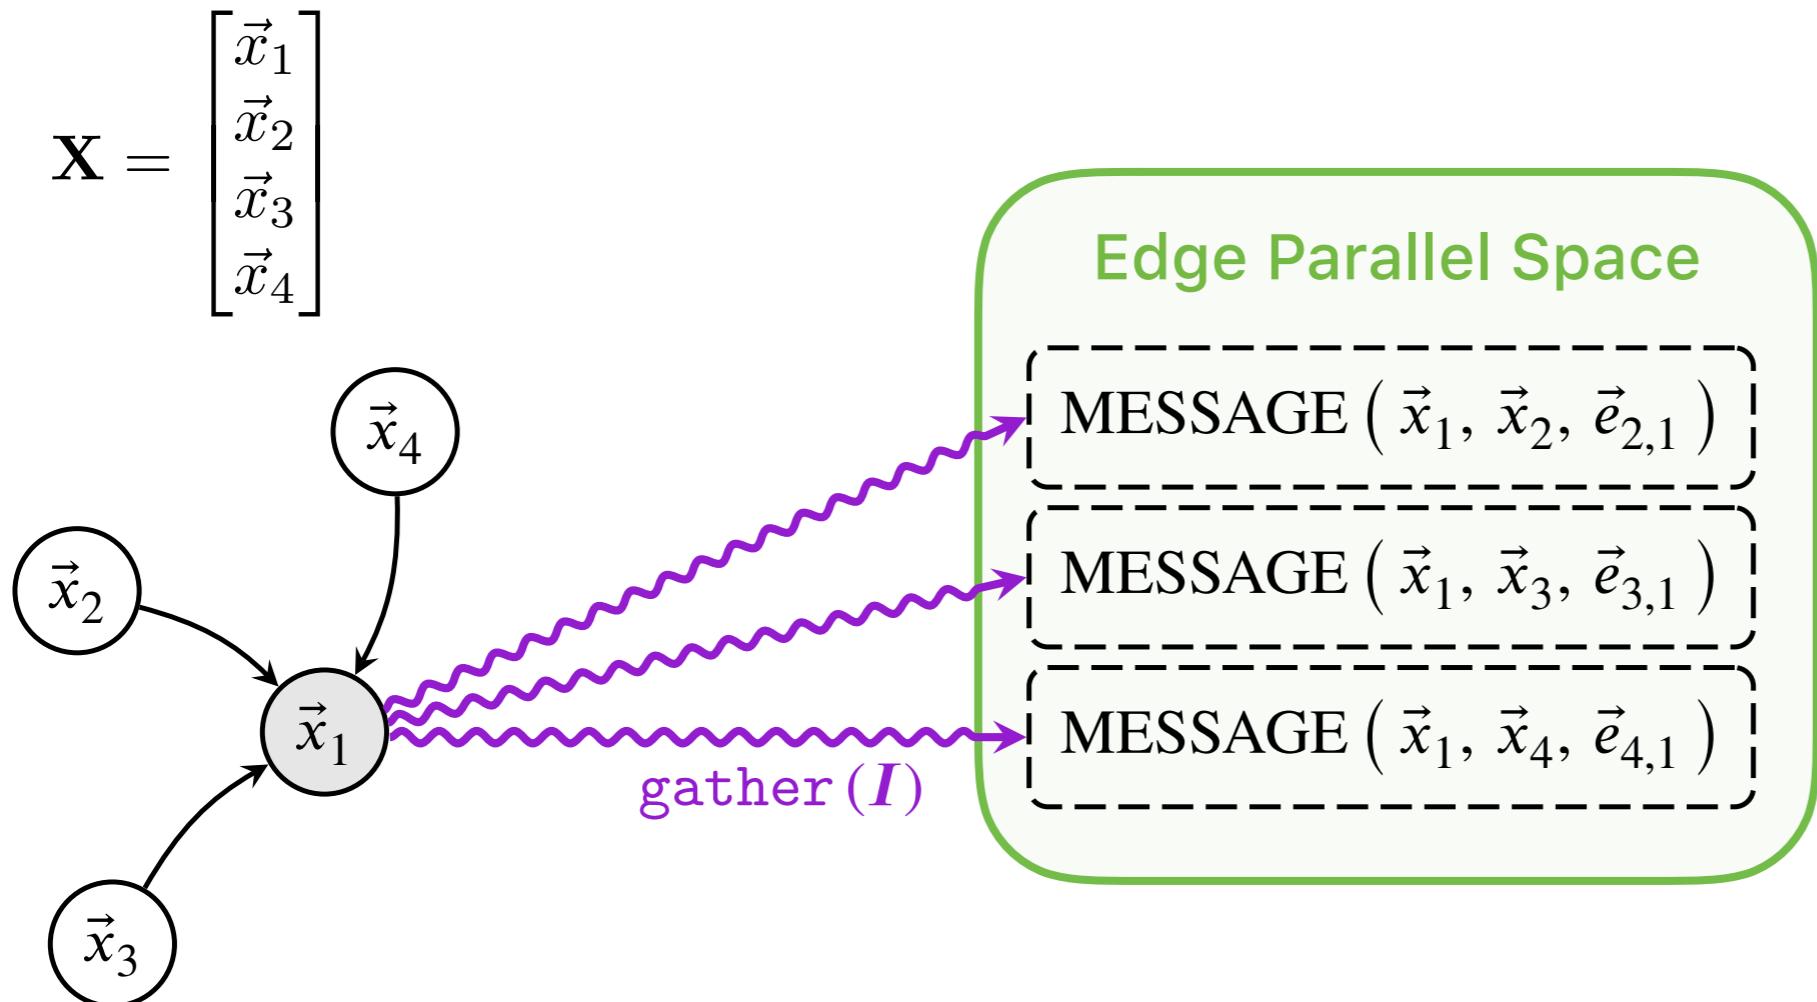
Given a sparse graph  $\mathcal{G} = (\mathbf{X}, (\mathbf{I}, \mathbf{E}))$

$$\mathbf{X} = \begin{bmatrix} \vec{x}_1 \\ \vec{x}_2 \\ \vec{x}_3 \\ \vec{x}_4 \end{bmatrix}$$



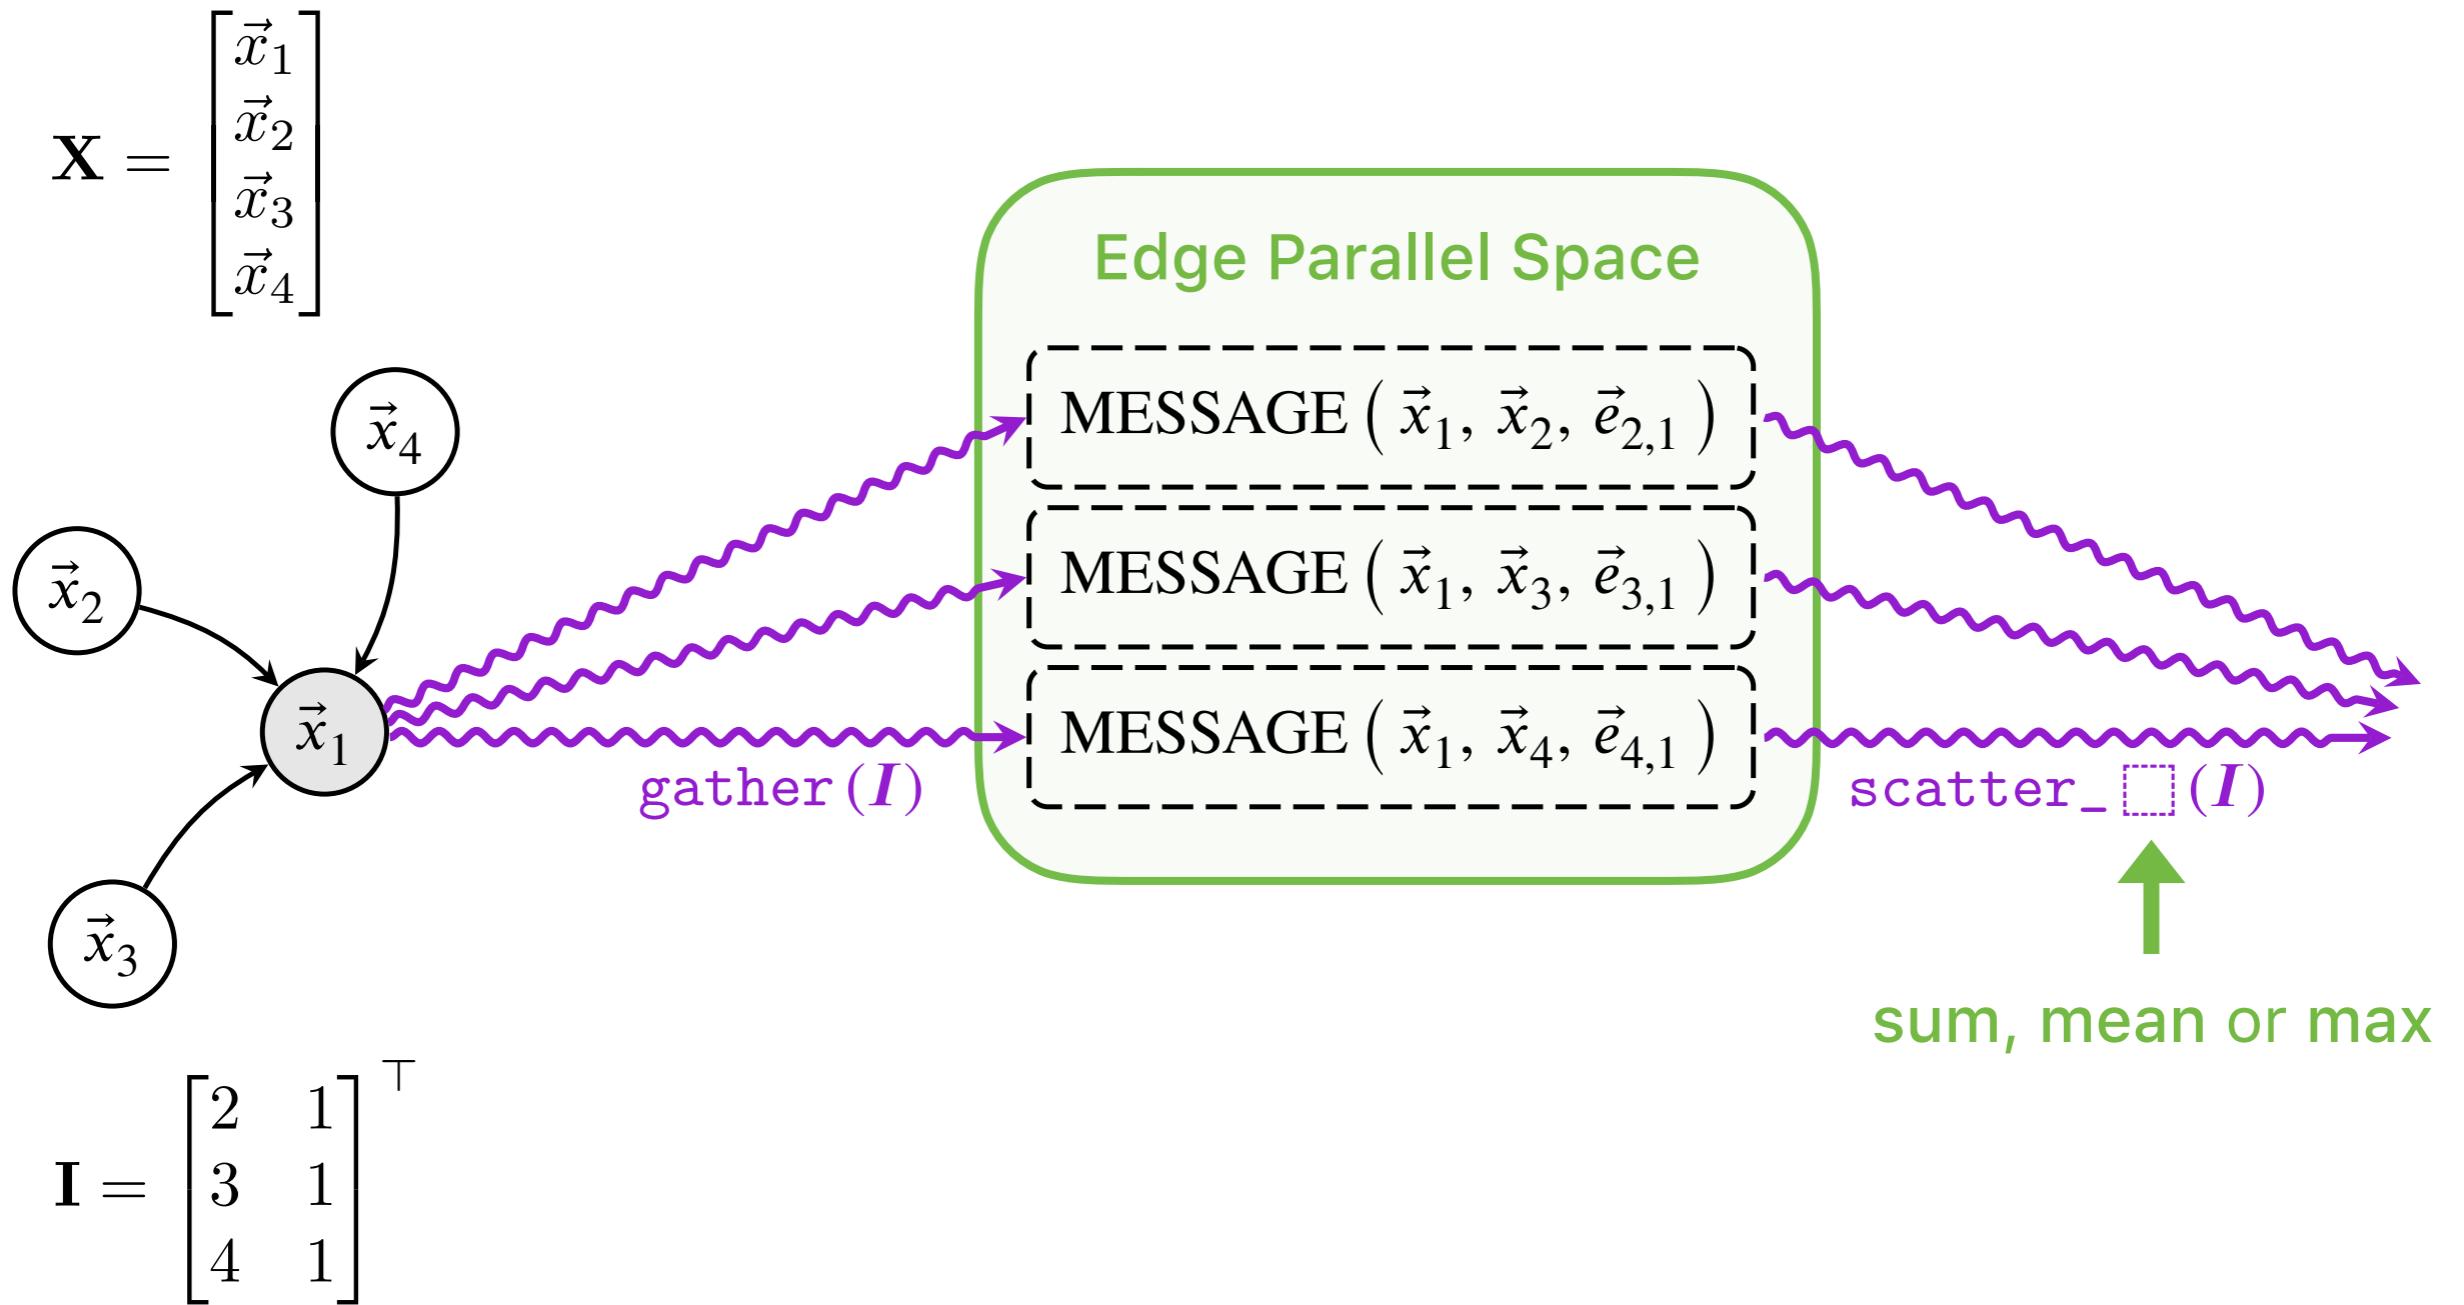
$$\mathbf{I} = \begin{bmatrix} 2 & 1 \\ 3 & 1 \\ 4 & 1 \end{bmatrix}^\top$$

Given a sparse graph  $\mathcal{G} = (\mathbf{X}, (\mathbf{I}, \mathbf{E}))$

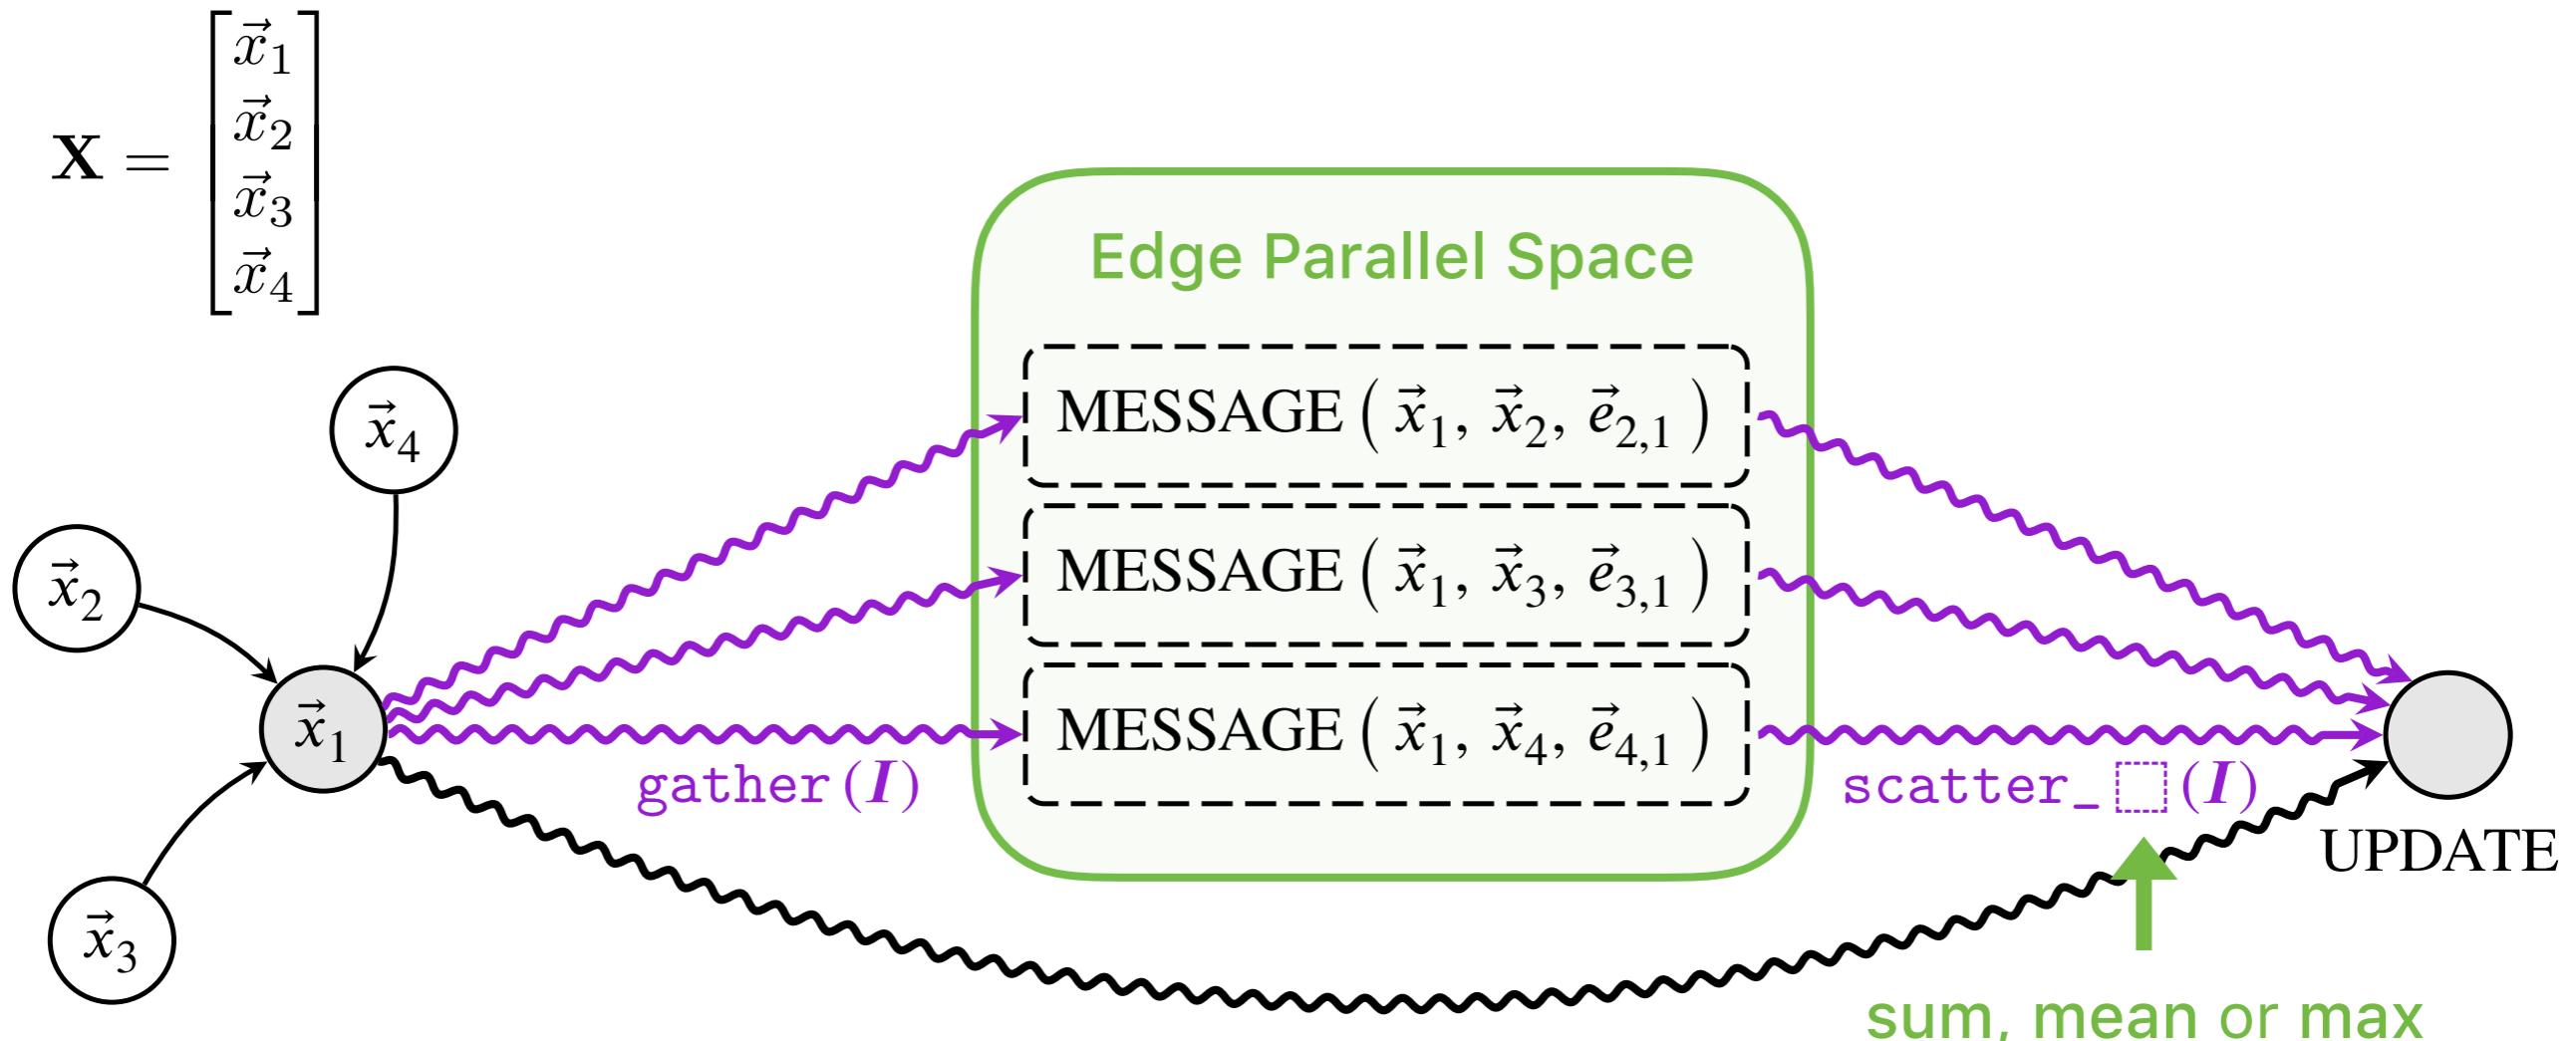


$$\mathbf{I} = \begin{bmatrix} 2 & 1 \\ 3 & 1 \\ 4 & 1 \end{bmatrix}^\top$$

Given a sparse graph  $\mathcal{G} = (\mathbf{X}, (\mathbf{I}, \mathbf{E}))$

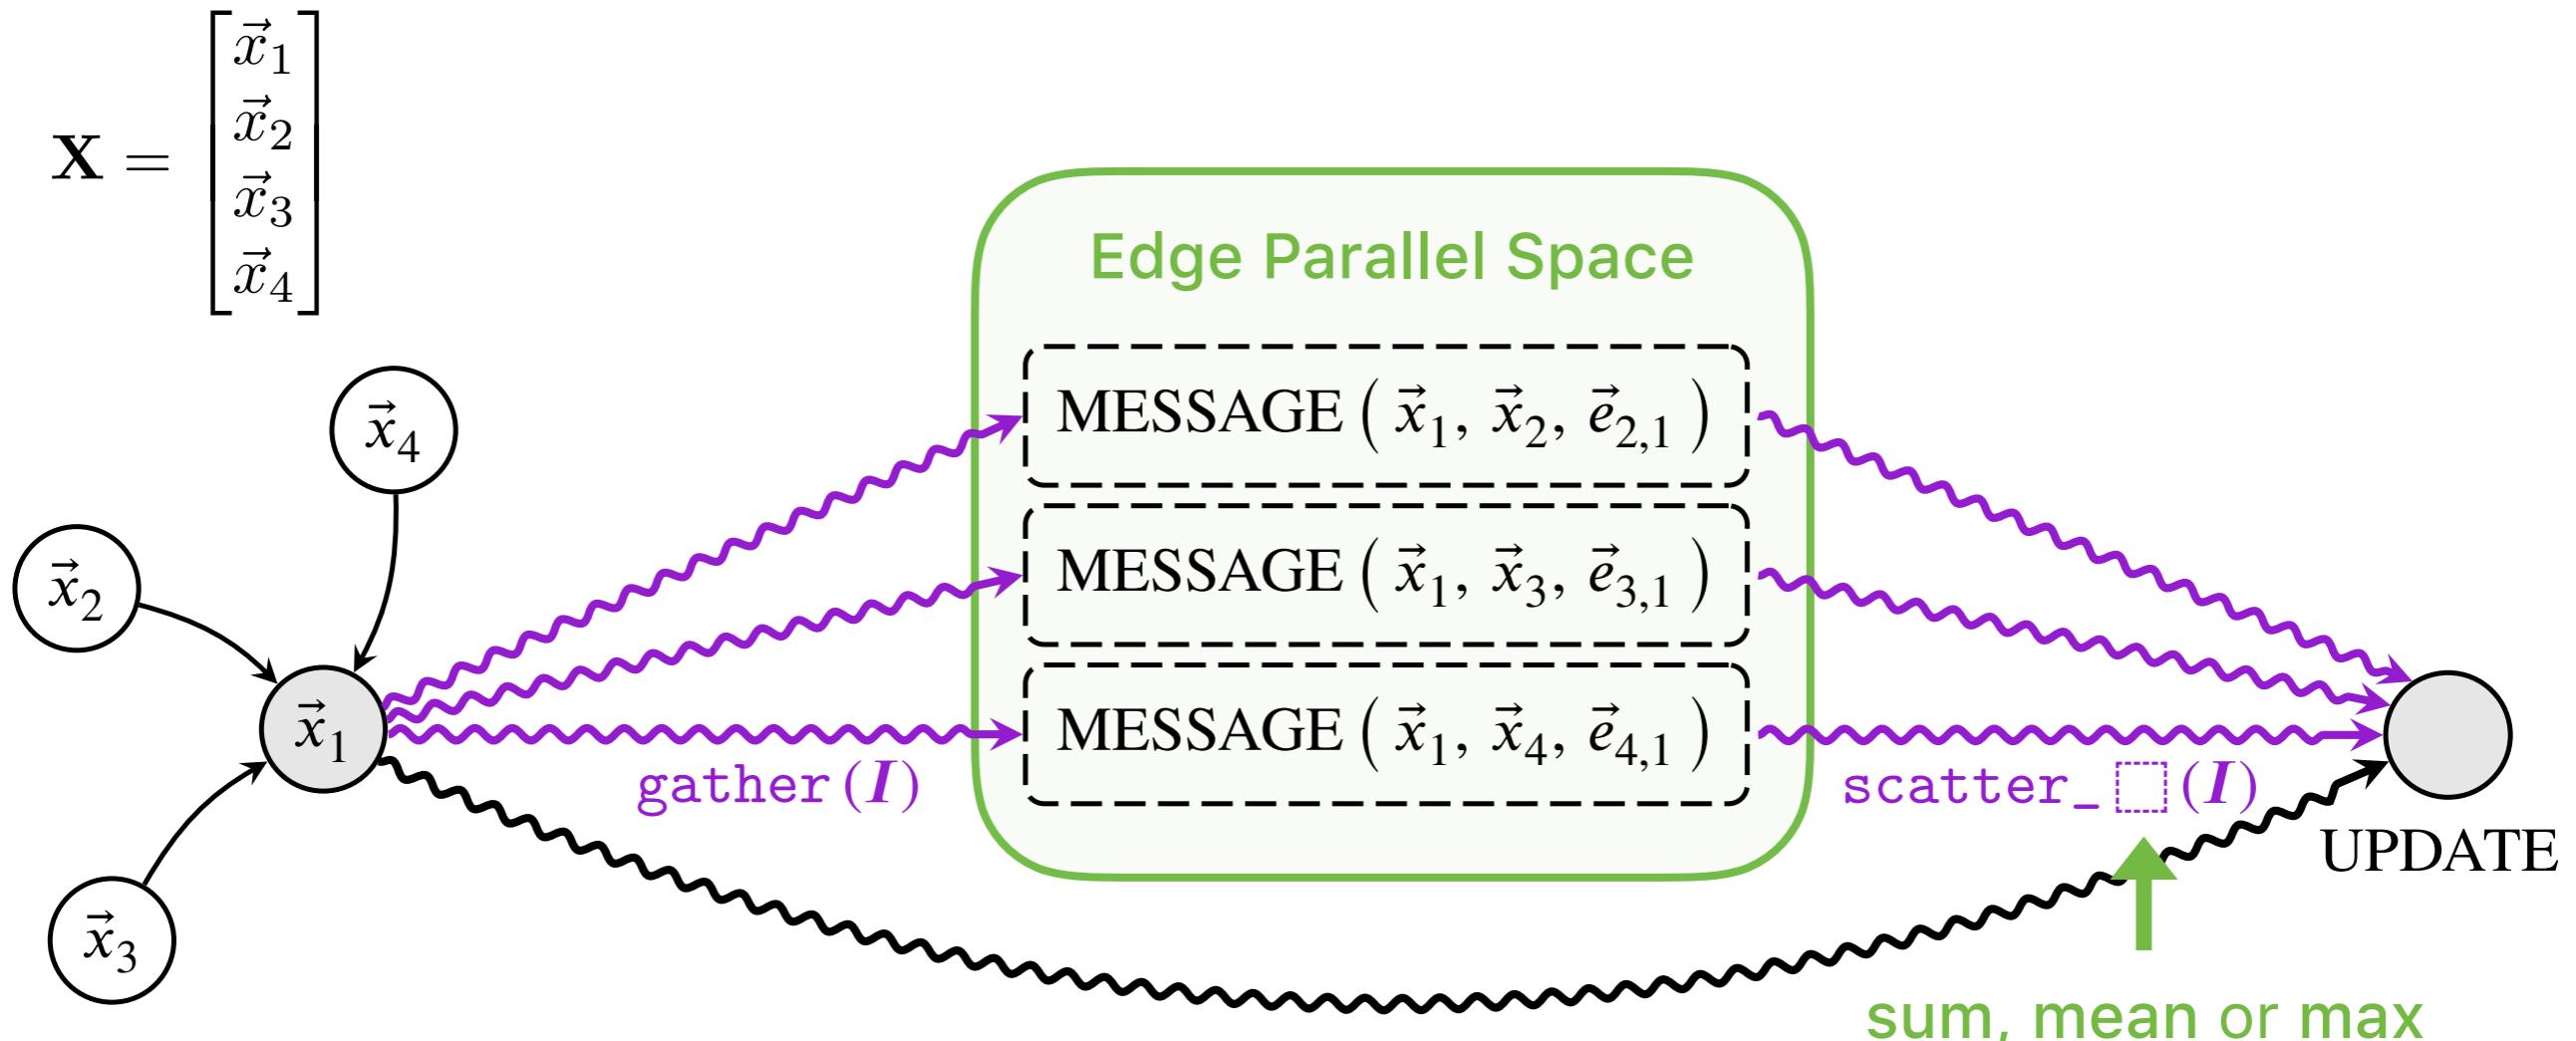


Given a sparse graph  $\mathcal{G} = (\mathbf{X}, (\mathbf{I}, \mathbf{E}))$



$$\mathbf{I} = \begin{bmatrix} 2 & 1 \\ 3 & 1 \\ 4 & 1 \end{bmatrix}^\top$$

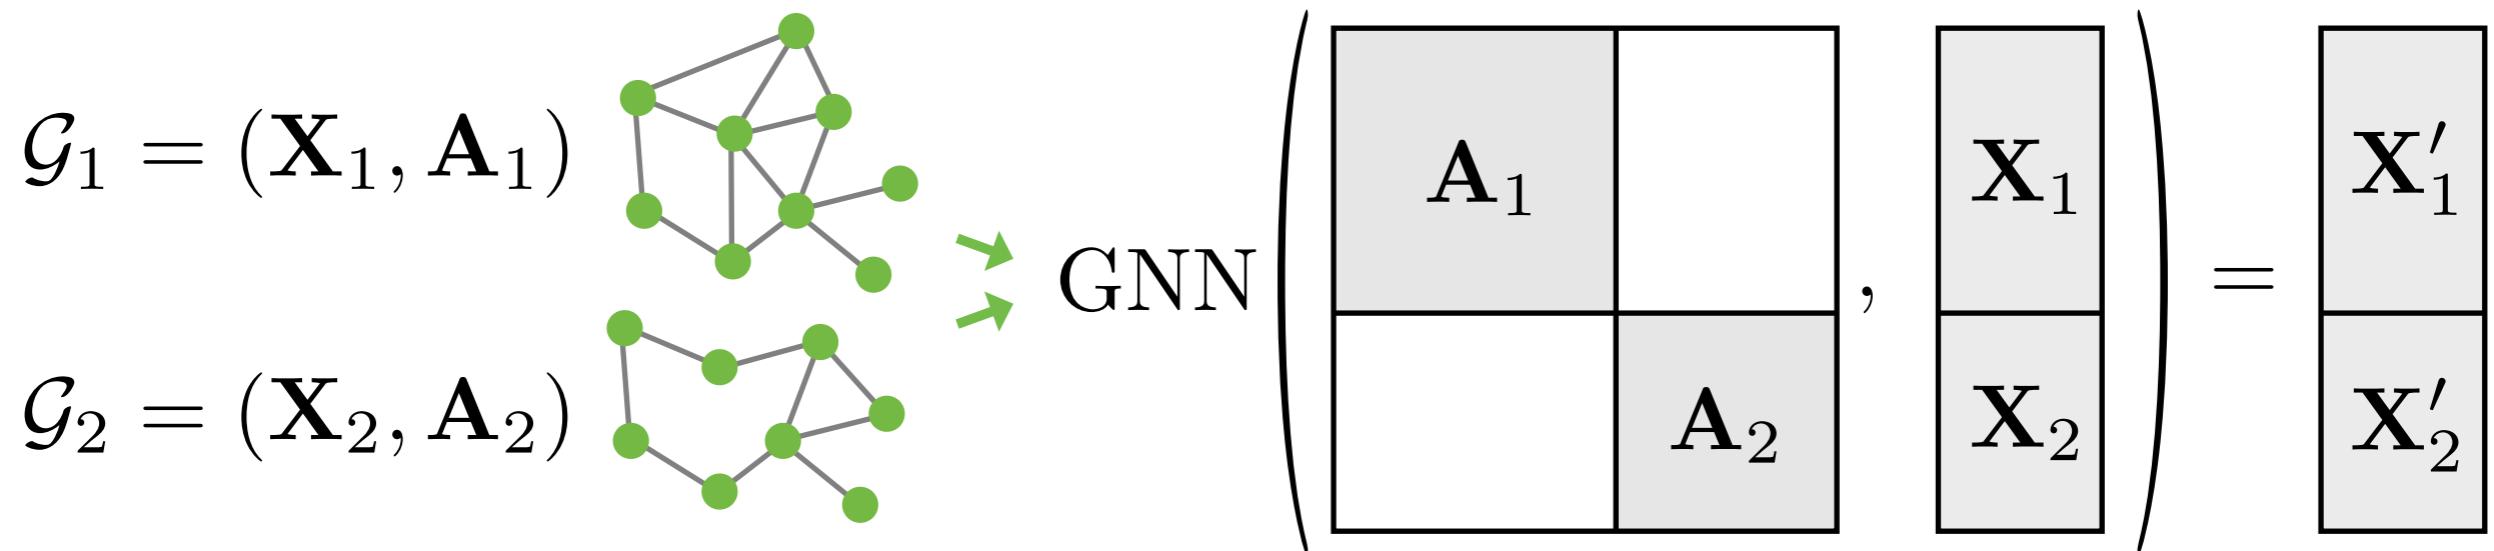
Given a sparse graph  $\mathcal{G} = (\mathbf{X}, (\mathbf{I}, \mathbf{E}))$



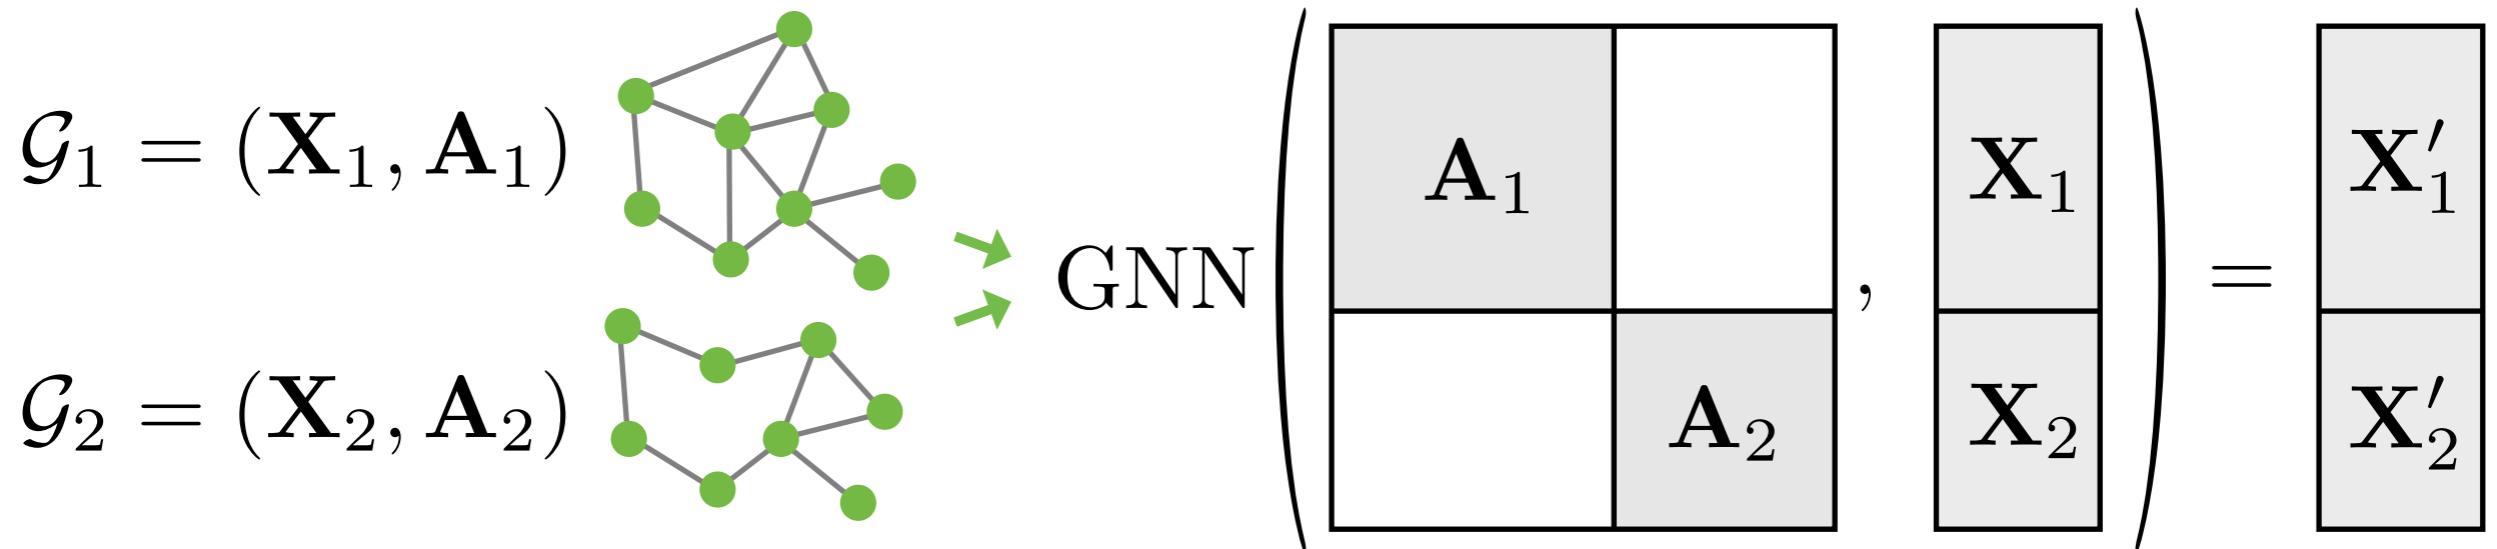
$$\mathbf{I} = \begin{bmatrix} 2 & 1 \\ 3 & 1 \\ 4 & 1 \end{bmatrix}^\top$$

Flexible implementation via  
gather/scatter operations!

## Mini-batch handling of different graphs:

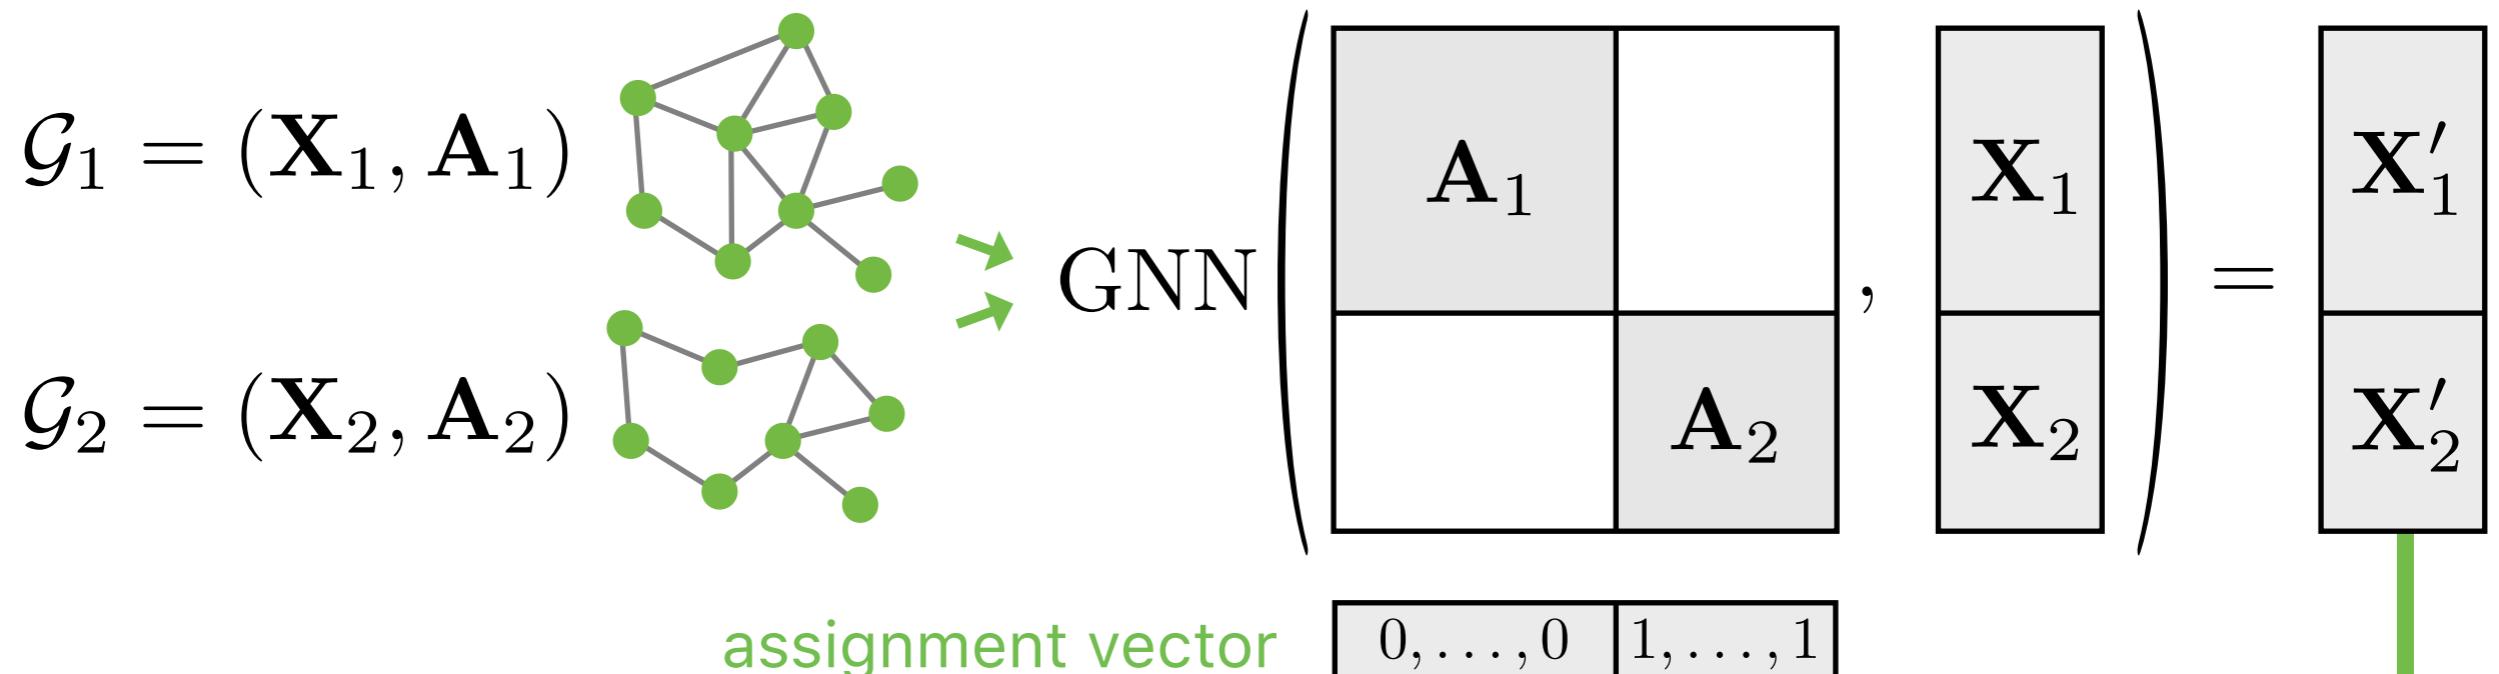


## Mini-batch handling of different graphs:

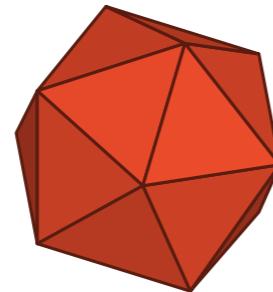


- ▶ no GNN modifications needed
- ▶ no memory/computation overhead
- ▶ supports examples of different size

## Mini-batch handling of different graphs:



- ▶ no GNN modifications needed
- ▶ no memory/computation overhead
- ▶ supports examples of different size



# PyTorch geometric

A Python library built upon PyTorch to enable deep learning on graphs:

- ▶ simplifies implementing and working with **Graph Neural Networks**
- ▶ omits the need of re-writing common logic (operators, datasets, batching, ...)
- ▶ bundles **fast implementations** from published papers



## GNNs

Cheby GCN SAGE PointNet MoNet MPNN GAT  
SplineCNN AGNN EdgeCNN S-GCN R-GCN PointCNN  
ARMA APPNP GIN GIN-E CG GatedGCN NMF TAG  
Signed-GCN DNA PPFNet FeaST Hyper-GCN GravNet

## Pooling

Set2Set SortPool DiffPool MinCUT Graclus  
VoxelGrid TopK SAG EdgePool ASAP

## Models

(V)GAE ARG(V)A DGI Node2Vec GraphUNet  
GeniePath SchNet DimeNet MetaPath2Vec ReNet

## Utilities

GNNExplainer SIGN GDC ClusterGCN DropEdge  
GraphSAINT NeighborSampling GraphSizeNorm JK





## 1. Data Loading Routine

```
from torch_geometric.datasets import PPI
dataset = PPI('~/datasets/PPI', split='train')

print(dataset)
>>> PPI(20)

print(dataset[0])
>>> Data(edge_index=[2, 32318], x=[1767, 50], y=[1767, 121])
```



## 1. Data Loading Routine

```
from torch_geometric.datasets import PPI
dataset = PPI('~/datasets/PPI', split='train')

print(dataset)
>>> PPI(20)

print(dataset[0])
>>> Data(edge_index=[2, 32318], x=[1767, 50], y=[1767, 121])

from torch_geometric.data import DataLoader
loader = DataLoader(dataset, batch_size=2, shuffle=True)

print(next(iter(loader)))
>>> Batch(batch=[3422], edge_index=[2, 82434], x=[3422, 50], y=[3422, 121])
```



## 2. Defining the Network Architecture

```
from torch_geometric.nn import GCNConv

class Net(torch.nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = GCNConv(dataset.num_features, 256)
        self.conv2 = GCNConv(256, dataset.num_classes)

    def forward(self, x, edge_index):
        x = self.conv1(x, edge_index)
        x = x.relu()
        x = self.conv2(x, edge_index)
        return x
```



### 3. Training the Model

```
model = Net().cuda()
optimizer = torch.optim.Adam(model.parameters(), lr=0.01)

def train():
    model.train()

    for data in loader:
        data = data.cuda()
        logits = model(data.x, data.edge_index)
        loss = compute_loss(logits, data.y)
        loss.backward()
        optimizer.step()
```



Training on giant graphs:



## Training on giant graphs:

```
def train(data): # Full-batch training  
    logits = model(data.x, data.edge_index)  
    loss = compute_loss(logits, data.y)
```

**RuntimError: CUDA error: out of memory**



## Training on giant graphs:

```
def train(data): # Full-batch training
    logits = model(data.x, data.edge_index)
    loss = compute_loss(logits, data.y)
```

RuntimeError: CUDA error: out of memory



```
+ loader = ClusterLoader( # Cluster-GCN
+     ClusterData(data, num_parts=128), batch_size=32)
```

```
~ def train(loader): # Mini-batch training
+     for data in loader:
        out = model(data.x, data.edge_index)
        loss = compute_loss(logits, data.y)
```



```
class MyOwnConv(MessagePassing):          add, mean or max  
    def __init__(self):  
        super(MyOwnConv, self).__init__(aggr='add')      aggregation
```



```
class MyOwnConv(MessagePassing):           add, mean or max  
    def __init__(self):                     aggregation  
        super(MyOwnConv, self).__init__(aggr='add')  
  
    def forward(self, x, edge_index, edge_weight):  
        return self.propagate(edge_index, x=x,  
                               edge_weight=edge_weight)  
  
                                         pass every-  
                                         thing needed for  
                                         propagation
```



```
class MyOwnConv(MessagePassing):                                add, mean or max
    def __init__(self):
        super(MyOwnConv, self).__init__(aggr='add')                aggregation

    def forward(self, x, edge_index, edge_weight):
        return self.propagate(edge_index, x=x,
                              edge_weight=edge_weight)

    def message(self, x_j, x_i, edge_weight):
        return x_j * edge_weight                                    pass every-
                                                               thing needed for
                                                               propagation
```

Node features get automatically lifted to source (\_j) and target (\_i) nodes



- ✓ uniform implementations of over 40 GNN operators/models
- ✓ extendable by using a simple message passing interface
- ✓ access to over 100 benchmark datasets
- ✓ automatic mini-batching for graphs with different sizes
- ✓ deterministic and differentiable pooling operators
- ✓ basic and more sophisticated readout functions
- ✓ useful transforms for augmentation, point sampling, ...
- ✓ leverages dedicated CUDA kernels

 [/rusty1s/pytorch\\_geometric](https://github.com/rusty1s/pytorch_geometric)

license [MIT](#)

PRs [welcome](#)

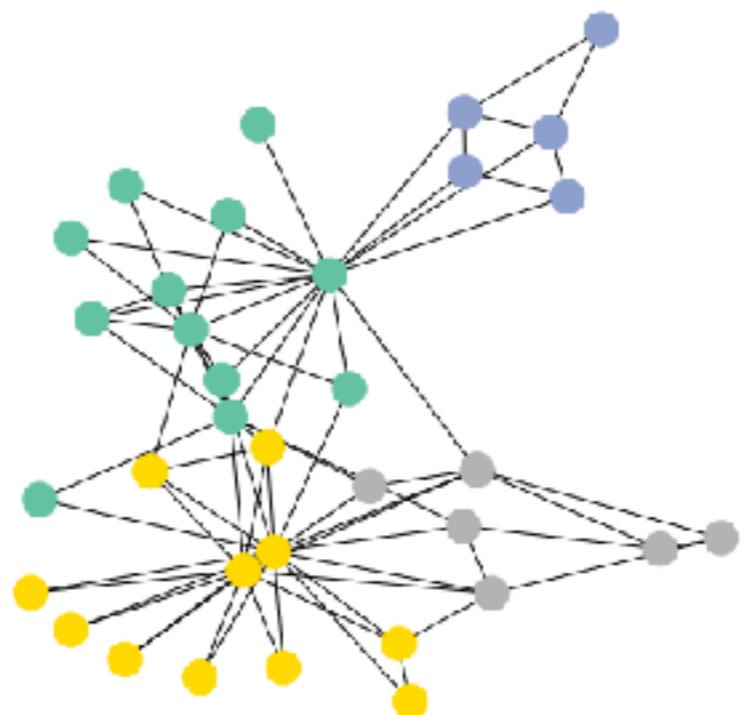


# Embedding the Karate Club Network!



pypi package 1.6.1 build passing docs passing codecov 87% contributions welcome

[Documentation](#) | [Paper](#) | [Colab Notebooks](#) | [External Resources](#) | [OGB Examples](#)





# Embedding the Karate Club Network!

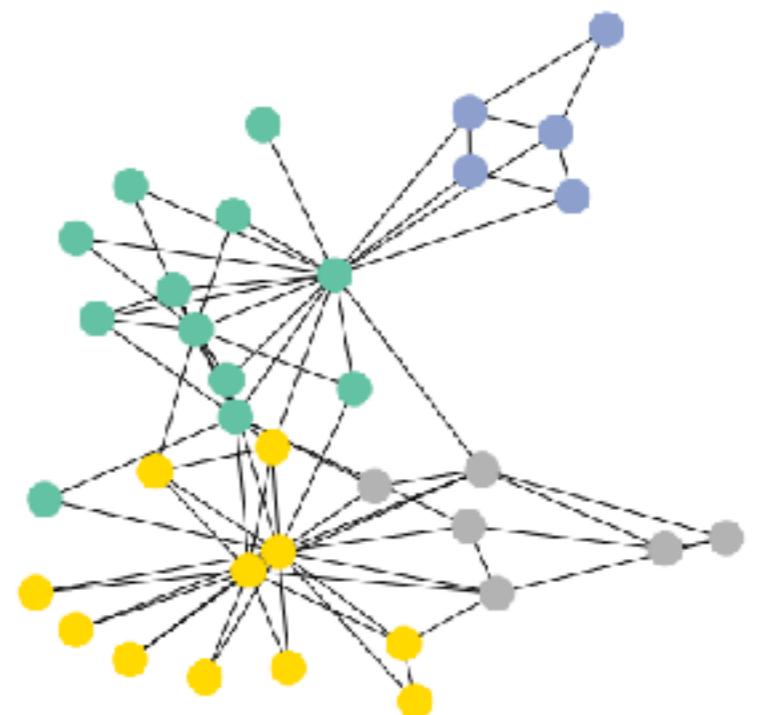


pypi package 1.6.1 build passing docs passing codecov 87% contributions welcome

[Documentation](#) | [Paper](#) | [Colab Notebooks](#) | [External Resources](#) | [OGB Examples](#)

## Exercise:

Run through and complete  
the Point Cloud Classification Colab





**Graph Neural Networks are one of the hottest research fields in deep learning!**

- ▶ GNNs learn **localized node embeddings** by following a simple message passing scheme
- ▶ Exciting research opportunities  
**Expressivity, Pooling, Scalability, ...**
- ▶ GNNs can help us to reason about **relational structure** in our world
- ▶ Exciting applications in  
**computer vision, NLP, biology, chemistry, physics, ...**



# Tutorial on Graph Neural Networks

Workshop on Geometry and Machine Learning with Applications to  
Biomedical Engineering - University College London



Matthias Fey

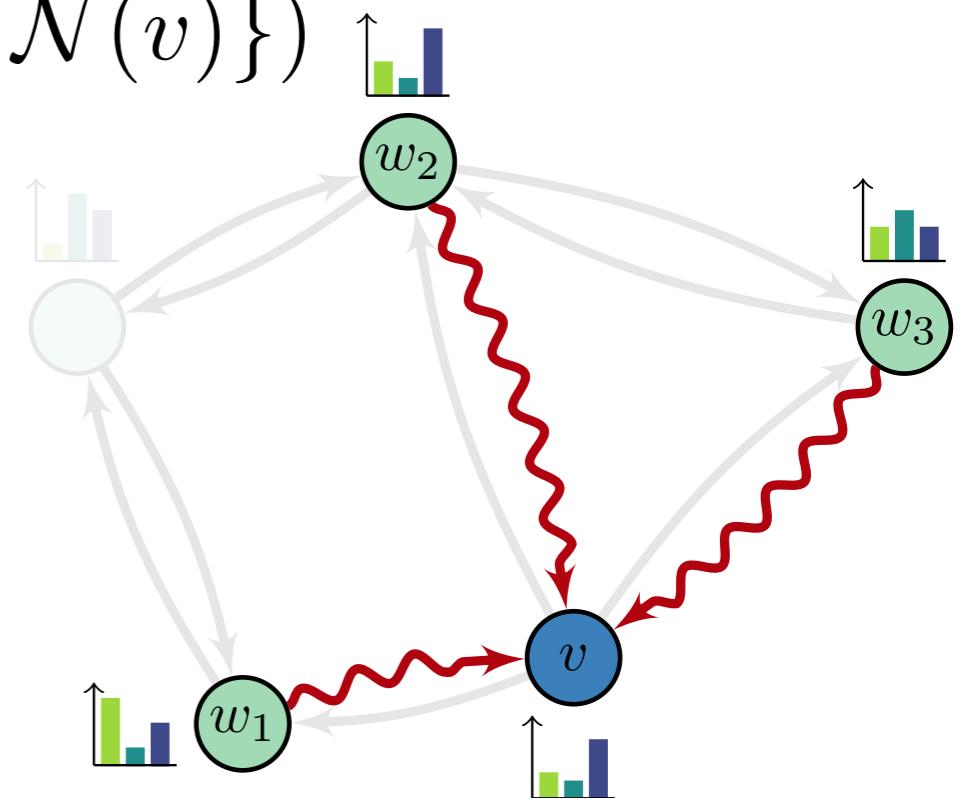
- ✉ matthias.fey@udo.edu
- ☁ rusty1s.github.io
- 🐱 🐦 /rusty1s

- ▶ GNNs learn localized node embeddings by following a simple message passing scheme

$$\mathbf{m}_{w,v}^{(\ell)} = \text{MESSAGE}(\mathbf{h}_v^{(\ell-1)}, \mathbf{h}_w^{(\ell-1)})$$

$$\mathbf{a}_v^{(\ell)} = \text{AGGREGATE}(\{\mathbf{m}_{w,v}^{(\ell)} : w \in \mathcal{N}(v)\})$$

$$\mathbf{h}_v^{(\ell)} = \text{UPDATE}(\mathbf{h}_v^{(\ell-1)}, \mathbf{a}_v^{(\ell)})$$

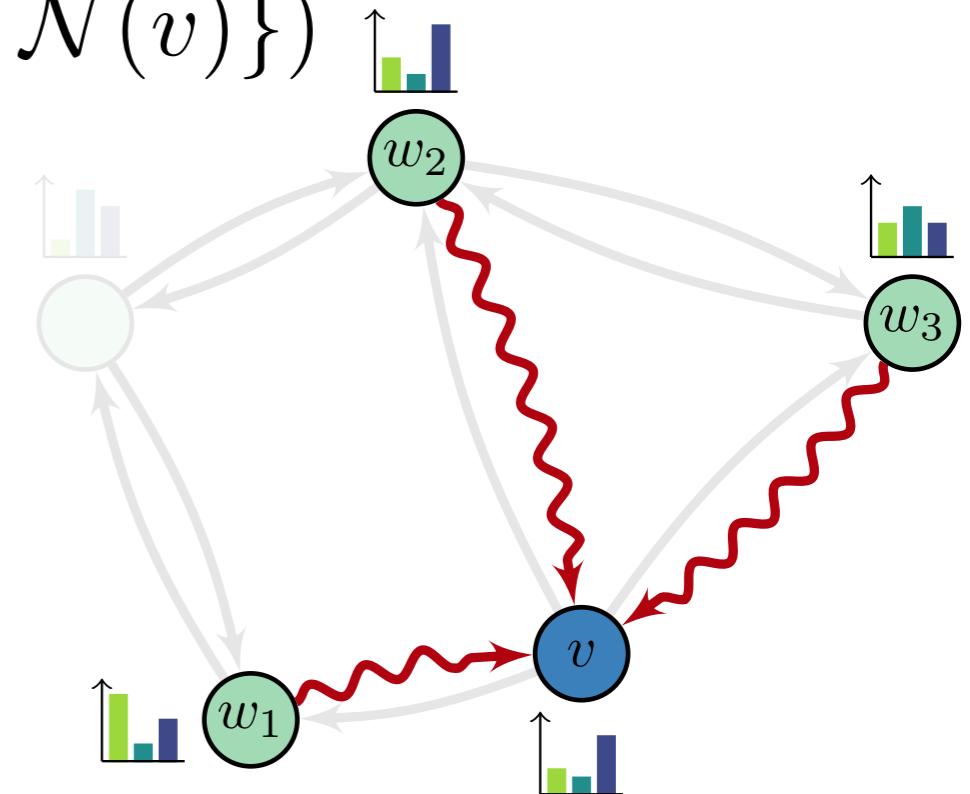
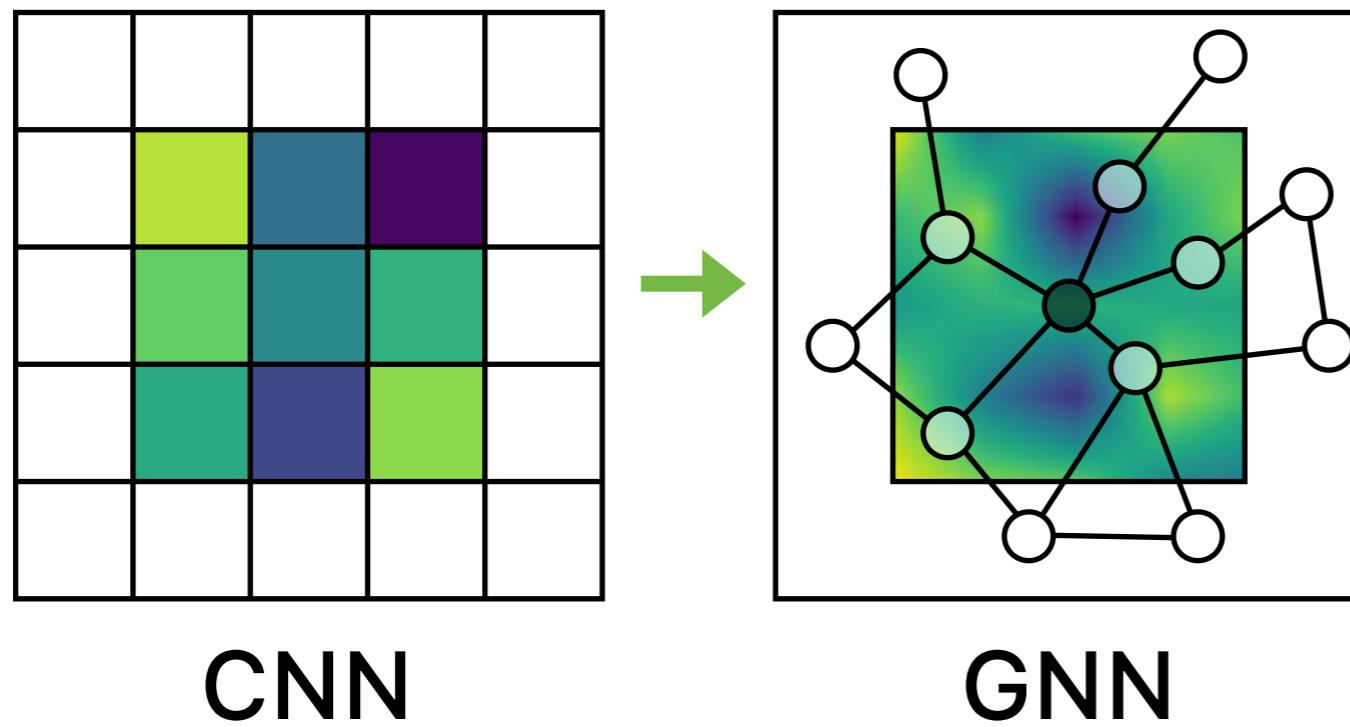


- ▶ GNNs learn localized node embeddings by following a simple message passing scheme

$$\mathbf{m}_{w,v}^{(\ell)} = \text{MESSAGE}(\mathbf{h}_v^{(\ell-1)}, \mathbf{h}_w^{(\ell-1)})$$

$$\mathbf{a}_v^{(\ell)} = \text{AGGREGATE}(\{\mathbf{m}_{w,v}^{(\ell)} : w \in \mathcal{N}(v)\})$$

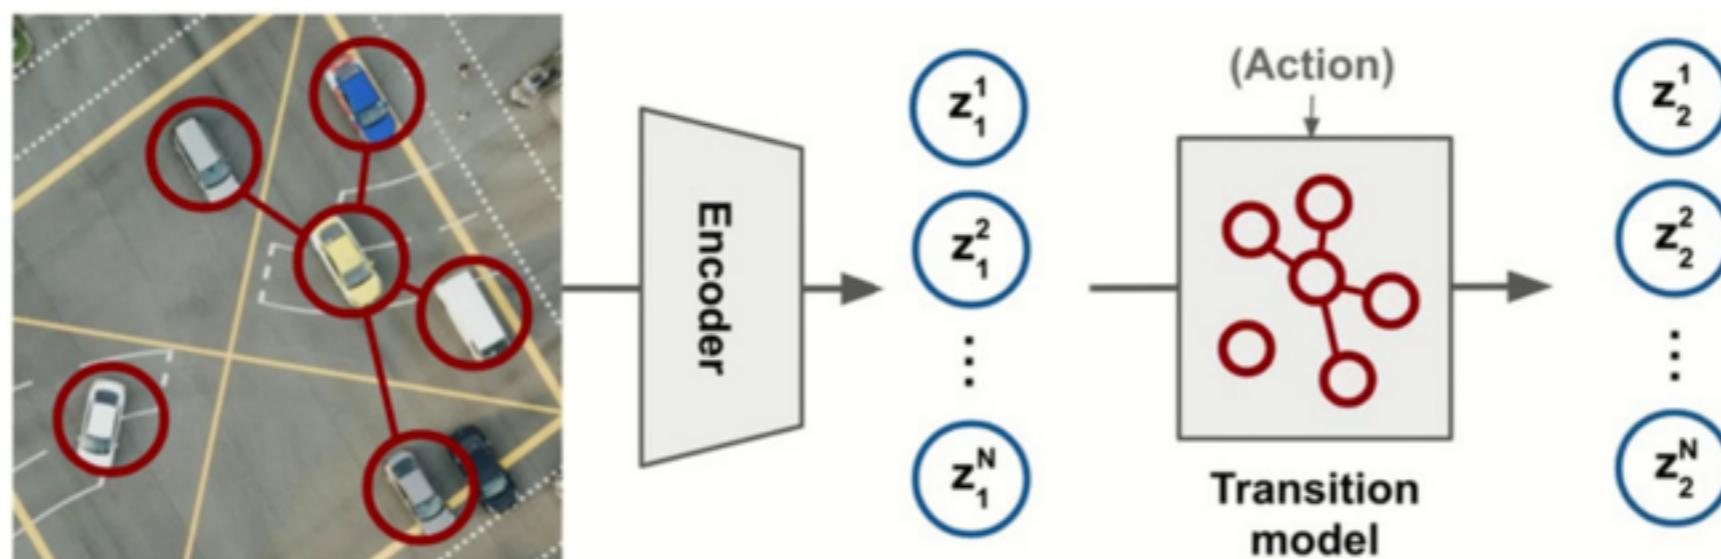
$$\mathbf{h}_v^{(\ell)} = \text{UPDATE}(\mathbf{h}_v^{(\ell-1)}, \mathbf{a}_v^{(\ell)})$$



- ▶ GNNs can help us to reason about relational structure in our world

*A small step towards more general AI?*

- ▶ Combination of CNNs and GNNs for "object-centric" learning





- ▶ Exciting research opportunities  
**Expressivity, Pooling, Scalability, ...**

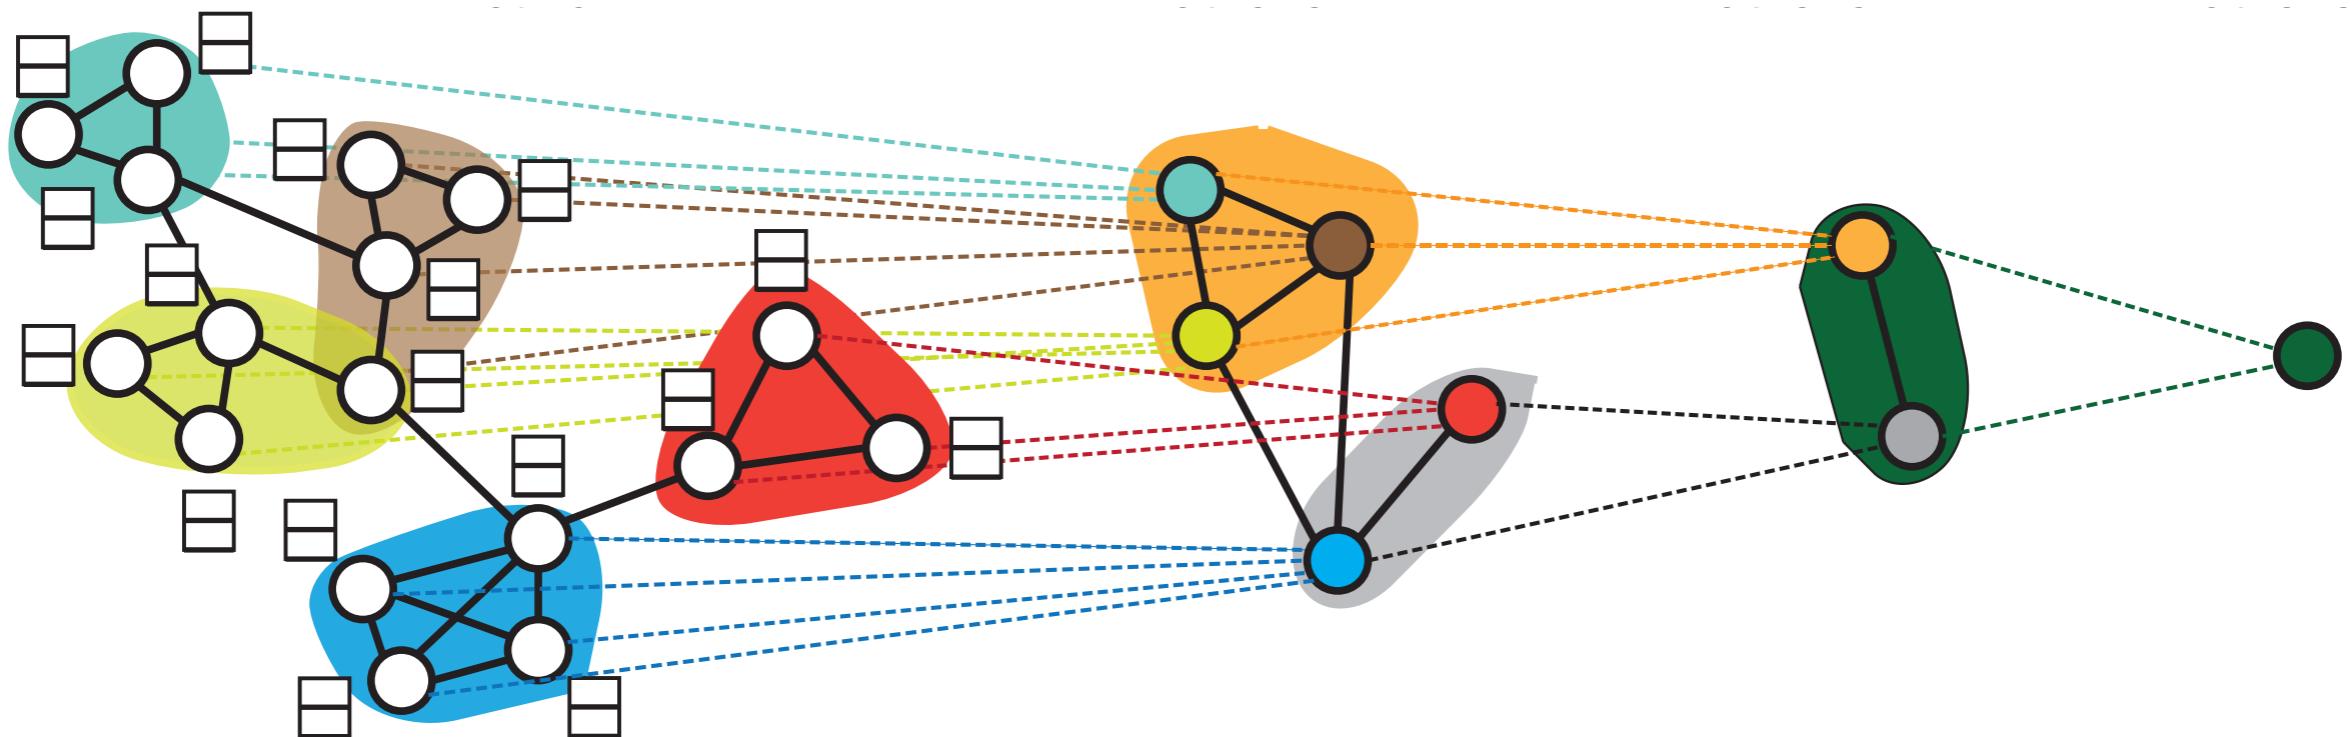
*How powerful are Graph Neural Networks in distinguishing graph structures?*

- ▶ Different aggregations yield different power
- ▶ Higher-order variants

- ▶ Exciting research opportunities  
Expressivity, Pooling, Scalability, ...

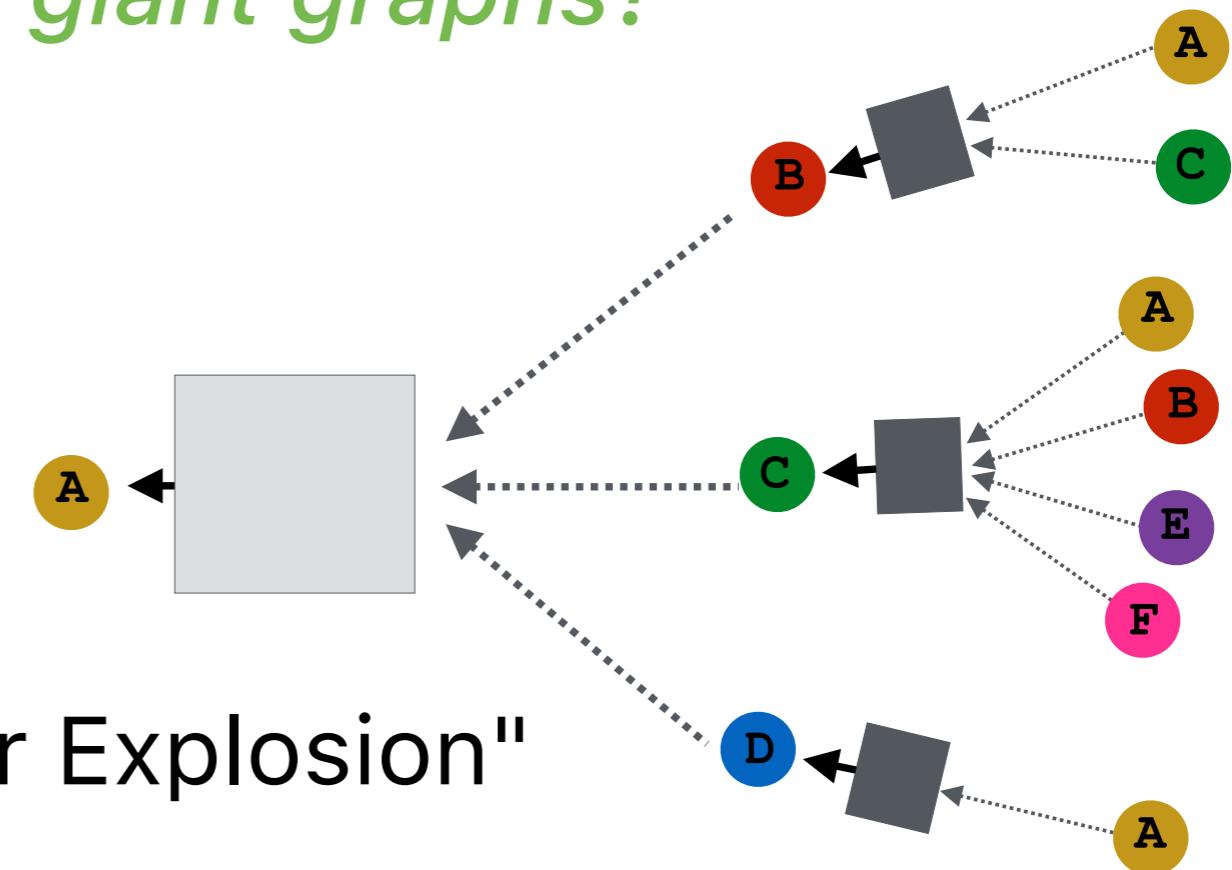
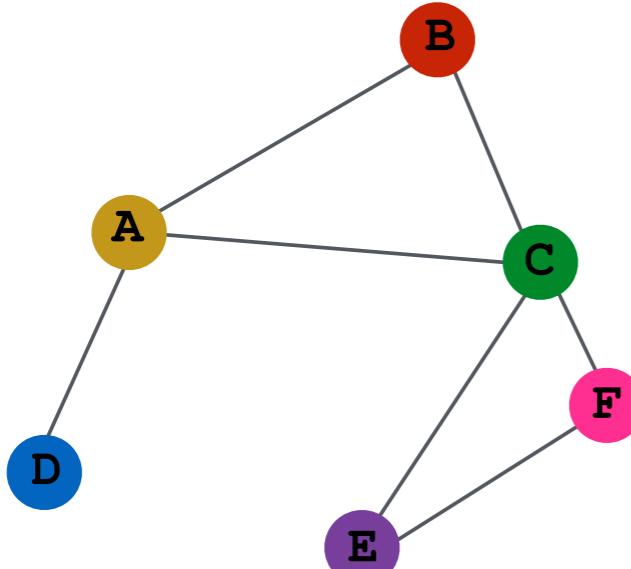
*Can we reason about hierarchies in graphs?*

*Can we generalize the concepts of pooling  
known from classical CNNs?*



- ▶ Exciting research opportunities  
Expressivity, Pooling, Scalability, ...

*How can we train GNNs in a mini-batch fashion  
on single giant graphs?*





- ▶ Exciting research opportunities  
Expressivity, Pooling, Scalability, ...

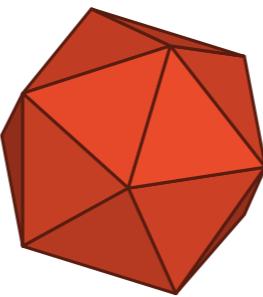
*There are still a lot of other things to explore!*

- ▶ Temporal GNNs
- ▶ Pre-Training of GNNs
- ▶ Generative models for graphs
- ▶ ...



- ▶ GNNs can be implemented efficiently

```
class GNN(torch.nn.Module):  
    def __init__(self):  
        super(GNN, self).__init__()  
        self.conv1 = GCNConv(128, 256)  
  
    def forward(self, x, edge_index):  
        x = self.conv1(x, edge_index)  
        x = x.relu()  
  
        ...
```



PyTorch  
geometric

colab

license MIT

PRs welcome

/rusty1s/pytorch\_geometric



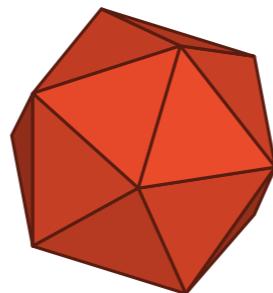
- ▶ GNNs can be implemented efficiently

```
class GNN(torch.nn.Module):  
    def __init__(self):
```

```
def
```

```
    x = x.relu()
```

```
    ...
```



PyTorch  
geometric

*Any Questions?*

license MIT

PRs welcome

[https://github.com/rusty1s/pytorch\\_geometric](https://github.com/rusty1s/pytorch_geometric)